

## NATIONAL UNIVERSITY OF SINGAPORE

SCHOOL OF COMPUTING

MIDTERM ASSESSMENT FOR  
Semester 2 AY2018/19

CS3243: INTRODUCTION TO ARTIFICIAL INTELLIGENCE

March 7, 2019

Time Allowed: 1 Hour 30 Minutes

---

**INSTRUCTIONS TO CANDIDATES**

1. This examination paper contains **THREE (3)** parts and comprises **13** printed pages, including this page.
2. Answer **ALL** questions as indicated. Unless explicitly said otherwise, you **must explain your answer**. Unexplained answers will be given a mark of 0.
3. Use the space provided to write down your solutions. If you need additional space to write your answers, we will provide you with draft paper. clearly write down your **student number** and **question number** on the draft paper, and attach them to your assessment paper. Make sure you clearly indicate on the **body of your paper as well** if you used draft paper to answer any question.
4. This is a **CLOSED BOOK** examination.
5. Please fill in your **Matriculation Number** below; **DO NOT WRITE YOUR NAME**.
6. For your convenience, we include an overview section of important definitions and algorithms from class at the end of this paper.

MATRICULATION NUMBER: \_\_\_\_\_

---

EXAMINER'S USE ONLY		
Part	Mark	Score
I	12	
II	18	
III	20	
TOTAL	50	

In Parts I, II, and III, you will find a series of short essay questions. For each short essay question, give your answer in the reserved space in the script.

## Part I

### Uninformed and Informed Search

(12 points) Short essay questions. Answer in the space provided on the script.

Consider the following game called **jumping Jack**: a player (called Jack) needs to jump from one platform to another, in order to reach a door. Jack may only jump to other platforms (or he will fall, which will make him lose!), and may do so under the following rules:

- Jack may only jump to platforms that are of distance 0 or 1 horizontally.
- Jack may only jump to platforms that are no more than one step higher than him; he can also jump to platforms on the same height, or that are arbitrarily lower than him.

See Figure I.1 for an illustration of legal jumps leading from Jack's initial position to a goal node. The cost of




Figure I.1: An illustration of the Jumping Jack game. The goal door is marked with  $G$ . The dotted arrows mark one legal path from Jack's initial position to the goal; note that this is not the shortest path!

jumping from one platform to another is always 1.

- (3 points) Formalize the Jumping Jack game as a search problem; what are the states? What are the actions? What is the cost function? You do not need to explain your answers.

**Solution:** A state should be the platform that Jack is on; actions are valid jumps from one platform to another; cost is 1 always.

- (9 points) Design two **non-trivial admissible** heuristics for the Jumping Jack game; by non-trivial we mean the heuristic you propose cannot be the all-zero heuristic, nor the optimal heuristic. **Your heuristics should be designed for a general instance of the game, not the specific illustration given in Figure I.1.**

Do your heuristics dominate one another? If so, explain why, if not - provide a counterexample.

**Solution:** We can take many options here, but one easy option is to play with horizontal distance, i.e.  $h(v)$  is the horizontal distance from  $v$  to the goal. Note that vertical distance is **not admissible** since Jack can drop arbitrarily far, and it is never mentioned that the goal must be above the initial position. A variation here can be that  $h(v)$  is the horizontal distance from  $v$  to the goal, plus 1 if they're not on the same height. A yet more refined variation can be that  $h(v)$  is the horizontal distance from  $v$  to the goal, plus 1 if  $v$  is above the goal, and plus the vertical distance if it's below the goal. We can also have rule relaxation heuristics here: heuristics that result from relaxing the rules of the game.

## Part II

### Adversarial Search

(18 points) Short essay questions. Answer in the space provided on the script.

1. (10 points) Consider the minimax search tree shown in Figure II.1. The MAX player controls the black nodes ( $s$ ,  $b_1$  and  $b_2$ ) and the MIN player controls the white nodes  $a_1$  and  $a_2$ . Square nodes are terminal nodes with utilities specified with respect to the MAX player. Suppose that we use the  $\alpha$ - $\beta$  pruning algorithm (reproduced in Figure IV.7), in the direction from **right to left** to prune the search tree. Complete the values for  $A$ ,  $B$  and




Figure II.1: Extensive form game with missing values.

$C$  in the terminal nodes in order to make the statement below true. You may assume that  $A$ ,  $B$  and  $C$  are **positive integers**.

- (a) (8 points) In order to ensure that **no arcs are pruned**, it must be the case that:

**Solution:**  $A > 7, B < C, C > 8$ . The MAX player can choose to be paid 7 by taking the rightmost arc. Therefore, in order to continue exploring  $a_2$ , it must be that  $A > 7$  (else the MIN player can guarantee  $\beta(a_2) \leq A \leq 7$  and the subtree rooted in  $a_2$  is pruned). Next, we note that  $\alpha(s) = 8$  after exploring the righthand side (assuming that  $A > 7$  and that  $A$  is an integer,  $A \geq 8$ ). Therefore, after exploring  $b_2$ , if we see  $C \leq 8$  there is no point in exploring the subtree rooted in  $a_1$  further. Finally, we note that after exploring the tree rooted in  $b_2$  we have  $\alpha(b_2) = C$ , and therefore  $\beta(a_1) = C$  as well; in order to check the terminal node of 11, we must have that  $B < C$ . If  $B \geq C$  then  $\alpha(b_1) \geq B \geq C$  and there is no point in checking the subtree rooted in  $b_1$  further.

- (b) (2 points) Under the values you found for  $A$ ,  $B$  and  $C$  above, what is the payoff to player 1 in a sub-perfect Nash equilibrium for this game, assuming that  $B \leq 11$ ? You do not need to explain your answer.

**Solution:** As computed above,  $A > 7, B < C, C > 8$ . We have  $\beta(a_1) = \min\{C, \max\{B, 11\}\}$ : either the MIN player picks  $b_2$  and pays  $C$ , or they pick  $b_1$  and pay the maximum of  $B$  and 11. Since  $B \leq 11$  this boils down to  $\min\{C, 11\}$ .





Figure II.2: Pursuer-evader graph; the pursuer starts in position  $s_1$  and the evader starts in position  $s_2$ .

2. (8 points) Consider a two player game described as follows: we have a *pursuer* (player 1, the MAX player) and an *evader* (player 2, the MIN player). The evader wants to not be caught, and the pursuer tries to catch them. Both players start in different positions on an undirected graph  $G = \langle V, E \rangle$ , denoted  $s_1$  (pursuer's start position) and  $s_2$  (evader's start position). Both players take turns to make moves, with the pursuer making the first move. If a player is in node  $v \in V$ , they can move to any of the neighbors of  $v$ . The pursuer catches the evader if they both share a node, in which case the pursuer gets 1 point and the evader gets  $-1$ . If the evader manages to not get caught after making  $k$  steps<sup>1</sup>, then the evader wins and gets 1 point (the pursuer gets  $-1$  points). **if the evader gets caught on the  $k$ -th step, they lose.** Consider the graph in Figure II.2, where the pursuer starts in position  $s_1$  and the evader starts in position  $s_2$ . Suppose that  $k = 2$ . Draw the extensive form game tree that results. In every node, write the position of the pursuer, the position of the evader, and the total number of steps that the evader has taken so far. Mark down the edges corresponding to actions in sub-perfect Nash equilibria on your game tree (by writing 'NE' next to them).

**Solution:** Here the orange nodes are controlled by the pursuer, and the blue by the evader. White nodes are terminal nodes.

<sup>1</sup>For example, if  $k = 3$  then the evader needs to move 3 times without being caught



Note that the game has *two SPNE*, since the evader is indifferent between getting caught in the beginning or getting caught later. The counter  $\ell$  is the number of steps that the evader has taken so far.

## Part III

### Constraint Satisfaction Problems

(20 points) Short essay questions. Answer in the space provided on the script.

1. (10 points) In the *vertex cover* problem we are given a graph  $G = \langle V, E \rangle$ . We say that a vertex  $v$  *covers* an edge  $e \in E$  if  $v$  is incident on the edge  $e$ . We are interested in finding a *vertex cover*; this is a set of vertices  $V' \subseteq V$  such that every edge is covered by some vertex in  $V'$ . In what follows you may **only** use variables of the form  $x_v$  where  $x_v = 1$  if  $v$  is part of the vertex cover, and is 0 otherwise. When writing the constraints you may **only use**

- Standard mathematical operators ( $+$ ,  $-$ ,  $\times$ ,  $\div$ ) and
- Standard logical and set operators ( $\forall$ ,  $\exists$ ,  $\vee$ ,  $\wedge$  and  $x \in X$ ,  $X \subseteq Y$ )

- (a) (5 points) Write down constraints that ensure that the condition “every edge  $e \in E$  is covered by at least one vertex  $v \in V$ ” holds, and a constraint that ensures that the total size of the vertex cover is no more than  $k$ . You may assume that edges are represented as sets of two vertices, i.e.  $e = \{u, v\}$  for some  $u, v \in V$ . You **need to justify** your answer here.

**Solution:**  $\forall \{u, v\} \in E : x_u + x_v \geq 1$  and  $\sum_{v \in V} x_v \leq k$ .

- (b) (5 points) Consider the graph in Figure III.1. In what follows, If you have two constraints of the form




Figure III.1: Graph for Vertex Cover CSP part (b)

$A \leq B$  and  $A \geq B$  you may merge them to a single constraint of the form  $A = B$ .

- (i) (2 points) Write down the vertex cover CSP induced by Figure III.1 with a vertex cover size  $k = 1$ . Use binary constraints to describe the constraint “the vertex cover must be of size 1”.

**Solution:** Since the size constraint is set to  $k = 1$  and we use binary constraints, it means that  $x_i + x_j \leq 1$  for all  $i, j \in \{1, \dots, 4\}$ . We get

$$x_1 + x_2 \geq 1$$

$$x_1 + x_3 \geq 1$$

$$x_1 + x_4 \geq 1$$

$$x_1 + x_2 \leq 1$$

$$x_1 + x_3 \leq 1$$

$$x_1 + x_4 \leq 1$$

$$x_2 + x_3 \leq 1$$

$$x_2 + x_4 \leq 1$$

$$x_3 + x_4 \leq 1$$

which narrows down to

$$x_1 + x_2 = 1$$

$$x_1 + x_3 = 1$$

$$x_1 + x_4 = 1$$

$$x_2 + x_3 \leq 1$$

$$x_2 + x_4 \leq 1$$

$$x_3 + x_4 \leq 1$$

- (ii) (3 points) Suppose that in our backtracking search we began by setting the value of  $X_1 = 0$  (i.e. we say that the vertex 1 is not part of our vertex cover), and then run the AC3 algorithm. What will the AC3 algorithm do? Assume that edges are processed in increasing lexicographic order, i.e. the queue is initialized to

$$(x_1, x_2), (x_1, x_3), (x_1, x_4), (x_2, x_1), (x_2, x_3), (x_2, x_4), \dots, (x_4, x_1), (x_4, x_2), (x_4, x_3).$$

You must explain your answer, but you do not have to provide a complete trace of the algorithm.

**Solution:** It is quite clear to us that the only solution here is to set  $x_1 = 1$  and the rest to 0. What will AC3 do if we assign  $x_1 = 0$ ? It begins by processing the edges of the form  $(x_1, x_j)$ . For each one of them, it does nothing, since the domains of  $x_2, x_3$  and  $x_4$  have values that make the constraint with  $x_1$  consistent. When reaching the edge  $(x_2, x_1)$  it will have the constraint  $x_1 + x_2 = 1$ , which can only be satisfied by reducing  $x_2$ 's domain to  $\{1\}$ . It will do nothing for  $(x_2, x_3); (x_2, x_4)$ . For  $(x_3, x_1)$  it will again reduce the domain of  $x_3$  to  $\{1\}$  for a similar reason (and add edges to the end of the queue). However, when reaching  $(x_3, x_2)$  we will have a problem: no value in the reduced domain of  $x_3$  will satisfy the constraint  $x_2 + x_3 \leq 1$ . Thus, the domain of  $x_3$  is reduced to  $D_3 = \emptyset$  after the  $\text{REVISE}(x_3, x_2)$  function is called, it will exit the  $\text{REVISE}(x_3, x_2)$  procedure; the subsequent check of  $D_3$ 's size will result in AC3 returning `false`.

2. (10 points) We have seen in class that if a CSP has a constraint graph that is a tree, it is possible to compute a satisfying assignment (or decide that one does not exist) in  $\mathcal{O}(nd^2)$  time, where the number of variables is  $n$ , and the maximal size of variable domains is  $d$ . Consider a binary CSP whose constraint graph is *nearly a tree*: there exists one variable (without loss of generality assume that it is  $X_1$ ) such that removing  $X_1$  from the constraint graph will result in a tree-structured constraint graph; in other words, the constraint graph induced by the variables  $X_2, \dots, X_n$  is a tree, but the one induced by the variables  $X_1, \dots, X_n$  is not. Prove that there is a polynomial-time algorithm that can find a satisfying assignment (or decide that one does not exist) for nearly-tree structured CSPs.

**Solution:** The solution is simple: the CSP will have a satisfying assignment if and only if for some value  $x \in X_1$  there is some satisfying assignment for the variables  $X_2, \dots, X_n$ . So, fixing the value of  $X_1$  to be some arbitrary  $x$ , we run the following procedure

1. For every  $x \in D_1$ :
  - (a) Reduce the domains  $D_2, \dots, D_n$  so that they contain only values that are consistent with  $X_1$ , i.e. for every  $X_j$ , the constraint  $C(x, y)$  is satisfied for every value  $y \in D_j$ . This can be done via forward checking in  $\mathcal{O}(nd)$  time.
  - (b) If none of the domains of  $X_2, \dots, X_n$  are empty, run the poly-time algorithm for tree structured CSPs on the reduced domains for  $X_2, \dots, X_n$ . This takes  $\mathcal{O}(nd^2)$  time.
  - (c) If the algorithm returns a satisfying assignment  $(y_2, \dots, y_n)$ , return  $(x, y_2, \dots, y_n)$
2. return `false`.

The algorithm's runtime is thus  $\mathcal{O}(nd^3)$ . If the algorithm returns an assignment, it must satisfy the CSP: all of the constraints of the form  $C(X_1, X_j)$  are guaranteed to be satisfied, because we only use values for  $X_j$  that satisfy  $C(x, X_j)$ . The other constraints only involve constraints of the form  $C(X_j, X_k)$  where  $j, k \neq 1$ , so we get a satisfying assignment by the tree algorithm. Suppose that there is some satisfying assignment to the CSP, say  $(x_1, x_2, \dots, x_n)$ ; our algorithm is guaranteed to find one: if it hasn't found any assignment up to the point that  $X_1$  was set to  $x_1$ , it will find one on the iteration where  $X_1$  is set to  $x_1$ : the tree algorithm will only consider variable assignments that agree with constraints involving  $X_1 = x_1$  (one of which is  $(x_2, \dots, x_n)$ ) and since the tree algorithm is guaranteed to find a satisfying assignment when one exists, it will output either  $(x_2, \dots, x_n)$  or some other assignment that agrees with all constraints when  $X_1 = x_1$ .



## Part IV

### Summary of Course Material

The following summary contains key parts from the course lecture notes. It **does not** offer a complete coverage of course materials. You may use any of the claims shown here without proving them, unless specified explicitly in the question.

## 1 Search Problems

### 1.1 Uninformed Search

We are interested in finding a solution to a fully observable, deterministic, discrete problem. A search problem is given by a set of *states*, where a *transition function*  $T(s, a, s')$  states that taking action  $a$  in state  $s$  will result in transitioning to state  $s'$ . There is a cost to this defined as  $c(s, a, s')$ , assumed to be non-negative. We are also given an initial state (assumed to be in our frontier upon initialization). The *frontier* is the set of nodes in a queue that have not yet been explored, but will be explored given the order of the queue. We also have a *goal test*, which for a given state  $s$  outputs “yes” if  $s$  is a goal state (there can be more than one). We have discussed two search

```

function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier

```

Figure IV.1: The tree search algorithm

```

function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    only if not in the frontier or explored set

```

Figure IV.2: The graph search algorithm

variants in class, tree search (Figure IV.1) and graph search (Figure IV.2). They differ in the fact that under graph search we do not explore nodes that we have seen before.

The main thing that differentiates search algorithms is the order in which we explore the frontier. In breadth-first search we explore the shallowest nodes first; in depth-first search we explore the deepest nodes first; in uniform-cost search (Figure IV.3) we explore nodes in order of cost.

```

function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier ← a priority queue ordered by PATH-COST, with node as the only element
  explored ← an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node ← POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child ← CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier ← INSERT(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child

```

Figure IV.3: The uniform cost search algorithm

Property	BFS	UCS	DFS	DLS	IDS
Complete	Yes	Yes	No	No	Yes
Optimal	No	Yes	No	No	No
Time	$\mathcal{O}(b^d)$	$\mathcal{O}(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$	$\mathcal{O}(b^m)$	$\mathcal{O}(b^\ell)$	$\mathcal{O}(b^d)$
Space	$\mathcal{O}(b^d)$	$\mathcal{O}(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$	$\mathcal{O}(bm)$	$\mathcal{O}(b\ell)$	$\mathcal{O}(bd)$

Table 1: Summary of search algorithms' properties

We have also studied variants where we only run DFS to a certain depth (Figure IV.4) and where we iteratively deepen our search depth (Figure IV.5). Table 1 summarizes the various search algorithms' properties.

```

function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff
return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff
if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
else if limit = 0 then return cutoff
else
  cutoff-occurred?  $\leftarrow$  false
  for each action in problem.ACTIONS(node.STATE) do
    child  $\leftarrow$  CHILD-NODE(problem, node, action)
    result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true
    else if result  $\neq$  failure then return result
  if cutoff-occurred? then return cutoff else return failure

```

Figure IV.4: The depth-limited search algorithm

```

function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
for depth = 0 to  $\infty$  do
  result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
  if result  $\neq$  cutoff then return result

```

Figure IV.5: The iterative deepening search algorithm

We also noted that no deterministic search algorithm can do well in general; in the worst case, they will all search through the entire search space.

## 1.2 Informed Search

Informed search uses additional information about the underlying search problem in order to narrow down the search scope. We have mostly discussed the  $A^*$  algorithm, where the priority queue holding the frontier is ordered by  $g(v) + h(v)$ , where  $g(v)$  is the distance of a node from the source, and  $h(v)$  is a *heuristic estimate* of the distance of the node from a goal node. There are two key types of heuristic functions

**Definition 1.1.** A heuristic  $h$  is *admissible* if it never overestimates the distance of a node from the nearest goal; i.e.

$$\forall v : h(v) \leq h^*(v),$$

where  $h^*(v)$  is the *optimal heuristic*, i.e. the true distance of a node from the nearest goal.

**Definition 1.2.** A heuristic  $h$  is *consistent* if it satisfies the triangle inequality; i.e.

$$\forall v, v' : h(v) \leq h(v') + c(v, v')$$

where  $c(v, v')$  is the cost of transitioning from  $v$  to  $v'$ .

We have shown that  $A^*$  with tree search is optimal when the heuristic is admissible, and is optimal with graph search when the heuristic is consistent. We also showed that consistency implies admissibility but not the other way around, and that running  $A^*$  with an admissible inconsistent heuristic with graph search may lead to a sub-optimal goal.

## 2 Adversarial Search

An extensive form game is defined by  $V$  a set of nodes and  $E$  a set of directed edges, defining a tree. The root of the tree is the node  $r$ . Let  $V_{\max}$  be the set of nodes controlled by the MAX player and  $V_{\min}$  be the set of nodes controlled by the MIN player. We often refer to the MAX player as player 1, and to the MIN player as player 2. Furthermore, let  $\text{desc}(v)$  be the descendants of a node  $v$  in the search tree. A *strategy* for the MAX player is a mapping  $s_1 : V_{\max} \rightarrow V$ ; similarly, a strategy for the MIN player is a mapping  $s_2 : V_{\min} \rightarrow V$ . In both cases,  $s_i(v) \in \text{chld}(v)$  is the choice of child node that will be taken at node  $v$ . We let  $\mathcal{S}_1, \mathcal{S}_2$  be the set of strategies for the MAX and MIN player, respectively.

The leaves of the minimax tree are *payoff nodes*. There is a payoff  $a(v)$  associated with each payoff node  $v$ , where  $a(v)$  is the amount that the MAX player receives and  $-a(v)$  is the amount that the MIN player receives if the node  $v$  is reached. The MAX player wants to maximize their payoff, and the MIN player wants to minimize the amount that they pay. More formally, the utility of the MAX player from  $v$  is  $u_{\max}(v) = a(v)$  and the utility of the MIN player is  $u_{\min}(v) = -a(v)$ . The utility of a player from a pair of strategies  $s_1 \in \mathcal{S}_1, s_2 \in \mathcal{S}_2$  is simply the utility they receive by the leaf node reached when the strategy pair  $(s_1, s_2)$  is played.

**Definition 2.1** (Nash Equilibrium). A pair of strategies  $s_1^* \in \mathcal{S}_1, s_2^* \in \mathcal{S}_2$  for the MAX player and the MIN player, respectively, is a *Nash equilibrium* if no player can get a strictly higher utility by switching their strategy. In other words:

$$\begin{aligned} \forall s \in \mathcal{S}_1 : u_1(s_1^*, s_2^*) &\geq u_1(s, s_2^*); \\ \forall s' \in \mathcal{S}_2 : u_2(s_1^*, s_2^*) &\geq u_2(s_1^*, s') \end{aligned}$$

**Definition 2.2** (Subgame). Given an extensive form game  $\langle V, E, r, V_{\max}, V_{\min}, \vec{a} \rangle$ , a subgame is a subtree of the original game, defined by some arbitrary node  $v$  set to be the root node, and all of its descendants (i.e. its children, its children's children etc.), denoted by  $\text{desc}(v)$ . Leaf nodes and node control are the same as in the original extensive form game.

**Definition 2.3** (Subgame-Perfect Nash Equilibrium (SPNE)). A pair of strategies  $s_1^* \in \mathcal{S}_1, s_2^* \in \mathcal{S}_2$  is a sub-perfect Nash equilibrium if it is a Nash equilibrium for any subtree of the original game tree.

<b>function</b> MINIMAX-DECISION( <i>state</i> ) <b>returns</b> an action <b>return</b> $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$
<b>function</b> MAX-VALUE( <i>state</i> ) <b>returns</b> a utility value <b>if</b> TERMINAL-TEST( <i>state</i> ) <b>then return</b> UTILITY( <i>state</i> ) $v \leftarrow -\infty$ <b>for each</b> $a$ <b>in</b> ACTIONS( <i>state</i> ) <b>do</b> $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$ <b>return</b> $v$
<b>function</b> MIN-VALUE( <i>state</i> ) <b>returns</b> a utility value <b>if</b> TERMINAL-TEST( <i>state</i> ) <b>then return</b> UTILITY( <i>state</i> ) $v \leftarrow \infty$ <b>for each</b> $a$ <b>in</b> ACTIONS( <i>state</i> ) <b>do</b> $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$ <b>return</b> $v$

Figure IV.6: The minimax algorithm (note a typo from the AIMA book -  $s$  should be *state*).

Figure IV.6 describes the minimax algorithm, which computes SPNE strategies for the MIN and MAX players, as we have shown in class. We discussed the  $\alpha$ - $\beta$  pruning algorithm as a method of removing subtrees that need not be explored (Figure IV.7).

## 3 Constraint Satisfaction Problems

A constraint satisfaction problem (CSP) is given by a set of variables  $X_1, \dots, X_n$ , each with a corresponding domain  $D_1, \dots, D_n$  (it is often assumed that domain sizes are constrained, i.e.  $|D_i| \leq d$  for all  $i \in [n]$ ). Constraints specify relations between sets of variables; we are given  $C_1, \dots, C_m$  constraints. The constraint  $C_j$  depends on a subset of variables and takes on a value of “true” if and only if the values assigned to these variables satisfy

```

function ALPHA-BETA-SEARCH(state) returns an action
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 



---


function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow -\infty$ 
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 



---


function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 

```

Figure IV.7: The  $\alpha$ - $\beta$  pruning algorithm.

$C_j$ . Our objective is to find an assignment  $(y_1, \dots, y_n) \in D_1 \times \dots \times D_n$  of values to the variables such that  $C_1, \dots, C_m$  are all satisfied. A binary CSP is one where all constraints involve at most two variables. In this case, we can write the relations between variables as a *constraint graph*, where there is an (undirected) edge between  $X_i$  and  $X_j$  if there is some constraint of the form  $C(X_i, X_j)$ .

In class we discussed *backtracking search* (Figure IV.8), which is essentially a depth-first search assigning variable values in some order. Within backtracking search, we can employ several heuristics to speed up our search.

```

function BACKTRACKING-SEARCH(csp) returns a solution, or failure
  return BACKTRACK( $\{\}$ , csp)

function BACKTRACK(assignment, csp) returns a solution, or failure
  if assignment is complete then return assignment
   $\text{var} \leftarrow \text{SELECT-UNASSIGNED-VARIABLE}(\text{csp})$ 
  for each  $\text{value}$  in ORDER-DOMAIN-VALUES( $\text{var}$ , assignment, csp) do
    if  $\text{value}$  is consistent with assignment then
      add  $\{\text{var} = \text{value}\}$  to assignment
       $\text{inferences} \leftarrow \text{INFERENCE}(\text{csp}, \text{var}, \text{value})$ 
      if  $\text{inferences} \neq \text{failure}$  then
        add  $\text{inferences}$  to assignment
         $\text{result} \leftarrow \text{BACKTRACK}(\text{assignment}, \text{csp})$ 
        if  $\text{result} \neq \text{failure}$  then
          return  $\text{result}$ 
      remove  $\{\text{var} = \text{value}\}$  and  $\text{inferences}$  from assignment
  return failure

```

Figure IV.8: Backtracking search

1. When selecting the next variable to check, it makes sense to choose:
  - (a) the most constrained variable (the one with the least number of legal assignable values).
  - (b) the most constraining variable (the one that shares constraints with the most unassigned variables)
2. When selecting what value to assign, it makes sense to choose a value that is least constraining for other variables.

It also makes sense to keep track of what values are still legal for other variables, as we run backtracking search.

**Forward Checking:** as we assign values, keep track of what variable values are allowed for the unassigned variables. If some variable has no more legal values left, we can terminate this branch of our search.

**Arc Consistency:** a generalization of forward checking.

**Definition 3.1.** Given two variables  $X_i, X_j$ ,  $X_i$  is consistent with respect to  $X_j$  (equivalently, the arc  $(X_i, X_j)$  is consistent) if for any value  $x \in D_i$  there exists some value  $y \in D_j$  such that the binary constraint on  $X_i X_j$  is satisfied with  $x, y$  assigned, i.e.  $C_{i,j}(x, y)$  is satisfied.

The AC3 algorithm (Figure IV.9) offers a nice way of iteratively reducing the domains of variables in order to ensure arc consistency at every step of our backtracking search. It can be done before we choose a value to assign to a variable, or after. Whenever we remove a value from the domain of  $X_i$  with respect to  $X_j$  (the REVISE operation in the AC3 algorithm), we need to recheck all of the neighbors of  $X_i$ , i.e. add all of the edges of the form  $(X_k, X_i)$  where  $k \neq j$  to the checking queue of the AC3 algorithm. We have seen in class that the AC3 algorithm runs in  $\mathcal{O}(n^2d^3)$  time. In general, finding a satisfying assignment (or deciding that

```

function AC-3(csp) returns false if an inconsistency is found and true otherwise
inputs: csp, a binary CSP with components  $(X, D, C)$ 
local variables: queue, a queue of arcs, initially all the arcs in csp

while queue is not empty do
   $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$ 
  if REVISE(csp,  $X_i, X_j$ ) then
    if size of  $D_i = 0$  then return false
    for each  $X_k$  in  $X_i.\text{NEIGHBORS} - \{X_j\}$  do
      add  $(X_k, X_i)$  to queue
return true



---


function REVISE(csp,  $X_i, X_j$ ) returns true iff we revise the domain of  $X_i$ 
  revised  $\leftarrow$  false
  for each  $x$  in  $D_i$  do
    if no value  $y$  in  $D_j$  allows  $(x, y)$  to satisfy the constraint between  $X_i$  and  $X_j$  then
      delete  $x$  from  $D_i$ 
      revised  $\leftarrow$  true
  return revised

```

Figure IV.9: The AC3 Algorithm

one does not exist) for a CSP with  $n$  variables and domain size bounded by  $d$  takes  $\mathcal{O}(d^n)$  via backtracking search; however, we have seen in class that if the constraint graph is a tree (or a forest more generally), we can find one in  $\mathcal{O}(nd^2)$  time.

---

**END OF EXAM PAPER**

---