

# Informed Search: Incorporating Domain Knowledge

---

CS3243: Introduction to Artificial Intelligence – Lecture 3

25 January 2023 (Video Lecture)

# Contents

---

1. Administrative Matters
2. Reviewing Uninformed Search
3. Greedy Best-First Search
4. A\* Search
5. Dominant Heuristics

# Administrative Matters

---

# Project 1

---

- Released this week (Week 3)
  - Due Week 6 Sunday, 19 February, 2359 hrs (about a month to work on it)
- Consultation session
  - Week 4 Friday, 3 February, Time TBC
  - Online (Zoom) – will be recorded
    - Meeting link = Canvas announcement
  - Reviews FAQ + any new questions raised
- Late submissions
  - Penalties
    - Within deadline +24 hours = 80% of score
    - Within deadline +48 hours = 50% of score
    - Beyond deadline +48 hours = 0% of score

Live Coding Session: TBC  
(Probably also in Week 4)

Start on Project 1 early!  
Start today!

DO NOT COPY/SHARE CODE!

# Upcoming...

---

- Tutorials begin this week
  - Only sessions on Monday/Tuesday rescheduled
    - T03 (MON 1400 hrs) → WED 1400 hrs - [T03 Zoom Link \(Week 3\)](#)
    - T04 (MON 1500 hrs) → WED 0800 hrs - [T04 Zoom Link \(Week 3\)](#)
    - T05 (TUE 0900 hrs) → FRI 1400 hrs - [T05 Zoom Link \(Week 3\)](#)
- Deadlines
  - TA1 (released last week)
    - *Due in your Week 3 tutorial session*
    - *Submit the a physical copy (more instructions on the Tutorial Worksheet)*
    - *Those with tutorials on Zoom this week → submit via email to tutors*
  - Remember to prepare for the tutorial
    - Participation marks = 5%

# Reviewing Uninformed Search

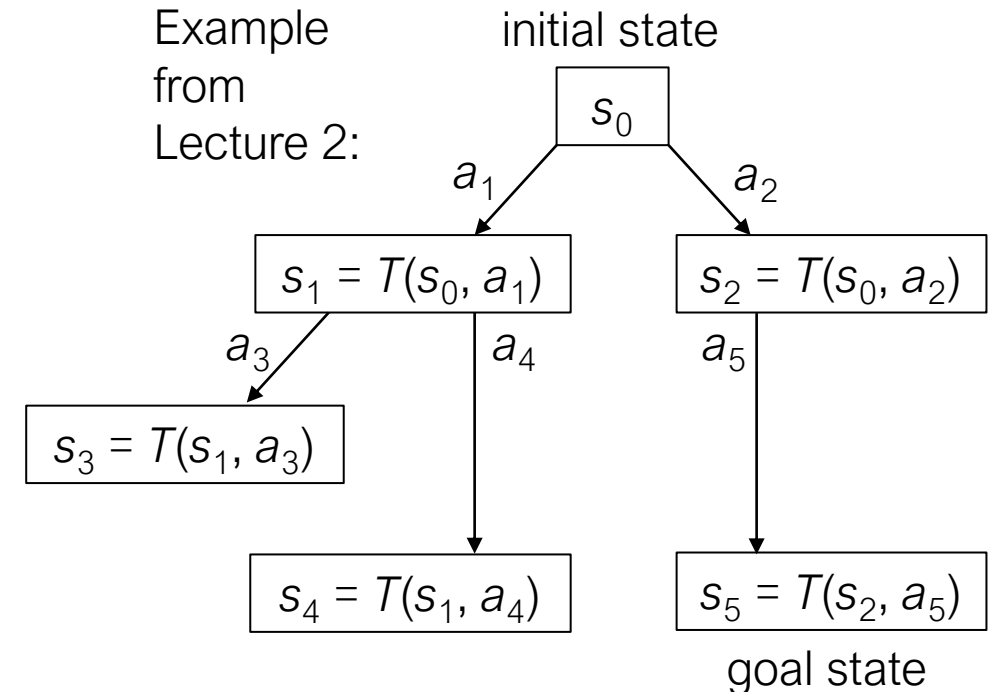
---

# Problem-solving Agent: General Path Planning

- Assumptions on problem environment
  - Fully observable
  - Deterministic
  - Discrete
  - Episodic
- General definition of a path planning problem
  - State representation,  $s_i$ 
    - Initial state ( $s_0$ )
  - Goal test,  $isGoal : s_i \rightarrow \{0, 1\}$   
where 1 denotes a goal state
  - Actions,  $actions : s_i \rightarrow A = \{a_1, \dots, a_k\}$
  - Action costs,  $cost : (s_i, a_j, s_i') \rightarrow v \geq 0$
  - Transition model,  $T : (s_i, a_j) \rightarrow s_i'$

General solution for ALL path planning problems – stronger AI (than reflex agent)

Example  
from  
Lecture 2:



# Tree Search Algorithm

---

```
frontier = {Node(initial state), NULL}
while frontier not empty:
    current = frontier.pop()
    if isGoal(current.state): return current.getPath()
    for a in actions(current.state):
        successor = Node(T(current.state, a), current)
        frontier.push(successor)
return failure
```

- Frontier
  - Paths to search
  - Considers ALL paths (including redundant ones)
- Concept of a Node for defining paths
  - Referenced *end point of path*, state *s*
  - Parent node, for *Linked List reference*
  - Other attributes
    - Action – facilitate *backtracking* in *DFS* – i.e.,  $O(m)$  space complexity
    - Path cost – for efficient *UCS* (avoids traversing path Linked List)
    - Depth – for efficient *DLS/IDS* (avoids traversing path Linked List)






# Uninformed Search Algorithms

- Systematic traversal of search space



- BFS: paths in order of number of actions taken
- DFS: paths to terminal state
- DLS/IDS: hybrid

- Systematic search allows completeness

- Assuming search space finite
- Assuming solution exists




...



# UCS and Optimality

- UCS traverses paths in order of path cost



- Ordered traversal of path costs required to ensure optimality

# Summary under tree search

- Performance of search algorithms

Criterion	BFS	UCS	DFS	DLS	IDS
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No <sup>3</sup>	No	Yes <sup>1</sup>
Optimal Cost?	Yes <sup>4</sup>	Yes	No	No	Yes <sup>4</sup>
Time	$O(b^d)$	$O(b^{1 + \lceil C^* / \epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1 + \lceil C^* / \epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$

1. Complete if  $b$  finite and either has a solution or  $m$  finite

2. Complete if all actions costs are  $> \epsilon > 0$

3. DFS is incomplete unless the search space is finite – i.e., when  $b$  finite and  $m$  finite

4. Cost optimal if action costs are all identical (and several other cases)

- Recall that an Early Goal Test is assumed for BFS
- UCS must perform a Late Goal Test to be optimal (this also accounts for the +1 in the index of its complexity)
- DFS is not complete (even under 1) as even if a solution exists, it may infinitely traverse a path without a solution (note the “or”)
- DFS space complexity may be improved to  $O(m)$  with backtracking (similar for DLS and IDS)

# Graph Search Algorithm (Version 1)

```
frontier = {Node(initial state), NULL}
reached = {initial state: Node(initial state)}
while frontier not empty:
    current = frontier.pop()
    if isGoal(current.state): return current.getPath()
    for a in actions(current.state):
        successor = Node(T(current.state, a), current)
        if successor.state not in reached:
            frontier.push(successor)
            reached.insert(successor.state: successor)
return failure
```

- Reached (sometimes also labelled: Visited)
  - Hash table that tracks all nodes already reached or visited via prior searching
- This version ensures that nodes are never revisited
  - Omits all redundant paths
  - May omit optimal path
  - No impact on completeness

# Graph Search Algorithm (Version 2)

---

```
frontier = {Node(initial state), NULL}
reached = {initial state: Node(initial state)}
while frontier not empty:
    current = frontier.pop()
    if isGoal(current.state): return current.getPath()
    for a in actions(current.state):
        successor = Node(T(current.state, a), current)
        if successor.state not in reached or
           successor.getCost() < reached[successor.state].getCost():
            frontier.push(successor)
            reached.insert(successor.state: successor)
return failure
```

- More relaxed constraint on paths that are considered
  - Also considers paths with lower path cost

# Tree Search Versus Graph Search

- Tree search will try ALL paths

- No paths excluded
- No issues with optimality


Completeness assumptions:

- $b$  finite, and  $m$  finite  
OR has a solution
- All action costs are  $> \epsilon > 0$

- What about graph search?

- Ensure optimal path not among excluded paths
- Consider this example

- $F = \{A(0)\}; R = \{A\}$ 
  - pop  $A(0)$ , push  $C(1)$  and  $G(3)$
- $F = \{C(1), G(3)\}; R = \{A, C, G\}$ 
  - pop  $C(1)$ , push  $G(2)$  since lower cost
- $F = \{G(2), G(3)\}; R = \{A, C, G\}$ 
  - pop  $G(2)$ , path is  $A \rightarrow C \rightarrow G$



Notice that without the update to  $G$  while it was on the frontier, we would not have returned the optimal path (similar to applying Early Goal Test)

# Graph Search Properties

- Performance under graph search implementations

Criterion	BFS	UCS	DFS	DLS	IDS
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No <sup>3</sup>	No	Yes <sup>1</sup>
Optimal Cost?	Yes <sup>4</sup>	Yes	No	No	Yes <sup>3</sup>
Time	$O( V  +  E )$				
Space					

1. Complete if  $b$  finite and either has a solution or  $m$  finite
2. Complete if all actions costs are  $> \epsilon > 0$
3. DFS is incomplete unless the search space is finite – i.e., when  $b$  finite and  $m$  finite
4. Cost optimal if action costs are all identical (and several other cases)

For CS3243, assume graph search Version 1 is used unless otherwise mentioned

**DONE? SOLVED?**

- Time and space complexities are now bounded by the size of the search space – i.e., the number of vertices (nodes) and edges,  $|V| + |E|$
- Note that we **do not** need to check for cheaper paths under graph search for BFS and DFS since costs play no part in those algorithms and they cannot guarantee an optimal solution anyway

# Informed Search

---




# Searching Less?

- Uninformed search algorithms are systematic
  - Search outward from the initial state
  - All directions
- Search spaces are LARGE
- Can we search more efficiently?


## General idea

Use domain knowledge about the problem environment to determine the cost required to go from a particular state to its nearest goal



# Path Costs & Heuristics


- UCS
  - Frontier = Priority Queue
  - Priority for node  $n = g(n)$ 
    - $g(n)$  quantifies the path cost from the initial state  $S$  to  $n.state$  as defined by  $n.path$
- General idea: use domain knowledge to estimate cost from  $n.state$  to  $G$
- Define a *heuristic* function  $h$ 
  - $h$  approximates the path cost from  $n.state$  to its nearest goal  $G$



$h(n) \geq 0$  requirement

# Ideas on Deriving Heuristic Functions

- Consider a Maze Puzzle problem
  - Layout known
  - Moves  $\leftarrow, \uparrow, \rightarrow, \downarrow$
  - Find path from **S** to **G**
- Example: Euclidean distance
  - $h(n)$  = Euclidean distance from  $n$  to **G**
- General requirements
  - Efficient – e.g., Euclidean distance is  $O(m)$ , where  $m$  = no. dimensions
  - More properties discussed later



$h(G) = 0$  requirement

General idea

Use domain knowledge about the costs (e.g., distances) between a given node and its closest goal – i.e., think about how to define the function  $h$

More on how to define  $h$  in the next lecture

# Implementation with Evaluation Functions

---

- Keep using a priority queue for frontier
  - Use different priorities
- Define an evaluation function  $f$ 
  - Priority for priority queue
  - Priority for node  $n = f(n)$
  - UCS:  $\text{priority} = f(n) = g(n)$
- Now consider different evaluation functions
  - *Greedy Best-First Search*:  $\text{priority} = f(n) = h(n)$
  - *A\* Search*:  $\text{priority} = f(n) = g(n) + h(n)$

# Best-First Search Algorithm

- General graph-search implementation

**function** BEST-FIRST-SEARCH(*problem*, *f*) **returns** a solution node or *failure*

*node*  $\leftarrow$  NODE(STATE=*problem*.INITIAL)

*frontier*  $\leftarrow$  a priority queue ordered by *f*, with *node* as an element

*reached*  $\leftarrow$  a lookup table, with one entry with key *problem*.INITIAL and value *node*

**while not** IS-EMPTY(*frontier*) **do**

*node*  $\leftarrow$  POP(*frontier*)

**if** *problem*.IS-GOAL(*node*.STATE) **then return** *node*

Late Goal Test

**for each** *child* **in** EXPAND(*problem*, *node*) **do**

*s*  $\leftarrow$  *child*.STATE

**if** *s* is not in *reached* **or** *child*.PATH-COST < *reached*[*s*].PATH-COST **then**

*reached*[*s*]  $\leftarrow$  *child*

add *child* to *frontier*

Graph-search (v2)

**return** *failure*

**function** EXPAND(*problem*, *node*) **yields** nodes

*s*  $\leftarrow$  *node*.STATE

**for each** *action* **in** *problem*.ACTIONS(*s*) **do**

*s'*  $\leftarrow$  *problem*.RESULT(*s*, *action*)

*cost*  $\leftarrow$  *node*.PATH-COST + *problem*.ACTION-COST(*s*, *action*, *s'*)

**yield** NODE(STATE=*s'*, PARENT=*node*, ACTION=*action*, PATH-COST=*cost*)


Utilises search problem definitions

# Greedy Best-First Search

---

# The Greedy Best-First Search Algorithm


- Implemented like UCS except
  - $f(n) = h(n)$
- General idea
  - Given all nodes along the frontier
  - Explore next reachable state that you estimate is closest to a goal



- keep picking states estimated to be closest to the goal (based on candidate paths on the frontier)

# The Greedy Best-First Search Algorithm

- Example (tree-search)



Notice that even with the perfect heuristic, we may not get the optimal solution. Why?

Algorithm never exploits information on path already travelled.

Assume this  $h$ :

$n$	$h(n)$	$h^*(n)$
$A$	3	3
$B$	2	2
$C$	1	1
$D$	1	1
$G$	0	0

Trace:

ITR1 =  $[A((-), 3)]$

ITR2 =  $[C((A), 1), B((A), 2)]$

ITR3 =  $[G((A, C), 0), B((A), 2)]$

ITR4 = DONE (A, C, G)

$h^*(n)$  = true path cost from  $n$  to nearest goal



# Completeness and Optimality

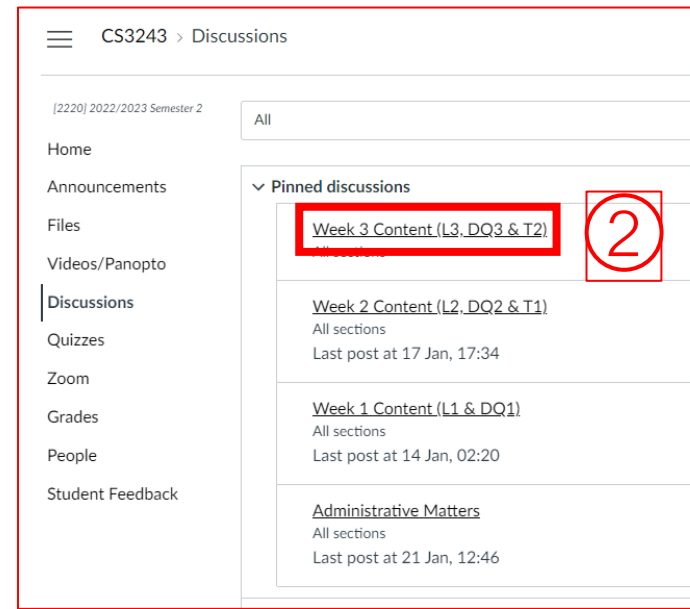
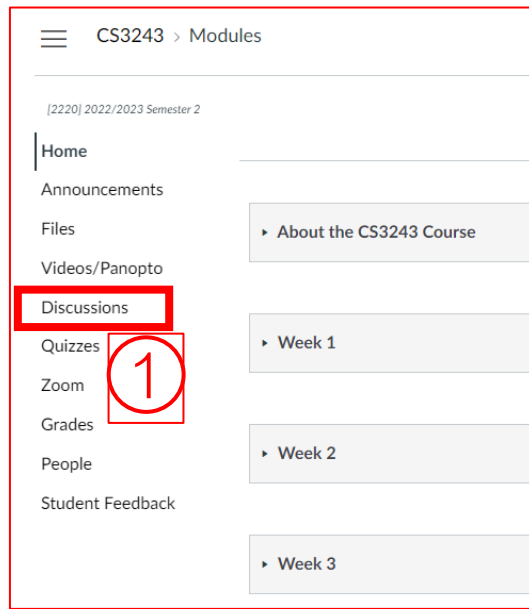
---

- Tree-search implementation is incomplete!
  - General idea
    - Can get stuck in a loop between nodes where  $h$  values are lowest
  - Prove with completeness counter example - T02 Q1a
- Graph-search implementation is complete if search space is finite
  - General idea
    - With no revisits, in finite state space, will visit entire space
  - Prove – T02 Q1b
- Not optimal under either tree search or graph search
  - As shown in example on previous slide
  - Find another example – T02 Q1c

Ordered traversal of path in order of path costs required to ensure optimality

# Questions about the Lecture?

- Do you need to clarify anything?
- Ask on Canvas > CS3243 > Discussion > Week 3 Content




# A\* Search

---

# The A\* Search Algorithm


- Greedy Best-First Search
  - With greedy,  $f(n) = h(n)$
  - Does not consider cost of path already taken
- Accounting for costs already incurred: A\*
  - With A\*,  $f(n) = g(n) + h(n)$ 
    - $g(n)$ : actual path cost from  $S$  to  $n$
    - $h(n)$ : estimated cheapest path cost from  $n$  to  $G$
- A\* priorities
  - Total path cost estimates from  $S$  to  $G$
  - Gets more accurate as paths get explored



# The A\* Search Algorithm

- Example (tree search)

same example as used on greedy (Slide 24)



A\* outputs the optimal solution,  
unlike the Greedy Best-First Search

Will it always be optimal?  
What about graph search?

Assume this  $h$ :

again, same as before (Slide 24)

$n$	$h(n)$	$h^*(n)$
$A$	3	3
$B$	2	2
$C$	1	1
$D$	1	1
$G$	0	0

Trace:

ITR1 = [A((-),0+3)]

ITR2 = [B((A),1+2), C((A),8+1)]

ITR3 = [D((A,B),2+1), C((A),8+1)]

ITR4 = [G((A,B,D),3+0), C((A),8+1)]

ITR5 = DONE (A,B,D,G)

$h^*(n)$  = true path cost from  $n$  to nearest goal

# Completeness and Optimality

---

- Completeness
  - Same criteria as UCS
    - $b$  finite, and  $m$  finite OR has a solution
    - All action costs  $> \epsilon > 0$
- Optimality
  - Depends on the properties of  $h$ 

Ordered traversal of path in order of path costs required to ensure optimality


# Admissible Heuristics

- $h(n)$  is *admissible* if  $\forall n, h(n) \leq h^*(n)$ 
  - $h(n)$  never overestimates the cost
    - Implications
      - Paths not ending at a goal are never over-estimated
        - Evaluation function for non-goal is never over-estimated
        - At non-goal  $n$ ,  $f(n) = g(n) + h(n) \leq g(n) + h^*(n)$
      - Paths ending at a goal are exact
        - Evaluation function of value of a goal is exact
        - At goal  $m$ ,  $f(m) = g(m) + h(m)$ , where  $h(m) = 0$
- Examples:
  - Euclidean distance in the maze environment (always underestimates)
- Theorem: If  $h(n)$  is admissible, then  $A^*$  using tree search is optimal

Main idea, by the time we visit a path to a goal,  $P$ , all paths with actual costs less than  $P$  must be searched

# Optimality of A\* Using Tree Search


- Proving the Theorem
  - T02 Q2a
- Consider the following
  - A\* is optimal  $\rightarrow$  returns optimal path
    - Let  $s$  to  $n$  to  $t^*$  be the optimal path
  - If not optimal:
    - Must explore path to  $t$  first
    - Where  $t$  is a goal
  - Not optimal  $\rightarrow$  explore  $t$  before  $n$





# Optimality of Admissible A\* Under Tree Search

- Assume non-optimal  $t$  expanded before  $n$  (on optimal path)
  - $f(t) < f(n)$
- Assuming tree-search
  - All paths searched
  - All sub-paths\* along a single path to a goal must be searched before path including that goal
    - For non-goal  $m$ ,  
 $f(m) \leq g(m) + h^*(m)$  (since admissible)
    - If goal  $m^*$  on path from  $m$ ,  
 $f(m) \leq f(m^*)$  (given  $\epsilon$ )
  - Since  $t^*$  is goal on optimal path
    - $f(n) < f(t^*) < f(t)$
    - **CONTRADICTION**



\* Consider a path,  $P$ , from an initial state  $s$  to a goal state  $t$ , to be  $s > n_1 > n_2 > \dots > n_k > t$   
Let a sub-path to  $P$ ,  $P'$  be any path  $s > n_1 > n_2 > \dots > n_i$ , where  $1 \leq i \leq k$

# A\* Using Graph Search

---

- Difference between tree search and graph search
  - Under admissibility and tree search
    - All nodes leading to a goal are expanded before the node representing the goal
    - Optimal path will be found
  - Under graph search we may skip some paths (due to no revisiting)
- Under graph search (Version 2), non-redundant paths never skipped
  - Graph search checks and allows revisits if path is non-redundant
    - As long as a path is cheaper, allow it onto the frontier even if must revisit
  - Still optimal since equivalent to tree search
- Under graph search (Version 1) may skip optimal path


# Ensuring A\* Optimal Under Graph Search (Version 1)?

- Will A\* be optimal under graph search Version 1?
- Example

T02 Q3b - construct an alternative example

Assume this admissible h:

$n$	$h(n)$	$h^*(n)$
S	8	8
A	7	7
B	0	6
G	0	0



Trace:

ITR1 = [S((-),0+8)]

ITR2 = [B((S),3+0), A((S),1+7)]

ITR3 = [A((S),1+7), G((S,B),9+0)]

ITR4 = [G((S,B),9+0)] as B popped before, do not revisit

ITR5 = DONE (S,B,G) not the optimal path!

We need a tighter constraint on h


Similar to what UCS offers →  
contours of search progression  
- i.e., stricter ordering of paths

# Why Not Optimal?

- Consider the previous example

Assume this admissible  $h$ :

$n$	$h(n)$	$h^*(n)$
$S$	8	8
$A$	7	7
$B$	0	6
$G$	0	0



Observe the sequence of  $f(n) = g(n) + h(n)$  values along each path:

path to $n$ (from $S$ )	$g(n) + h(n)$	$g(n) + h^*(n)$
$S$	$0+8$	$0+8$
$S > A$	$1+7$	$1+7$
$S > A > B$	$2+0$	$2+6$
$S > A > B > G$	$8+0$	$8+0$

Dip!

path to $n$ (from $S$ )	$g(n) + h(n)$	$g(n) + h^*(n)$
$S$	$0+8$	$0+9$
$S > B$	$3+0$	$3+6$
$S > B > G$	$9+0$	$6$

Dip!

Want  $h$  that ensures paths traversed in order of true path cost!

# Consistent Heuristics

- Forming contours

Under tree search  $g$  costs are monotonically increasing

$f$  costs are not since we can underestimate  $h$  by differing amounts

But ALL paths with path costs less than the optimal path cost will be explored first

- For  $f$  costs to be monotonically increasing along a path

- Assume  $n$  is a predecessor of  $n'$  along a path
- We need:  $g(n) + h(n) \leq g(n) + \text{cost}(n, a, n') + h(n')$
- And thus  $h(n) \leq \text{cost}(n, a, n') + h(n')$

- We will use the above requirement

- $h(n)$  is **consistent** if  $\forall n$ , and successor of  $n, n'$ ,

$$h(n) \leq \text{cost}(n, a, n') + h(n')$$


- Theorem: If  $h(n)$  is consistent, then  $A^*$  using graph search is optimal

Since true path to  $n$  must be less than true path cost to  $n'$ , we want to ensure  $h$  maintains this ordering

Note that:  
consistency  $\Rightarrow$  admissibility  
• Proof – T02 Q3a

# Still an Issue!

- Refer to the UCS example on Slide 14
- Assume consistent heuristic  $h^*$ 
  - With this example, we have
    - $F = \{A(g(A)=0 + h(A)=2)(-)\}; R = \{A\}$ 
      - pop  $A(0+2)(-)$ , push  $C(1+1)(A)$  and  $G(3+0)(A)$
    - $F = \{C(2)(A), G(3)(A)\}; R = \{A, C, G\}$ 
      - pop  $C(2)(A)$ , cannot push  $G(2+0)(A,C)$  since already in  $R$ !
    - $F = \{G(3)(A)\}; R = \{A, C, G\}$ 
      - pop  $G(3)(A)$ , path is  $A \rightarrow G$



Considering  $G$  as reached too early under graph search for  $A^*$  is similar to applying Early Goal Test to UCS!

# Graph Search Algorithm (Version 3)

---

```
frontier = {Node(initial state), NULL}
reached = {}
while frontier not empty:
    current = frontier.pop()
    reached.insert(current.state: current)
    if isGoal(current.state): return current.getPath()
    for a in actions(current.state):
        successor = Node(T(current.state, a), current)
        if successor.state not in reached:
            frontier.push(successor)
return failure
```

- Only adds a node to reached when it is popped
- Proof requires this version (given counterexample for Version 1 on Slide 38)
- Prove this in a similar manner to the UCS proof (contours) – T02 Q2b

# Dominant Heuristics

---



# Efficiency & Dominance

---

- Efficiency of  $A^*$  depends on the accuracy of its heuristics
  - Higher heuristic accuracy means we need to try fewer paths
  - Specifics not covered in CS3243
- Which heuristics are better?
- If  $h_1(n) \geq h_2(n)$  for all  $n$ , then  $h_1$  ***dominates***  $h_2$ 
  - If  $h_1$  is also *admissible*
    - $h_1$  must be closer to  $h^*$  than  $h_2$
    - $h_1$  must be more efficient than  $h_2$

Note: with typical interpretations, dominance requires admissibility. We apply this in CS3243.

# Questions about the Lecture?

- Do you need to clarify anything?
- Ask on Canvas > CS3243 > Discussion > Week 3 Content

