


National University of Singapore
School of Computing
CS3243 Introduction to AI

Tutorial 2: Informed Search (Solutions)

1. (a) Provide a counter-example to show that the **tree search** implementation for the **Greedy Best-First Search** algorithm is **incomplete**.

Solution: Consider the following search space with initial state s_0 , goal state g , and where $h(s_0) = 3$, $h(s_1) = 4$, $h(s_2) = 5$, and $h(g) = 0$.




Each time s_0 is explored, we add s_1 to the front of the frontier, and each time s_1 is explored, we add s_0 to the front of the frontier. Notice that s_2 is never at the front of the frontier. This causes the greedy best-first search algorithm to continuously loop over s_0 and s_1 .

- (b) Briefly explain why the **graph search** implementation for the **Greedy Best-First Search** algorithm is **complete**.

Solution: Assuming a finite search space, a graph search implementation of the greedy best-first search algorithm will eventually visit all states within the search space. As such, the algorithm would either find a goal state and return the path to it, or else, indicate that there is no solution if a goal is not found.

- (c) Provide a counter-example to show that neither the **tree search** nor the **graph search** implementations for the **Greedy Best-First Search** algorithm are **optimal**.


Solution: Consider the following search space with initial state s_0 , goal states g_1 and g_2 , and where $h(s_0) = 9$, $h(g_1) = 0$, $h(s_1) = 1$, and $h(g_2) = 0$.



With either implementation, when s_0 is explored, g_1 would be added to the front of the frontier and then explored next, resulting in the algorithm returning the non-optimal $s_0 \rightarrow g_1$ path.

2. (a) Prove that the **tree search** implementation of the **A* Search** algorithm is optimal when an **admissible heuristic** is utilised.

Solution: Assume the following search tree below.



Let s be the initial state, n be an intermediate state along the optimal path, t be a sub-optimal goal state (i.e., a goal state reached via a suboptimal path), and t^* be the goal along the optimal path.

An optimal solution implies that n must be expanded before t .

Proof by contradiction:

- Let us assume that a suboptimal solution is found – i.e., that t is expanded before n , which implies that (A): $f(t) \leq f(n)$
- In other words, given the above frontier, only when $f(t) < f(n)$ would we expand t before n
- However, since t is not on the optimal path but t^* is, we have:

$$\begin{array}{ll}
 f(t) > f(t^*) & \\
 f(t) > g(t^*) & \text{since } h(t^*) = 0 \\
 f(t) > g(n) + p(n, t^*) & \text{where } p(n, t^*) \text{ is the actual cost from } n \text{ to } t^* \\
 f(t) > g(n) + h(n) & \text{asserting admissibility*} \\
 f(t) > f(n) & \text{this contradicts (A)}
 \end{array}$$

- Note: we do not consider $f(t) = f(n)$ since that would mean $f(t)$ is equally optimal.

- (b) Prove that the **graph search** implementation of the **A* Search** algorithm is optimal when a **consistent heuristic** is utilised. Assume graph search **Version 3**.

Solution: Similar to the UCS proof of optimality under graph search, we must show that when a node n is popped from the frontier, we have found the optimal path to it.

Let $f(s_k) = g(s_k) + h(s_k)$ be the minimum f value for s_k we have observed when s_k is popped.

Let the optimal path from the start node, s_0 , to any node, s_g , be $P = s_0, s_1, \dots, s_{g-1}, s_g$. We must show that when we pop s_g , $f(s_g) = g(s_g) + h(s_g) = g^*(s_g) + h(s_g)$, where $g^*(s_g)$ denotes the optimal path cost from s_0 to s_g via P .

Base case: $f(s_0) = g(s_0) + h(s_0) = g^*(s_0) + h(s_0) = h(s_0)$ as s_0 is the start node.

Induction step: Assume that for all s_0, s_1, \dots, s_k , when we pop s_i , $f(s_i) = g(s_i) + h(s_i) = g^*(s_i) + h(s_i)$, or rather, $g(s_i) = g^*(s_i)$.

Since $g^*(s_{k+1})$ is the minimum path cost from s_0 to s_{k+1} , we know that:

$$\begin{aligned} g(s_{k+1}) + h(s_{k+1}) &\geq g^*(s_{k+1}) + h(s_{k+1}) \\ g(s_{k+1}) &\geq g^*(s_{k+1}) \end{aligned} \quad (\text{A})$$

To make sure that each s_{k+1} is only popped after we pop s_k , the condition $f(s_k) \leq f(s_{k+1})$, or rather $h(s_k) \leq c(s_k, s_{k+1}) + h(s_{k+1})$, where $c(s_k, s_{k+1})$ is the action cost from s_k to s_{k+1} , is required, which leads us to assert that h is consistent.

Consequently, just after s_k is popped, we have:

$$\begin{aligned} g(s_{k+1}) + h(s_{k+1}) &= \min\{g(s_{k+1}) + h(s_{k+1}), g(s_k) + c(s_k, s_{k+1}) + h(s_{k+1})\} \\ g(s_{k+1}) &= \min\{g(s_{k+1}), g(s_k) + c(s_k, s_{k+1})\} \\ &\leq g(s_k) + c(s_k, s_{k+1}) \\ &= g^*(s_k) + c(s_k, s_{k+1}) && \text{from our inductive hyp.} \\ &= g^*(s_{k+1}) && (\text{B}) \end{aligned}$$

From (A) and (B), we obtain $g(s_{k+1}) = g^*(s_{k+1})$. Hence, by induction, whenever we pop a node from the frontier, the optimal path to the node would have been found.

Also, given that graph search Version 3 is utilised (i.e., only nodes popped from the frontier are added to reached), the optimal path would not be excluded.

3. (a) Given that a **heuristic** h is such that $h(t) = 0$, where t is any goal state, prove that if h is **consistent**, then it must be **admissible**.

Solution: The proof is by induction on $k(n)$, which denotes the number of actions required to reach the goal from a node n to the goal node t .

Base case ($k = 1$, i.e., the node n is one step from t):

Since the heuristic function h is consistent,


$$h(n) \leq c(n, a, t) + h(t)$$

Since $h(t) = 0$,

$$h(n) \leq c(n, a, t) = h^*(n)$$

Therefore, h is admissible.

Induction case:



Suppose that our assumption holds for every node that is $k - 1$ actions away from t , and let us observe a node n that is k actions away from t ; that is, the least-actions optimal path from n to t has $k > 1$ steps.

We write the optimal path from n to t as

$$n \rightarrow n_1 \rightarrow n_2 \rightarrow \cdots \rightarrow n_{k-1} \rightarrow t.$$

Since h is consistent, we have

$$h(n) \leq c(n, a, n_1) + h(n_1).$$

Now, note that since n_1 is on a least-cost path to t from n , we must have that the path $n_1 \rightarrow n_2 \rightarrow \cdots \rightarrow n_{k-1} \rightarrow t$ is a minimal-cost path from n_1 to t as well. By our induction hypothesis we have

$$h(n_1) \leq h^*(n_1)$$

Combining the two inequalities we have that


$$h(n) \leq c(n, a, n_1) + h^*(n_1)$$

Note that $h^*(n_1)$ is the cost of the optimal path from n_1 to t ; by our previous observation (that $n_1 \rightarrow n_2 \rightarrow \cdots \rightarrow n_{k-1} \rightarrow t$ is an optimal cost path from n_1 to t), we have that the cost of the optimal path from n to t — i.e. $h^*(n)$ — is exactly $c(n, a, n_1) + h^*(n_1)$, which concludes the proof.

- (b) Give an example of an **admissible heuristic** that is **not consistent**.

Solution: An example of an admissible heuristic function that is not consistent is as follows.

Consider a heuristic function h , such that $h(s_0) = 3$, $h(s_1) = 1$, and $h(t) = 0$ for the following graph.



h is admissible since

$$h(s_0) \leq h^*(s_0) = 1 + 2 = 3$$

$$h(s_1) \leq h^*(s_1) = 2$$

However, h is not consistent since $3 = h(s_0) > c(s_0, s_1) + h(s_1) = 1 + 1 = 2$.

4. We have seen various search strategies in class, and analysed their worst-case running time. Prove that **any deterministic search algorithm** will, in the worst case, **search the entire state space**. More formally, prove the following theorem

Theorem 1. *Let \mathcal{A} be some complete, deterministic search algorithm. Then for any search problem defined by a finite connected graph $G = \langle V, E \rangle$ (where V is the set of possible states and E are the transition edges between them), there exists a choice of start node s_0 and goal node g so that \mathcal{A} searches through the entire graph G .*


Solution: Let us begin by running \mathcal{A} on the graph G , without setting any goal node at all: that is, there are no goal nodes at all in G . In this case, the algorithm \mathcal{A} will return “False” when it explores the entire set V . Let $H_t(\mathcal{A}, s_0) \subseteq V$ be the set of nodes that \mathcal{A} explores if it starts at s_0 , and does not encounter a goal node at steps $1, \dots, t$ (at $t = 1$ we have $H_1(\mathcal{A}, s_0) = \{s_0\}$). We also let v_t be the node that \mathcal{A} selects at time t given that it has observed the set $H_{t-1}(\mathcal{A}, s_0)$ so far. We note that it is entirely possible that \mathcal{A} selects $v_t \in H_{t-1}(\mathcal{A}, s_0)$; however, we make a simple observation: the sequence $(H_t(\mathcal{A}, s_0))_{t=1}^{\infty}$ is weakly increasing in size, and there exists some time t^* such that for all $t > t^*$, $H_t(\mathcal{A}, s_0) =$

V ; in other words, since \mathcal{A} is a complete search algorithm, it will continue exploring the nodes in G until all nodes have been exhausted. Let us assume that t^* is the first time step for which $H_t(\mathcal{A}, s_0) = V$. In other words, at time $t^* - 1$, $|H_{t^*-1}(\mathcal{A}, s_0)| = |V| - 1$.

We now set the goal node to be v_{t^*} . From our previous argument, we know that when \mathcal{A} starts at s_0 it will explore a set of size $|V| - 1$ before reaching v_{t^*} , realizing that it is a goal node and terminating. In other words, for any node s_0 , if we select a goal node according to the above procedure, the algorithm \mathcal{A} will exhaustively search through the entire graph before reaching a goal node.

Here is another, inductive proof. let us set the goal node to some arbitrary node g_1 . If \mathcal{A} searches through the entire graph G when g_1 is the goal, we are done; otherwise, let U_1 be the set of unsearched nodes when g_1 is the goal node. We take an arbitrary node g_2 in U_1 to be the goal; since \mathcal{A} is deterministic and complete it will run the same search order that it did when g_1 was the goal, and then search through the nodes in U_1 until it reaches g_2 . If it searched through all the nodes in U_1 as well, we are done, otherwise repeat. In general, suppose that we have set g_t to be the goal node and that \mathcal{A} did not search through the entire graph until it reached g_t ; let U_t be the set of unsearched nodes when g_t is the goal node. We set g_{t+1} to be some arbitrary node in U_t and rerun \mathcal{A} ; since \mathcal{A} is deterministic we know that when g_{t+1} is the goal we have $U_{t+1} \subset U_t$. Since $U_1 \supset U_2 \supset \dots \supset U_t$ and the number of nodes in G is finite, there exists some iteration t^* such that $U_{t^*} = \emptyset$; thus g_{t^*} is a goal node for which \mathcal{A} searches through the entire graph.

5. (a) In the search problem below, we have listed 5 heuristics. Indicate whether each **heuristic** is **admissible** and/or **consistent** in the table below.



	S	A	B	G	Admissible	Consistent
h_1	0	0	0	0		
h_2	8	1	1	0		
h_3	9	3	2	0		
h_4	6	3	1	0		
h_5	8	4	2	0		

Solution:

	S	A	B	G	Admissible	Consistent
h_1	0	0	0	0	True	True
h_2	8	1	1	0	True	False
h_3	9	3	2	0	True	True
h_4	6	3	1	0	True	False
h_5	8	4	2	0	False	False

- (b) Write out the order of the nodes that are explored by the **A* Search** algorithm. Assume a **graph search** implementation that utilises heuristic h_4 . Further, assume graph search **Version 3**.

You should express your answer in the form $A-B-C$ (i.e., no spaces, all uppercase letters, delimited by the dash (–) character), which, for example, corresponds to the order A , B , and C .

Solution: $S-A-B-G$

- (c) Which heuristic would you use? Explain why.

Solution: The heuristic h_3 corresponds to the exact cost from each node to the goal node (i.e., $h_3 = h^*$), and therefore it is the optimal heuristic.

- (d) Prove or disprove the following statement:

The heuristic $h(n) = \max\{h_3(n), h_5(n)\}$ is admissible.

Solution: This is false, since $4 = h(A) > h^*(A) = 3$.