

# Adversarial Search: Playing Games

---

CS3243: Introduction to Artificial Intelligence – Lecture 7

6 March 2023

# Contents

---

1. Administrative Matters
2. Reviewing Search Problems
3. Adversarial Search Problems (i.e., Games)
4. Optimal Decisions via Minimax
5.  $\alpha$ - $\beta$  Pruning
6. Heuristic Minimax

Reference: AIMA 4<sup>th</sup> Edition, Section 6.1-6.3

# Administrative Matters

---

# Upcoming...

---

- Deadlines
  - None this week!
  - TA6 (released today)
    - *Due in your Week 6 tutorial session*
    - *Submit the a physical copy (more instructions on the Tutorial Worksheet)*
  - Prepare for the tutorial!
    - Participation marks = 5%
  - Midterm Marks Appeal
    - *Due next week (i.e., Week 9) on the day of your assigned tutorial session, 2359 hrs*
  - Project 2
    - *Due next week, Sunday (19 March), 2359 hrs*

# Reviewing Search Problems

---

# So Far...

---

- Path search (path planning)
  - Search for a path from start to goal
    - Complete: finds a solution or says when there isn't one
    - Optimal: path cost of path found is minimal
  - Uninformed
    - Systematically search all paths via general search problem formulation
  - Informed
    - Uses a heuristic to search less of the search space
- Goal search
  - Focus on goal and ignore path
    - Completeness consideration only
  - Local search
    - Uses heuristic to guide search to goal (uses restarts; many variants)
  - Constraint satisfaction problem
    - Uses specific search problem formulation and shrinks search space via inference

# Games

---

# Games and Search

---

- Can we solve games using existing methods?
  - In our searching thus far, we control all actions
    - All actions taken are determined by our agent
  - With games, your opponent decides actions too...
    - Multi-agent problem
    - Conventional planning  $\Rightarrow$  wasted computation since opponent can spoil your plans



# Games and Search

---

- What is a game anyway?
  - Assume two players
  - Zero-sum game
    - Winner gets paid, and loser pays
  - We define
    - MAX player – player 1, who wants to maximise value (agent)
    - MIN player – player 2, who want to minimise value (opponent who wants agent to lose)
- General idea behind the search problem
  - Simulate play against utility maximising opponent
  - Find a strategy – i.e., define a move for every possible opponent response


# Search Problem Formulation for Games

---

# Formulating Games

- State representation

- As per general formulation



- TO-MOVE(s)

- Returns  $p$ , the player to move in state  $s$

- ACTIONS(s)

- Legal moves in state  $s$

- RESULT(s, a)

- The transition model; returns resultant state when taking action  $a$  at state  $s$

- IS-TERMINAL(s)

- Returns TRUE when game is over and FALSE otherwise
  - States where game has ended are called terminal states

- UTILITY(s, p)

- Defines the final numeric value to player  $p$  when the game ends in terminal state  $s$

# Note on Utility

- Given zero-sum games
  - At terminal state  $s$ 
    - $UTILITY(MAX, s) + UTILITY(MIN, s) = 0$

- Tic-Tac-Toe example

- X (agent) wins

- $UTILITY(s_i, MAX) = 1$
    - $UTILITY(s_i, MIN) = -1$

$s_i$  :

X		O
O	X	
X	O	X

- O (opponent) wins

- $UTILITY(s_j, MAX) = -1$
    - $UTILITY(s_j, MIN) = 1$

$s_j$  :

X		O
X	O	X
O	X	

- Draw

- $UTILITY(s_k, MAX) = 0$
    - $UTILITY(s_k, MIN) = 0$

$s_k$  :

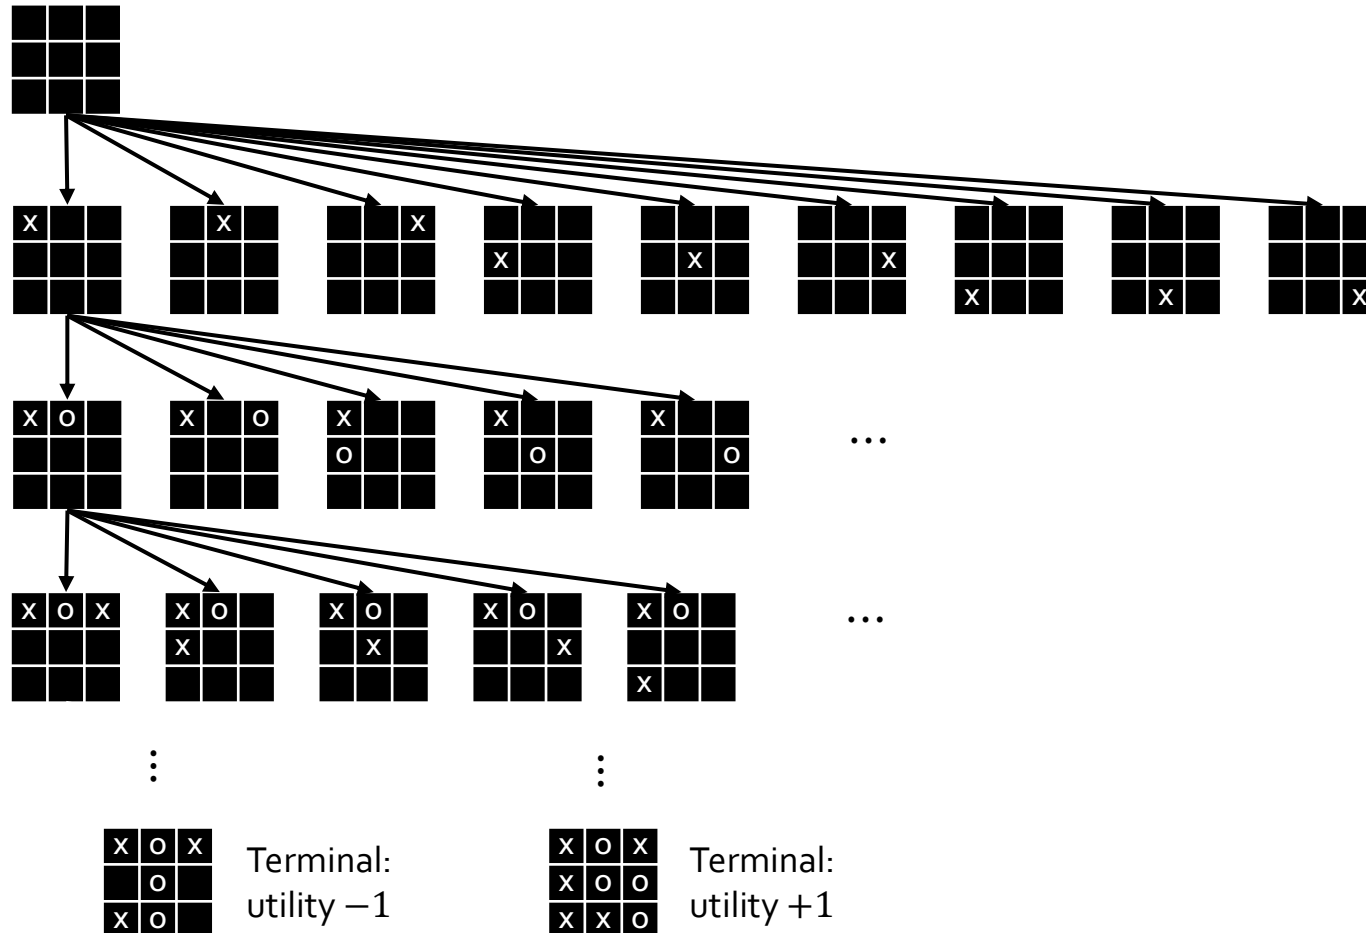
X	X	O
O	O	X
X	O	X

Notice that here the utilities are relative to the player – however, we will standardise the scores such that they reflect only the MAX player (i.e., agent) scores later

# Game Trees

---

# Example Game Tree



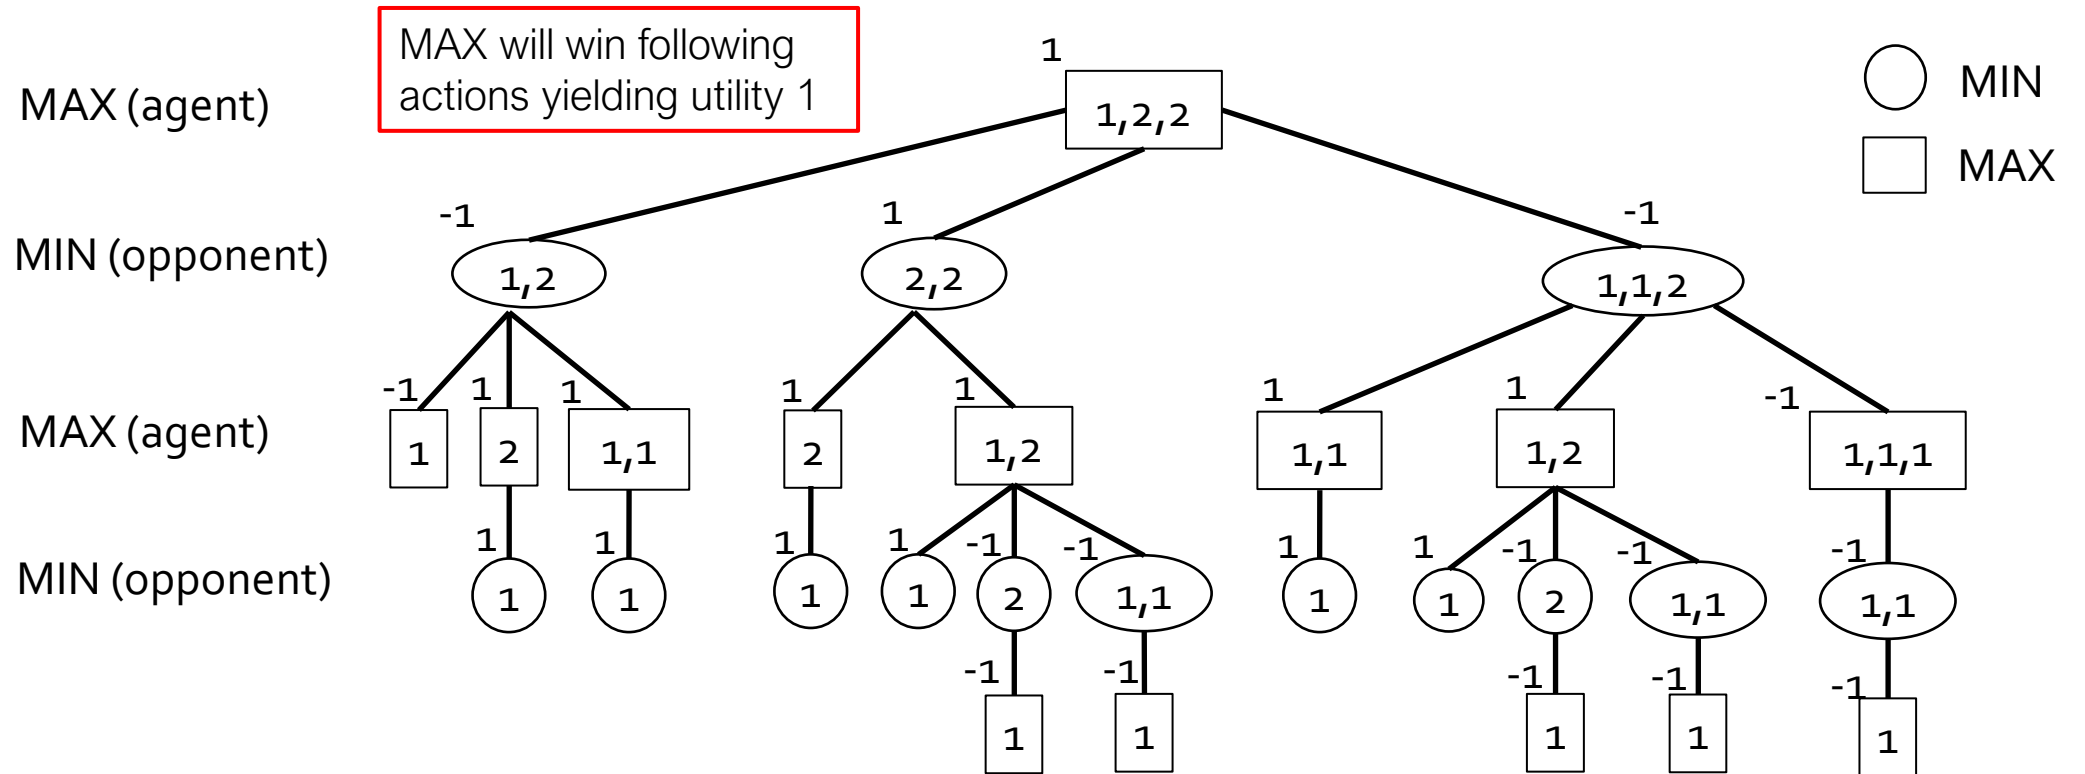
- More on environment characteristics
  - 2-player
  - Deterministic
  - Turn-taking
- Zero-sum implications
  - Loser for every winner
  - Agent utilises sum to zero
  - May also consider constant-sum game
  - Completely adversarial game

# Another Example: Game of NIM

---

- Several piles of sticks are given
  - Represent the configuration of piles by a monotone sequence of integers
    - Example: (1,3,5)
  - With each turn, a player may remove any number of sticks from **ONE** pile
    - Example:
      - Remove 4 sticks from last pile (of 5 sticks)  $\Rightarrow$  (1,3,5) becomes (1,1,3)
  - The player who takes the last stick loses
- Let's try...
  - Represent the NIM game (1,2,2) as a game tree

# Game of NIM: (1,2,2) Game Tree



DFS Traversal with Backwards induction on utility



# Strategies

---

# Player Strategies

---

- A strategy  $s$  for *player  $i$*  :
  - What will *player  $i$*  do at every node of the game tree that they make a move in?
    - Need to specify behaviour in states that may never be reached!

- Winning strategy

A strategy  $s_1^*$  for **Player 1** is called winning if  
for any strategy  $s_2$  by **Player 2**,  
the game ends with **Player 1** as the winner.

- Non-losing strategy

A strategy  $t_1^*$  for **Player 1** is called non-losing if  
for any strategy  $s_2$  by **Player 2**,  
the game ends in either a tie or a win for **Player 1**.


# Optimal Strategy at Node - Minimax

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s, \text{MAX}) & \text{if Is-Terminal}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MIN} \end{cases}$$


- Intuitively
  - MAX chooses move to maximise the minimum payoff
    - MIN chooses at successors
  - MIN chooses move to minimise the maximum payoff
    - MAX chooses at successors

Remember that  $\text{Result}(s, a)$  outputs the utility in terms of MAX player's utility (i.e., MIN wants to minimise this value, while MAX wants to maximise it)

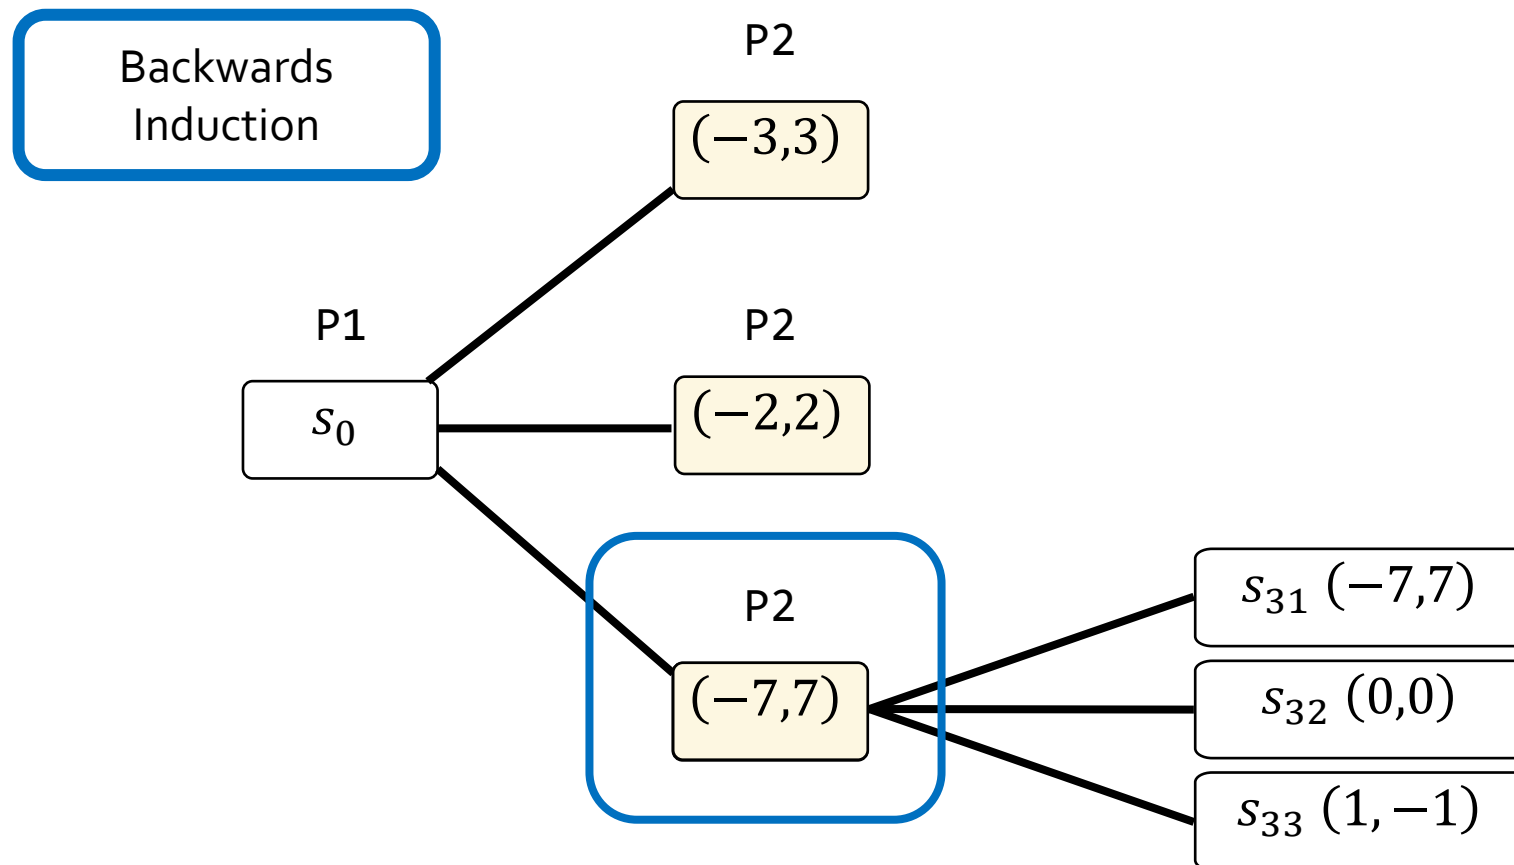
# Minimax Play – Subgame Perfect Nash Equilibrium



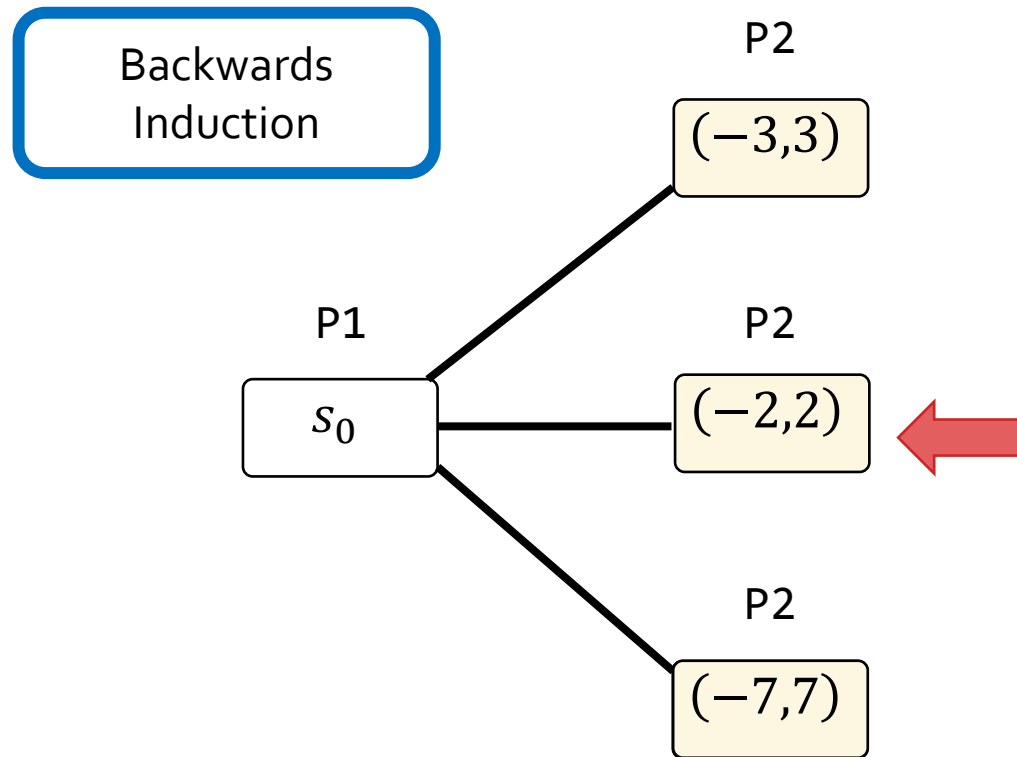
# Minimax Play – Subgame Perfect Nash Equilibrium



# Minimax Play – Subgame Perfect Nash Equilibrium

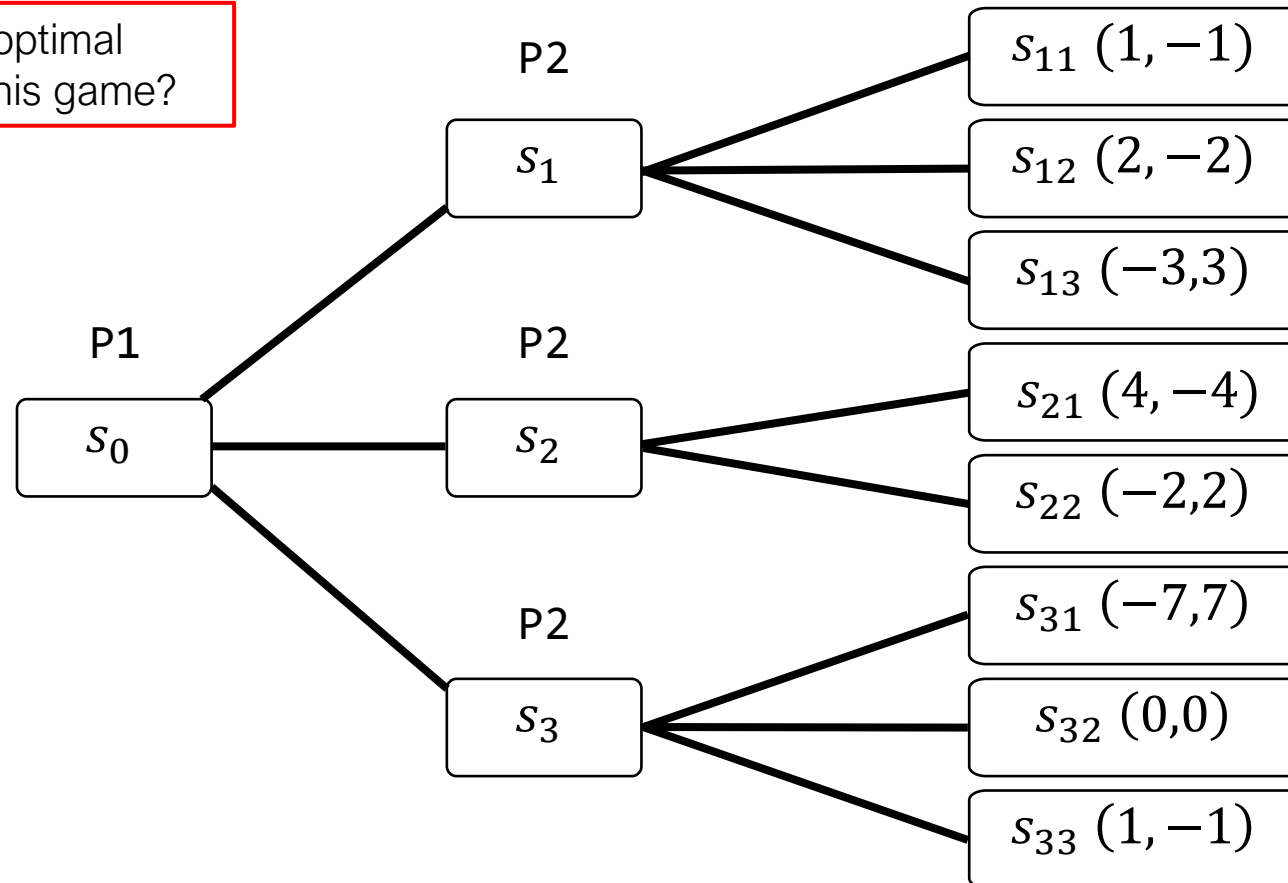


# Minimax Play – Subgame Perfect Nash Equilibrium



# Minimax Play – Subgame Perfect Nash Equilibrium

What are the optimal strategies in this game?





# Minimax Algorithm Properties

---

- Complete?
  - Yes (if game tree is finite)
- Optimal?
  - Yes (optimal gameplay)
- Time
  - $O(b^m)$
- Space
  - $O(bm)$
- Minimax runs in time polynomial in tree size
- Returns a subgame perfect Nash Equilibrium
  - I.e., the best action at every node

Are we done?

# Backwards Induction

---

- Game trees are massive
  - Chess has a massive game tree
    - $10^{123}$  nodes
  - In comparison, planet Earth has about  $10^{50}$  atoms ...
- Impossible to expand the entire tree
- Have to find ways to shrink the search tree
  - We've seen this before
  - Common theme in search

# Questions about the Lecture?

---

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
  - Specify a question
  - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)  
<https://archipelago.rocks/app/resend-invite/64238273017>

# $\alpha$ - $\beta$ Pruning


---

# $\alpha$ - $\beta$ Pruning - General Idea

---


- Basic idea
  - Don't explore moves that would never be considered
- Maintain bounds on values seen thus far while searching
  - $\alpha$  bounds MAX's values
  - $\beta$  bounds MIN's values
- Prune subtrees that will never affect Minimax decision

# $\alpha$ - $\beta$ Pruning – Example Trace




Note that from this point, all utility values will be specified in terms of the MAX player.

# $\alpha$ - $\beta$ Pruning – Example Trace



Note that from this point, all utility values will be specified in terms of the MAX player.


# $\alpha$ - $\beta$ Pruning – Example Trace



Note that from this point, all utility values will be specified in terms of the MAX player.




# $\alpha$ - $\beta$ Pruning – Example Trace




Note that from this point, all utility values will be specified in terms of the MAX player.

# $\alpha$ - $\beta$ Pruning – Example Trace



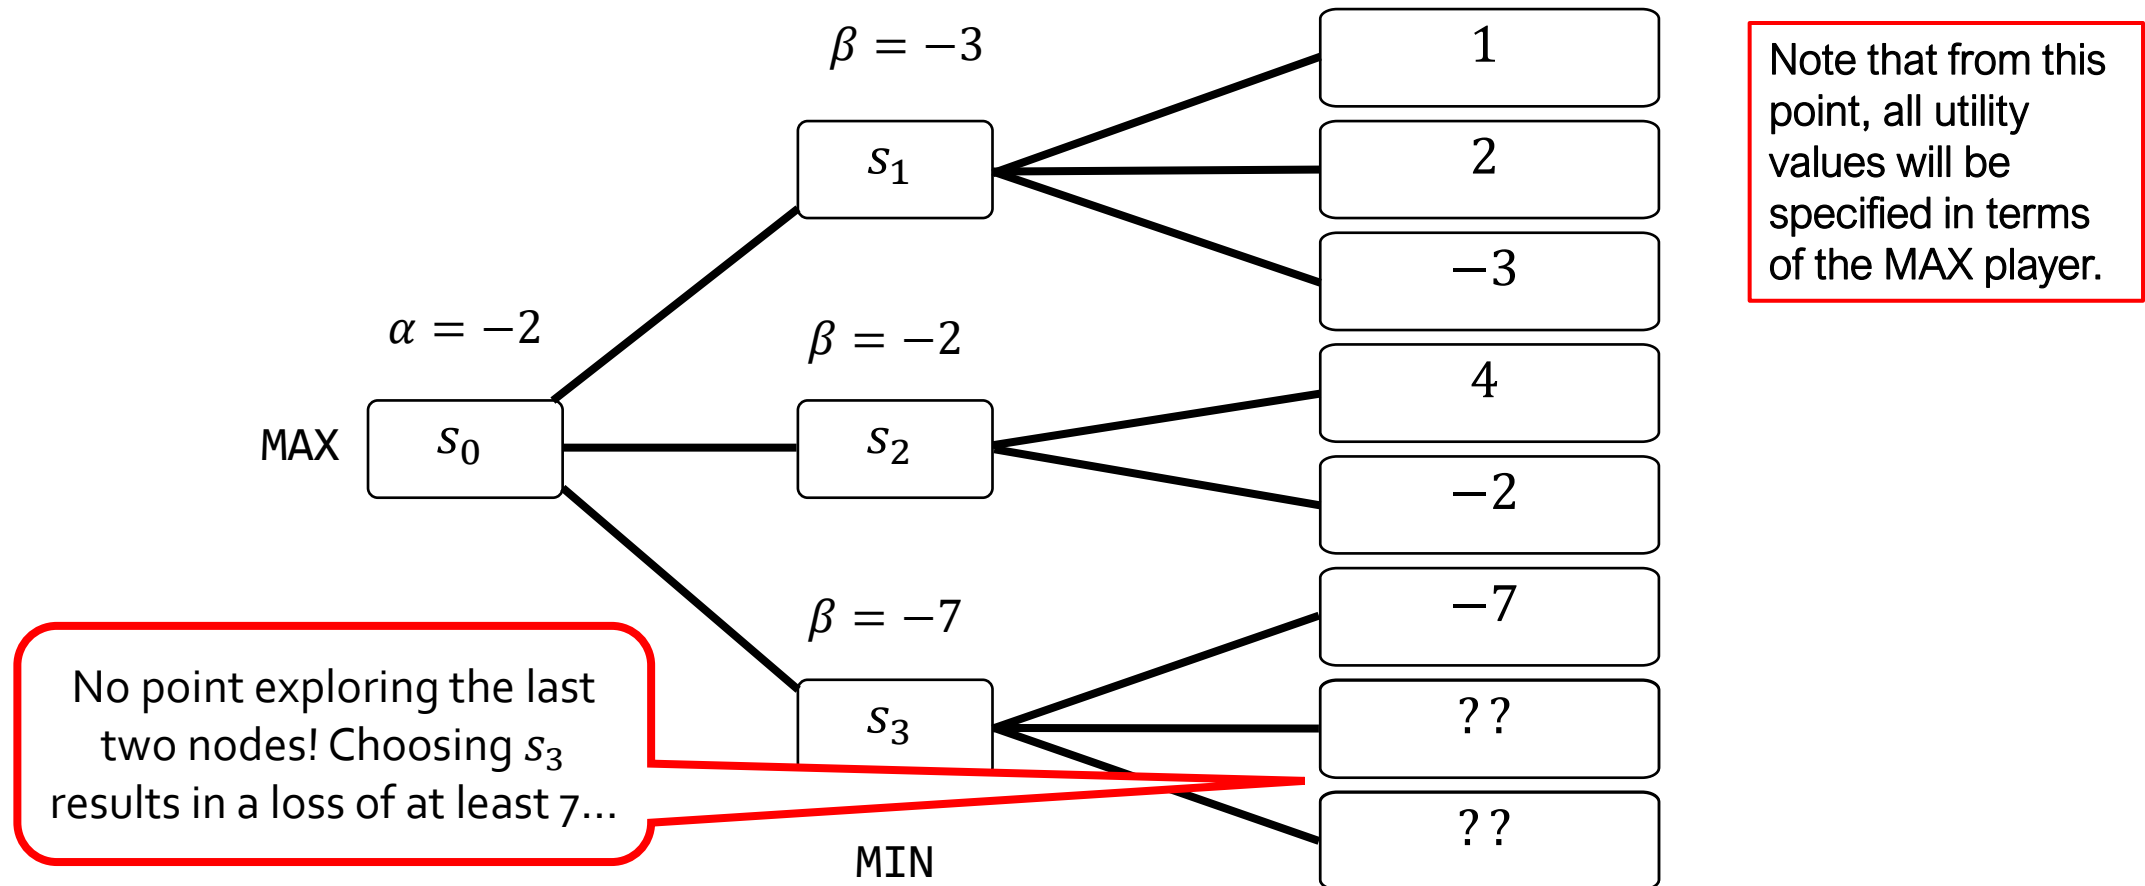
Note that from this point, all utility values will be specified in terms of the MAX player.

# $\alpha$ - $\beta$ Pruning – Example Trace

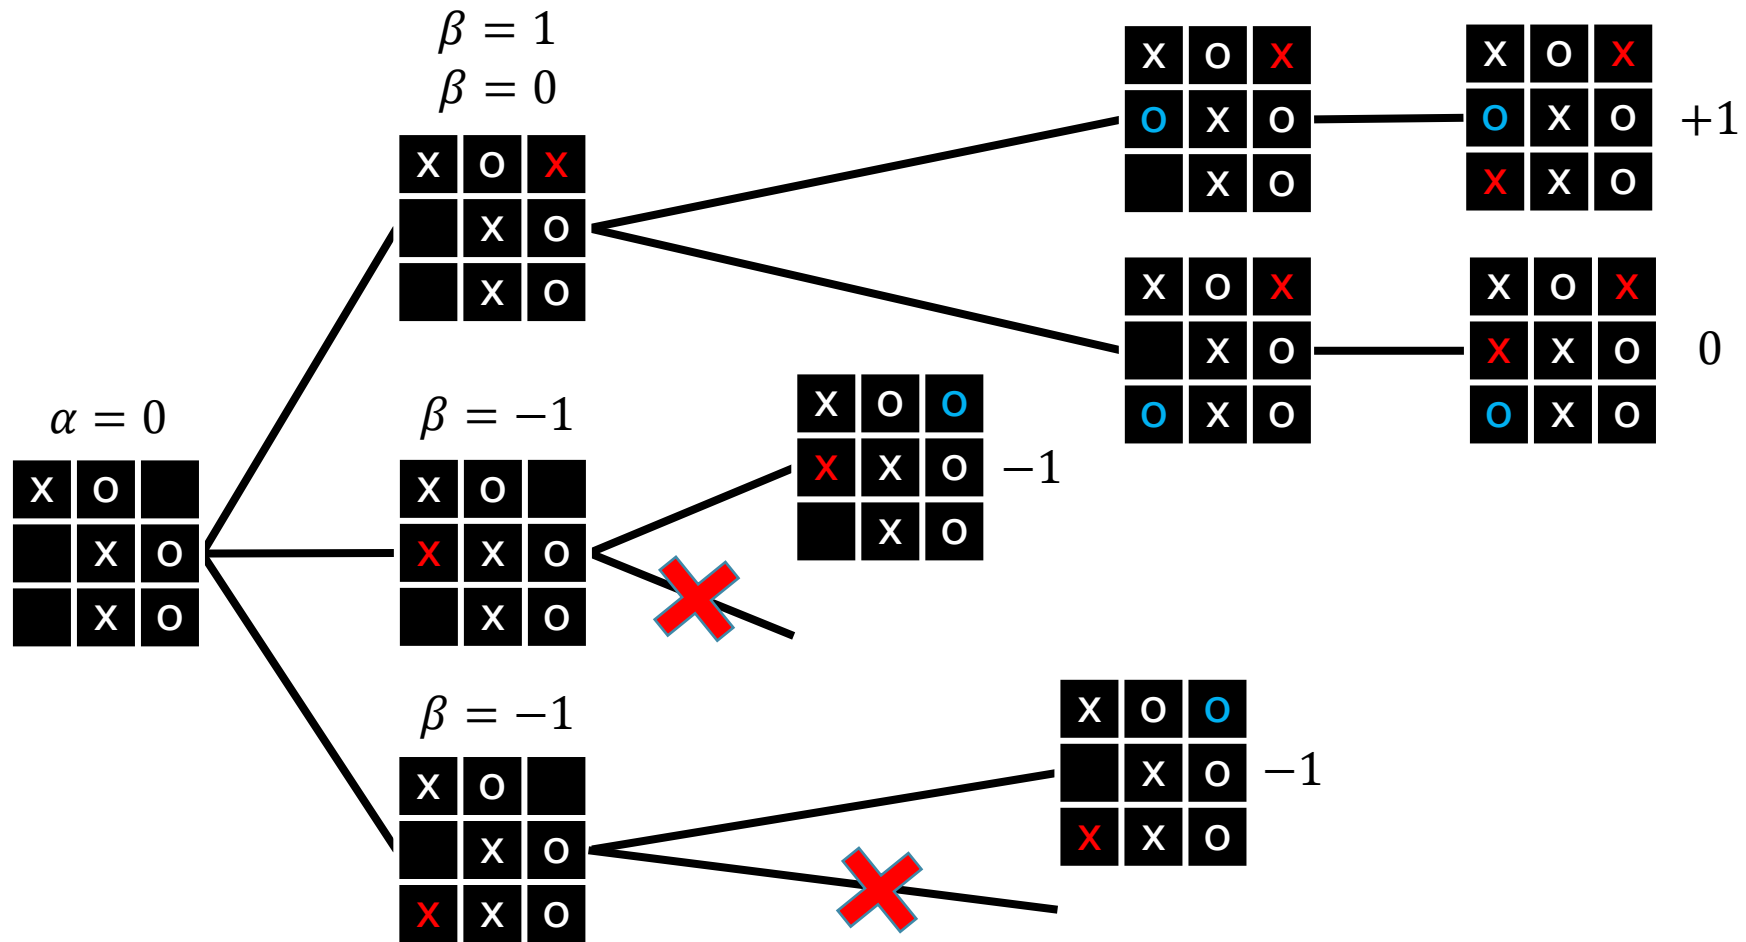


Note that from this point, all utility values will be specified in terms of the MAX player.

# $\alpha$ - $\beta$ Pruning – Example Trace



# $\alpha$ - $\beta$ Pruning – Tic-Tac-Toe Example



# $\alpha$ - $\beta$ Pruning

- MAX node  $n$ 
  - $\alpha(n)$  = highest observed value found on path from  $n$
  - Initially  $\alpha(n) = -\infty$
- MIN node  $n$ 
  - $\beta(n)$  = lowest observed value found on path from  $n$
  - Initially  $\beta(n) = +\infty$
- Pruning rules
  - Given a MIN node  $n$ , stop searching below  $n$  if
    - Some MAX ancestor  $i$  (of  $n$ ) with  $\alpha(i) \geq \beta(n)$
  - Given a MAX node  $n$ , stop searching below  $n$  if
    - Some MIN ancestor  $i$  (of  $n$ ) with  $\beta(i) \leq \alpha(n)$

MIN will choose  $\beta(n)$  or lower at  $n$ , but ancestor MAX will NEVER choose the subtree at  $n$  since at  $i$ , there is a better option with higher value  $\alpha(i)$

MAX will choose  $\alpha(n)$  or higher at  $n$ , but ancestor MIN will NEVER choose the subtree at  $n$  since at  $i$ , there is a better option with lower value  $\beta(i)$

# $\alpha$ - $\beta$ Pruning Analysis

---

- Pruning a branch never affects the final outcome
- Good move ordering improves effectiveness of pruning
  - “Perfect” ordering
    - Time complexity  $O(b^{m/2})$
    - Good pruning strategies allow us to search twice as deep!
  - Example: Chess
    - Simple ordering gets you close to best-case result
      - Checks
      - Take pieces
      - Forward moves
      - Backwards move
- Expansion-order heuristics will improve the search
- Random ordering gives complexity  $O(b^{3m/4})$  for  $b < 1000$

# Issue with $\alpha$ - $\beta$ Pruning

---

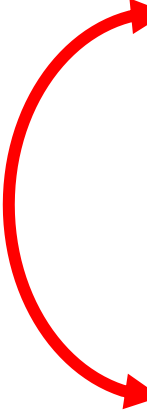
- Original Problem
  - Most games have very large game trees
- Solution
  - $\alpha$ - $\beta$  pruning can remove large parts of search trees
- Unresolved Issue
  - Maximum depth of tree
    - Backwards induction works backwards from terminal states
    - Still have to traverse to a terminal states
  - Standard solution – Heuristic Minimax
    - Cutoff test – e.g., depth limit (DLS)
    - Evaluation function – estimates expected utility of state



# Heuristic Minimax

---

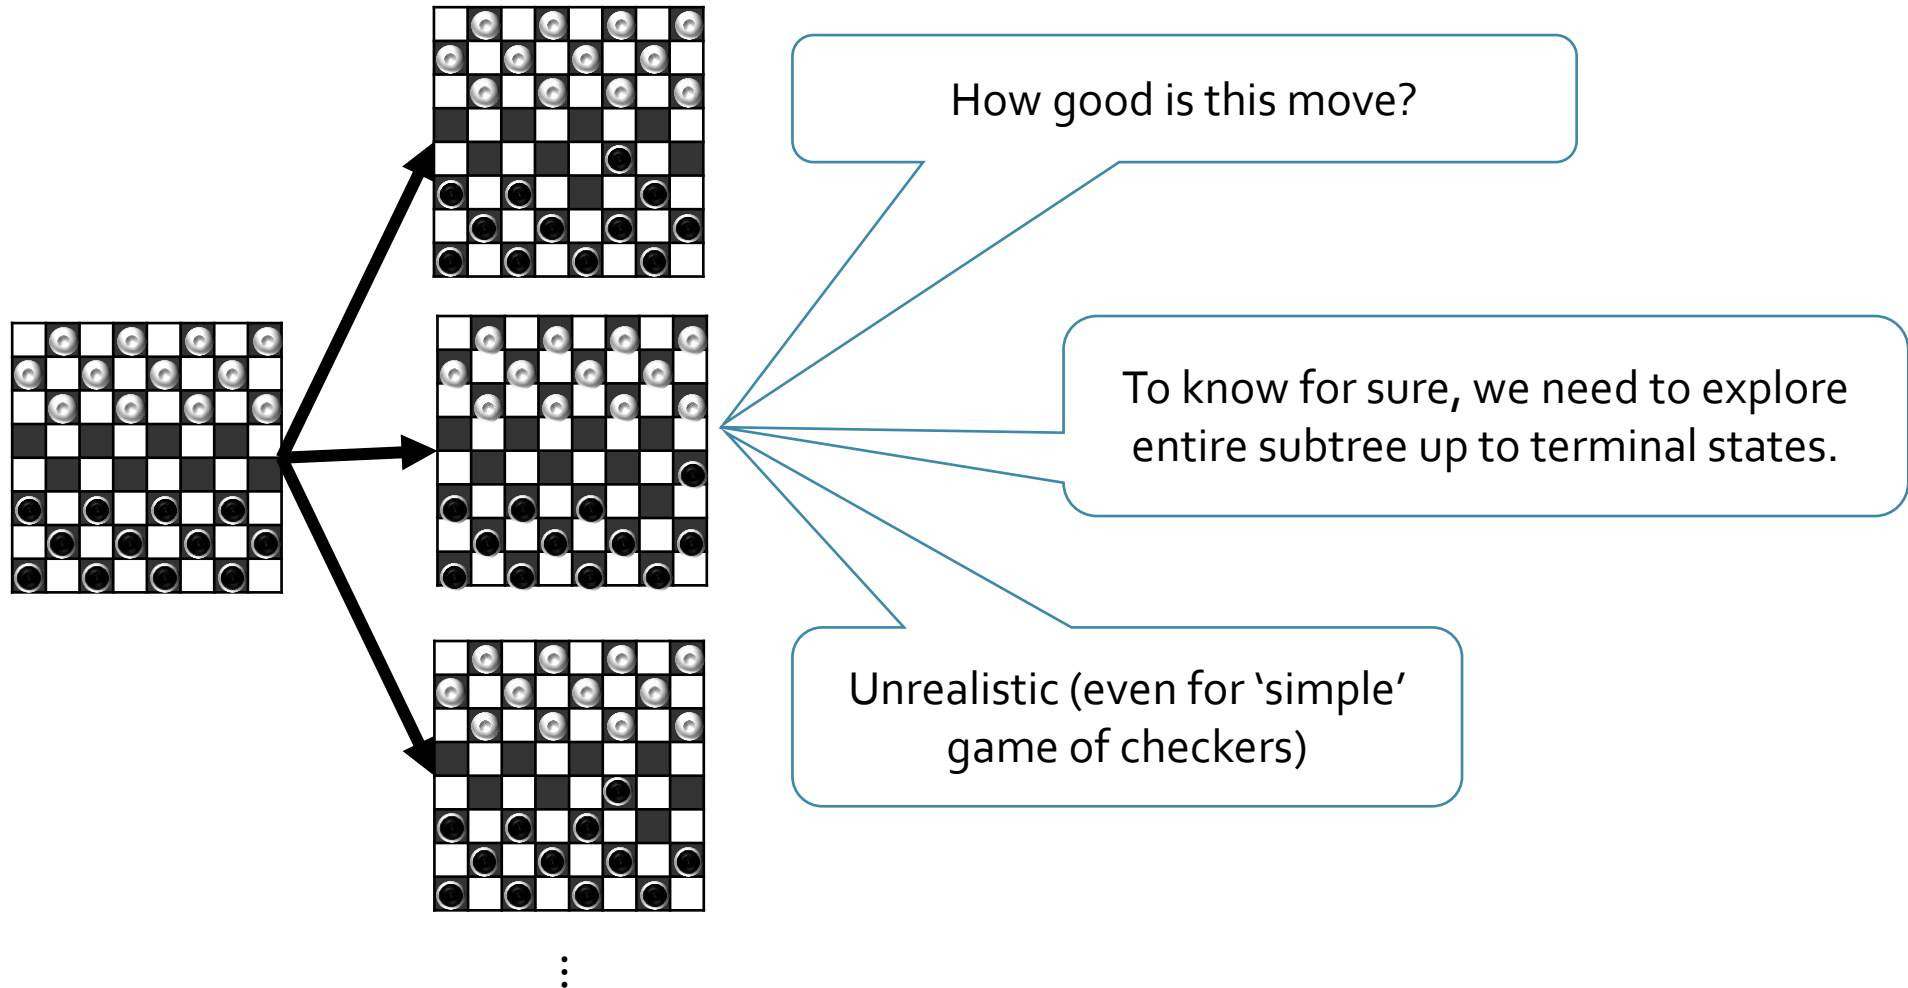
# Heuristic Minimax Value


$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s, \text{MAX}) & \text{if Is-Terminal}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MIN} \end{cases}$$


$$\text{H-Minimax}(s, d) = \begin{cases} \text{Eval}(s, \text{MAX}) & \text{if Cutoff-Test}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-Minimax}(\text{Result}(s, a), d + 1) & \text{if To-Move}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-Minimax}(\text{Result}(s, a), d + 1) & \text{if To-Move}(s) = \text{MIN} \end{cases}$$

Run Minimax until depth d; then start using the evaluation function to choose nodes


# Evaluation Functions – Checkers Example



# Evaluation Functions – Checkers Example



# Evaluation Functions – Checkers Example



# Evaluation Functions

- An evaluation function is a mapping from game states to real values
  - $f: \mathcal{S} \rightarrow \mathbb{R}$

- Default evaluation function:

$$f(s) = \begin{cases} \text{Utility}(s, \text{MAX}) & \text{if } \text{Is-Terminal}(s) \\ 0 & \text{otherwise} \end{cases}$$

No information on quality  
of non-terminal nodes

- Determine a function to estimate value that is strongly correlated to actual chances of winning
  - Modelling problem (similar to heuristic design problem from informed/local search)

# Evaluation Functions

---

- Determine important features/variables
- Chess example
  - # of pieces (*NPcs*)
  - # of queens (*NQns*)
  - # of controlled squares (*CtlSqs*)
  - # of threatened opponent pieces (*ThrPcs*)
  - ...
- $f(n) = w_1 \times (NPcs) + w_2 \times (NQns) + w_3 \times (CtlSqs) + w_4 \times (ThrPcs)$

Determine values for  $w_1, \dots, w_4$

# Cutting Off Search

---

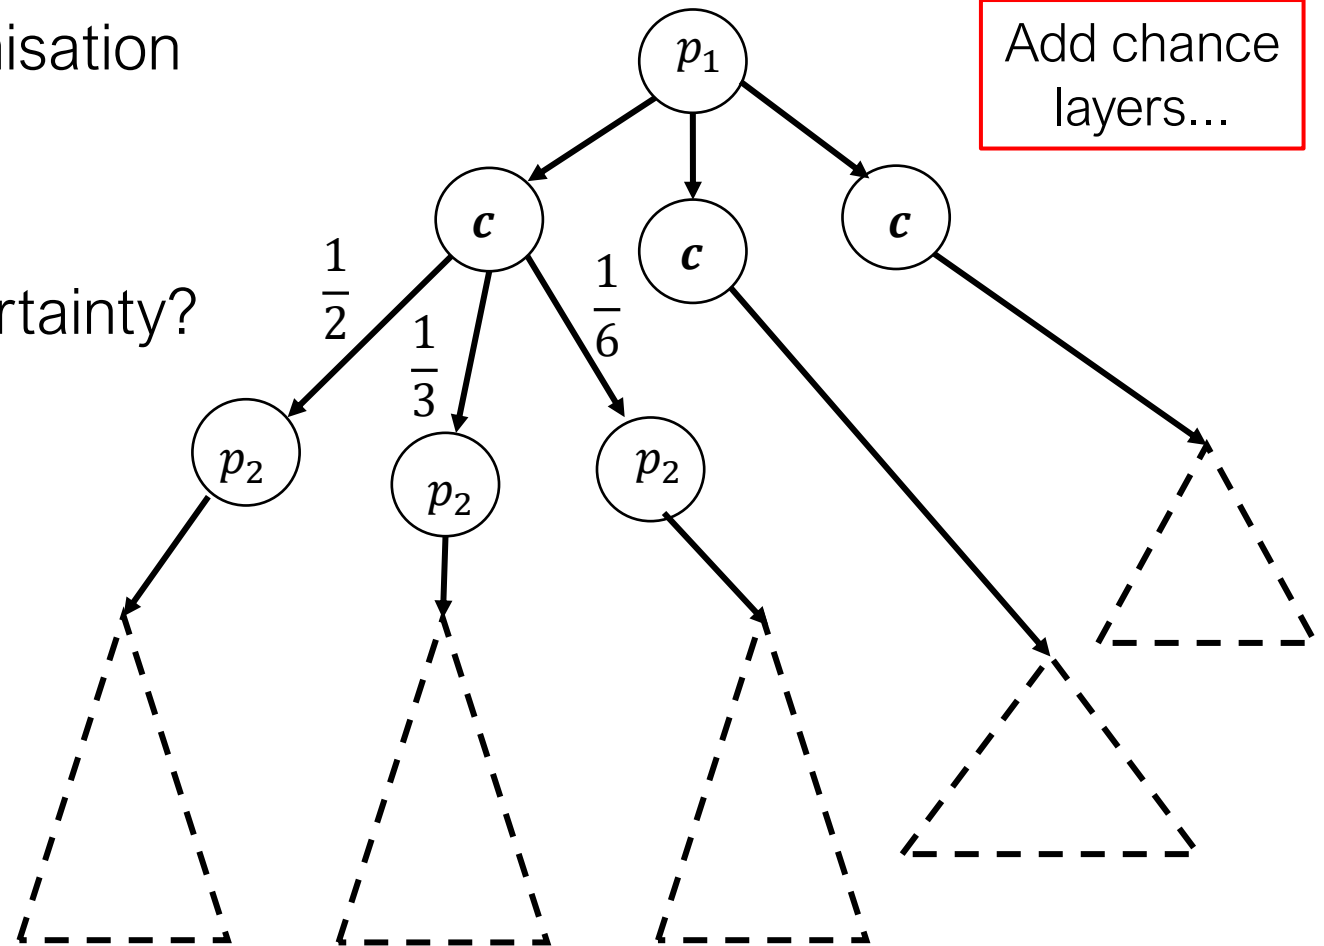
- Modify Minimax or  $\alpha$ - $\beta$  Pruning algorithms by replacing
  - Is-Terminal( $s$ ) with Cutoff-Test( $s, d$ )
  - Utility( $s, p$ ) with Eval( $s, p$ )
- Can replace DLS strategy with IDS



# Stochastic Games

- Many games have randomisation
  - Settlers of Catan
  - Poker
- How do we deal with uncertainty?
  - Can still use **Minimax**
  - Search tree is larger

Calculate expected value of a state (MUCH harder than deterministic games)



# Questions about the Lecture?

---

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
  - Specify a question
  - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)  
<https://archipelago.rocks/app/resend-invite/64238273017>