

Adversarial Search: Playing Games

CS3243: Introduction to Artificial Intelligence – Lecture 7

6 March 2023

Contents

1. Administrative Matters
2. Reviewing Search Problems
3. Adversarial Search Problems (i.e., Games)
4. Optimal Decisions via Minimax
5. α - β Pruning
6. Heuristic Minimax

Reference: AIMA 4th Edition, Section 6.1-6.3

Administrative Matters

Upcoming...

- Deadlines
 - None this week!
 - TA6 (released today)
 - *Due in your Week 6 tutorial session*
 - *Submit the a physical copy (more instructions on the Tutorial Worksheet)*
 - Prepare for the tutorial!
 - *Participation marks = 5%*
 - Midterm Marks Appeal
 - *Due next week (i.e., Week 9) on the day of your assigned tutorial session, 2359 hrs*
 - Project 2
 - *Due next week, Sunday (19 March), 2359 hrs*

Reviewing Search Problems

So Far...

- Path search (path planning)
 - Search for a path from start to goal
 - Complete: finds a solution or says when there isn't one
 - Optimal: path cost of path found is minimal
 - Uninformed
 - Systematically search all paths via general search problem formulation
 - Informed
 - Uses a heuristic to search less of the search space
- Goal search
 - Focus on goal and ignore path
 - Completeness consideration only
 - Local search
 - Uses heuristic to guide search to goal (uses restarts; many variants)
 - Constraint satisfaction problem
 - Uses specific search problem formulation and shrinks search space via inference

Games

Games and Search

- Can we solve games using existing methods?
 - In our searching thus far, we control all actions
 - All actions taken are determined by our agent
 - With games, your opponent decides actions too...
 - Multi-agent problem
 - Conventional planning \Rightarrow wasted computation since opponent can spoil your plans

Games and Search

- What is a game anyway?
 - Assume two players
 - Zero-sum game
 - Winner gets paid, and loser pays
 - We define
 - MAX player – player 1, who wants to maximise value (agent)
 - MIN player – player 2, who want to minimise value (opponent who wants agent to lose)
- General idea behind the search problem
 - Simulate play against utility maximising opponent
 - Find a strategy – i.e., define a move for every possible opponent response

Search Problem Formulation for Games

Formulating Games

- State representation

- As per general formulation



- TO-MOVE(s)

- Returns p , the player to move in state s

- ACTIONS(s)

- Legal moves in state s

- RESULT(s, a)

- The transition model; returns resultant state when taking action a at state s

- IS-TERMINAL(s)

- Returns TRUE when game is over and FALSE otherwise
 - States where game has ended are called terminal states

- UTILITY(s, p)

- Defines the final numeric value to player p when the game ends in terminal state s


Note on Utility

- Given zero-sum games
 - At terminal state s
 - $UTILITY(MAX, s) + UTILITY(MIN, s) = 0$

- Tic-Tac-Toe example


- X (agent) wins

- $UTILITY(s_i, MAX) = 1$
 - $UTILITY(s_i, MIN) = -1$



- O (opponent) wins

- $UTILITY(s_j, MAX) = -1$
 - $UTILITY(s_j, MIN) = 1$



- Draw


- $UTILITY(s_k, MAX) = 0$
 - $UTILITY(s_k, MIN) = 0$



Notice that here the utilities are relative to the player – however, we will standardise the scores such that they reflect only the MAX player (i.e., agent) scores later

Game Trees

Example Game Tree

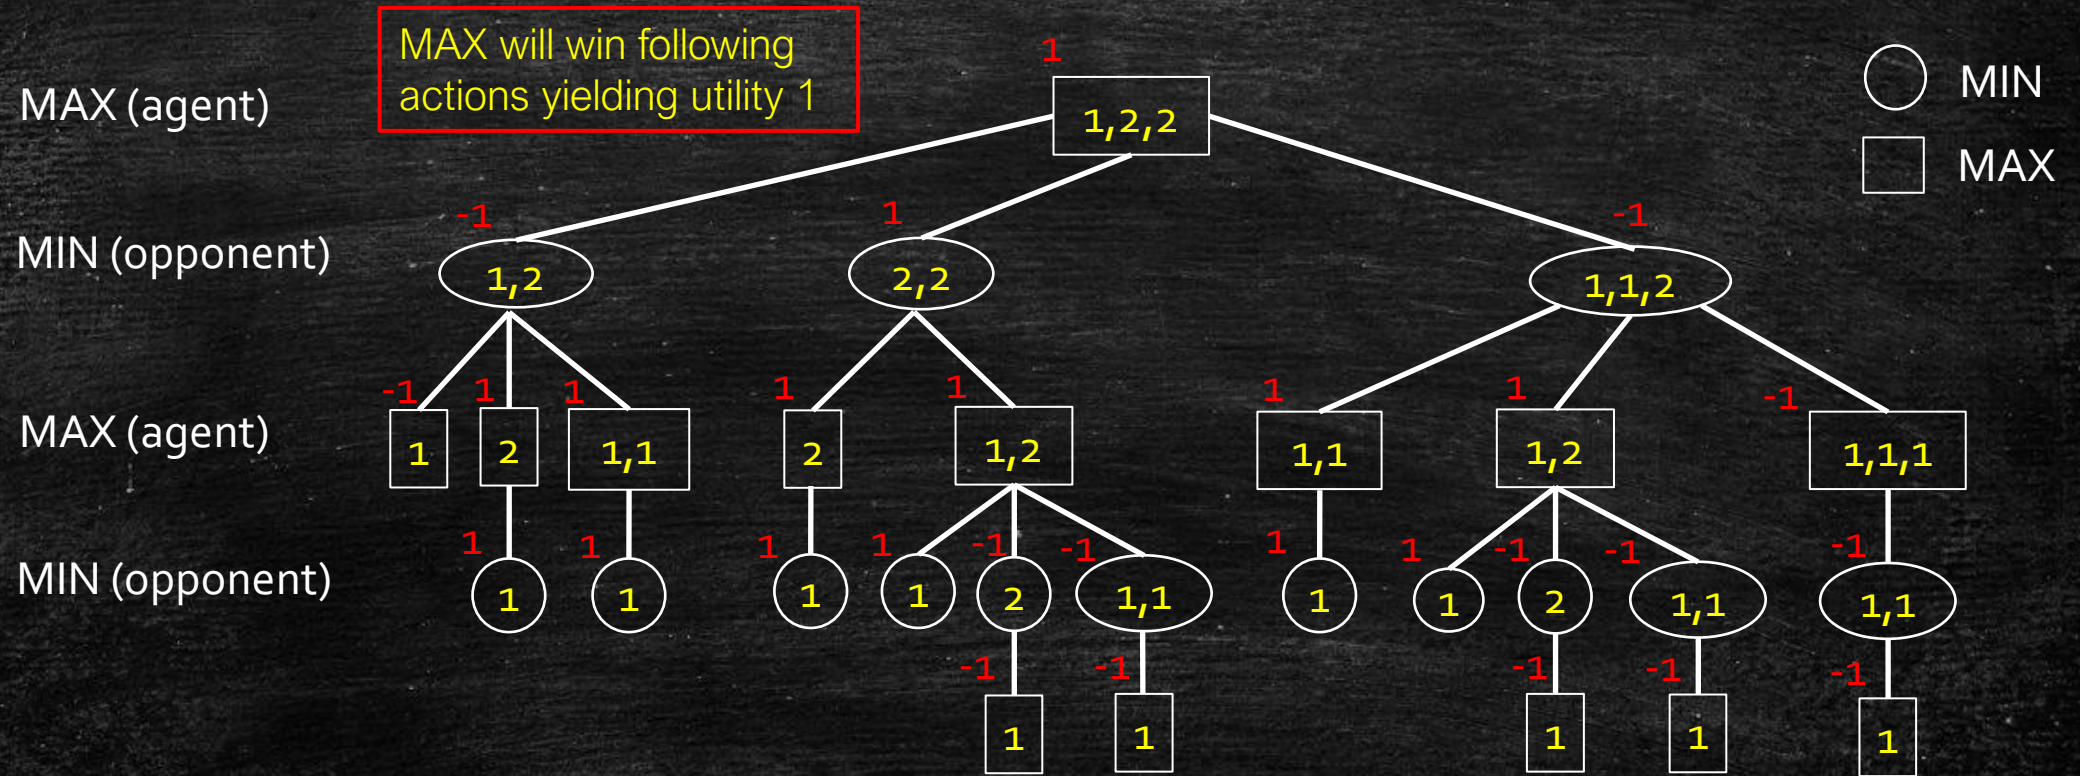


- More on environment characteristics
 - 2-player
 - Deterministic
 - Turn-taking
- Zero-sum implications
 - Loser for every winner
 - Agent utilises **sum to zero**
 - May also consider **constant-sum game**
 - Completely **adversarial game**

Another Example: Game of NIM

- Several piles of sticks are given
 - Represent the configuration of piles by a monotone sequence of integers
 - Example: (1,3,5)
 - With each turn, a player may remove any number of sticks from ONE pile
 - Example:
 - Remove 4 sticks from last pile (of 5 sticks) \Rightarrow (1,3,5) becomes (1,1,3)
 - The player who takes the last stick loses
- Let's try...
 - Represent the NIM game (1,2,2) as a game tree

Game of NIM: (1,2,2) Game Tree



DFS Traversal with Backwards induction on utility

Strategies

Player Strategies

- A strategy s for *player i* :
 - What will *player i* do at every node of the game tree that they make a move in?
 - Need to specify behaviour in states that may never be reached!

- Winning strategy

A strategy s_1^* for **Player 1** is called winning if for any strategy s_2 by **Player 2**, the game ends with **Player 1** as the winner.

- Non-losing strategy

A strategy t_1^* for **Player 1** is called non-losing if for any strategy s_2 by **Player 2**, the game ends in either a tie or a win for **Player 1**.


Optimal Strategy at Node - Minimax

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s, \text{MAX}) & \text{if Is-Terminal}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MIN} \end{cases}$$


- Intuitively
 - MAX chooses move to maximise the minimum payoff
 - MIN chooses at successors
 - MIN chooses move to minimise the maximum payoff
 - MAX chooses at successors

Remember that $\text{Result}(s, a)$ outputs the utility in terms of MAX player's utility (i.e., MIN wants to minimise this value, while MAX wants to maximise it)

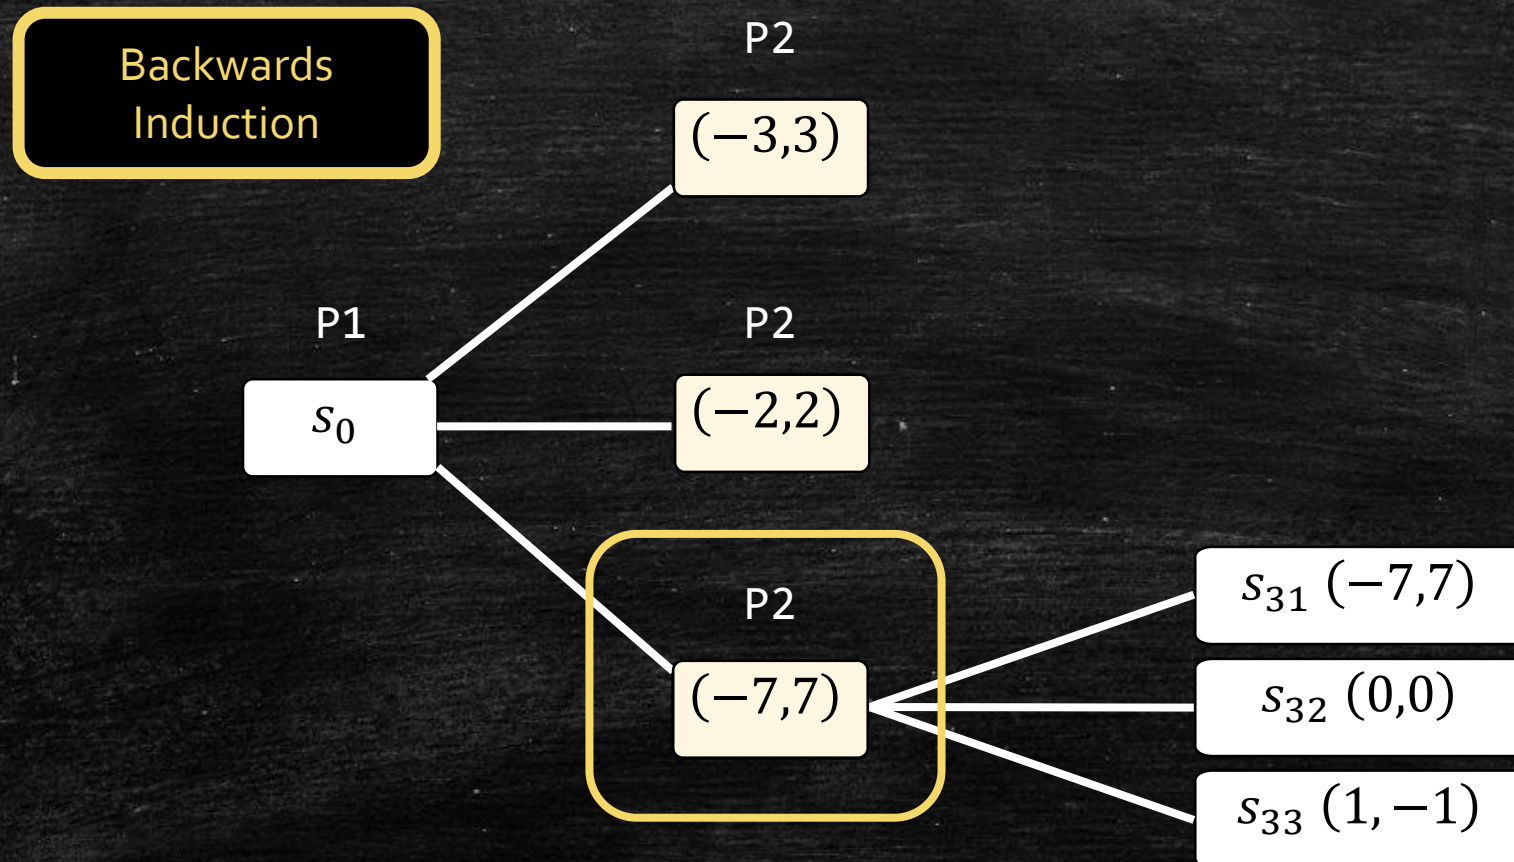
Minimax Play – Subgame Perfect Nash Equilibrium



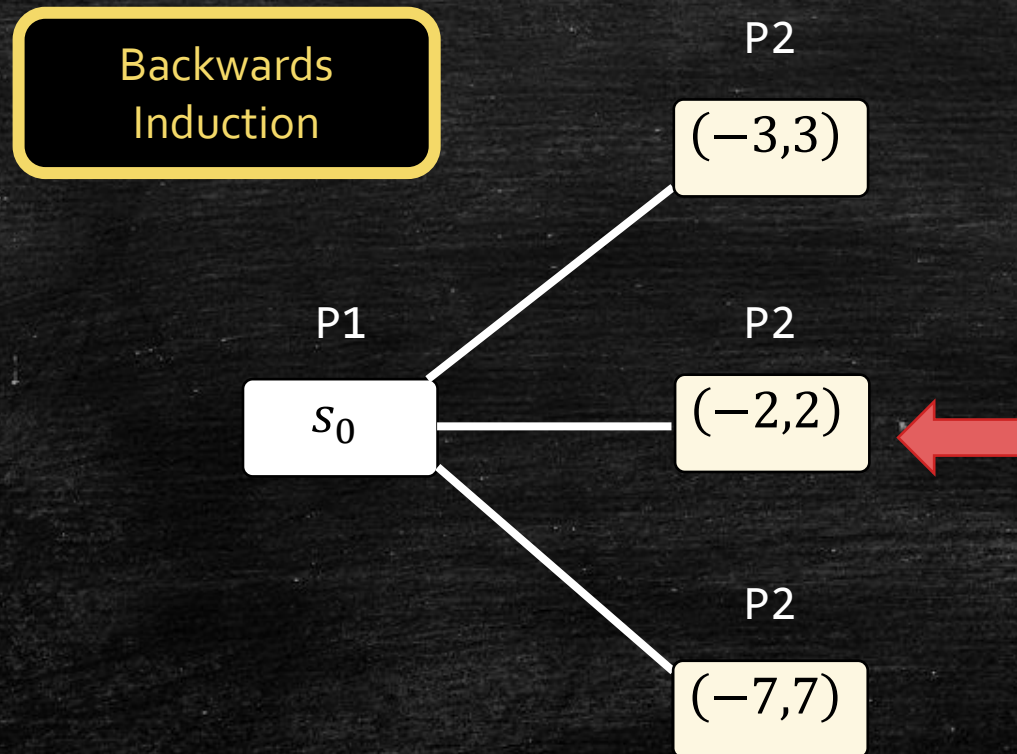
Minimax Play – Subgame Perfect Nash Equilibrium



Minimax Play – Subgame Perfect Nash Equilibrium

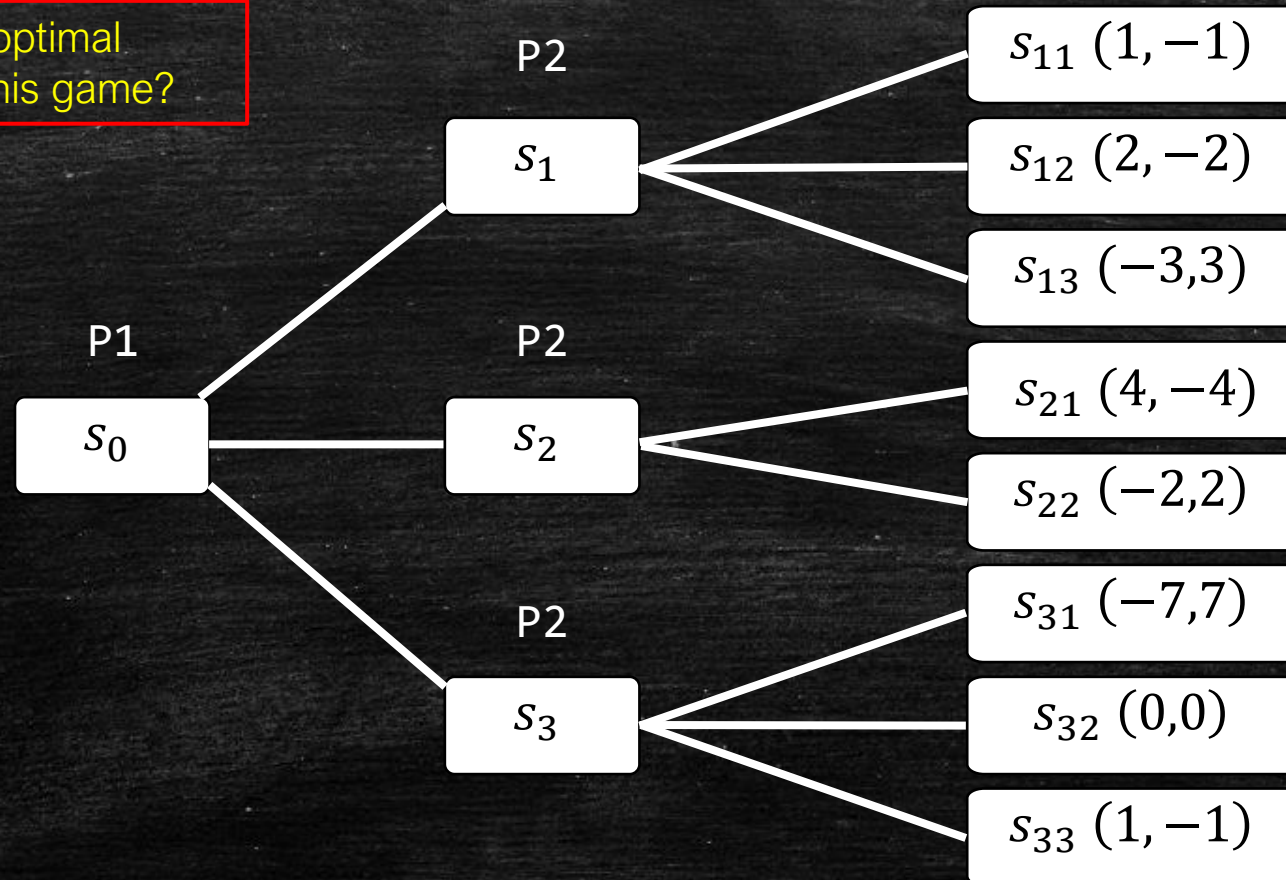


Minimax Play – Subgame Perfect Nash Equilibrium



Minimax Play – Subgame Perfect Nash Equilibrium

What are the optimal strategies in this game?



Minimax Algorithm Properties

- **Complete?**
 - Yes (if game tree is finite)
- **Optimal?**
 - Yes (optimal gameplay)
- **Time**
 - $O(b^m)$
- **Space**
 - $O(bm)$
- Minimax runs in time polynomial in tree size
- Returns a subgame perfect Nash Equilibrium
 - i.e., the best action at every node

Are we done?

Backwards Induction

- Game trees are massive
 - Chess has a massive game tree
 - 10^{123} nodes
 - In comparison, planet Earth has about 10^{50} atoms ...
- Impossible to expand the entire tree
- Have to find ways to shrink the search tree
 - We've seen this before
 - Common theme in search

Questions about the Lecture?

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
 - Specify a question
 - Upvote someone else's question




Invitation Link (Use NUS Email --- starts with E)
<https://archipelago.rocks/app/resend-invite/64238273017>

α - β Pruning

α - β Pruning - General Idea


- Basic idea
 - Don't explore moves that would never be considered
- Maintain bounds on values seen thus far while searching
 - α bounds MAX's values
 - β bounds MIN's values
- Prune subtrees that will never affect Minimax decision

α - β Pruning – Example Trace




Note that from this point, all utility values will be specified in terms of the MAX player.

α - β Pruning – Example Trace




Note that from this point, all utility values will be specified in terms of the MAX player.

α - β Pruning – Example Trace




Note that from this point, all utility values will be specified in terms of the MAX player.

α - β Pruning – Example Trace



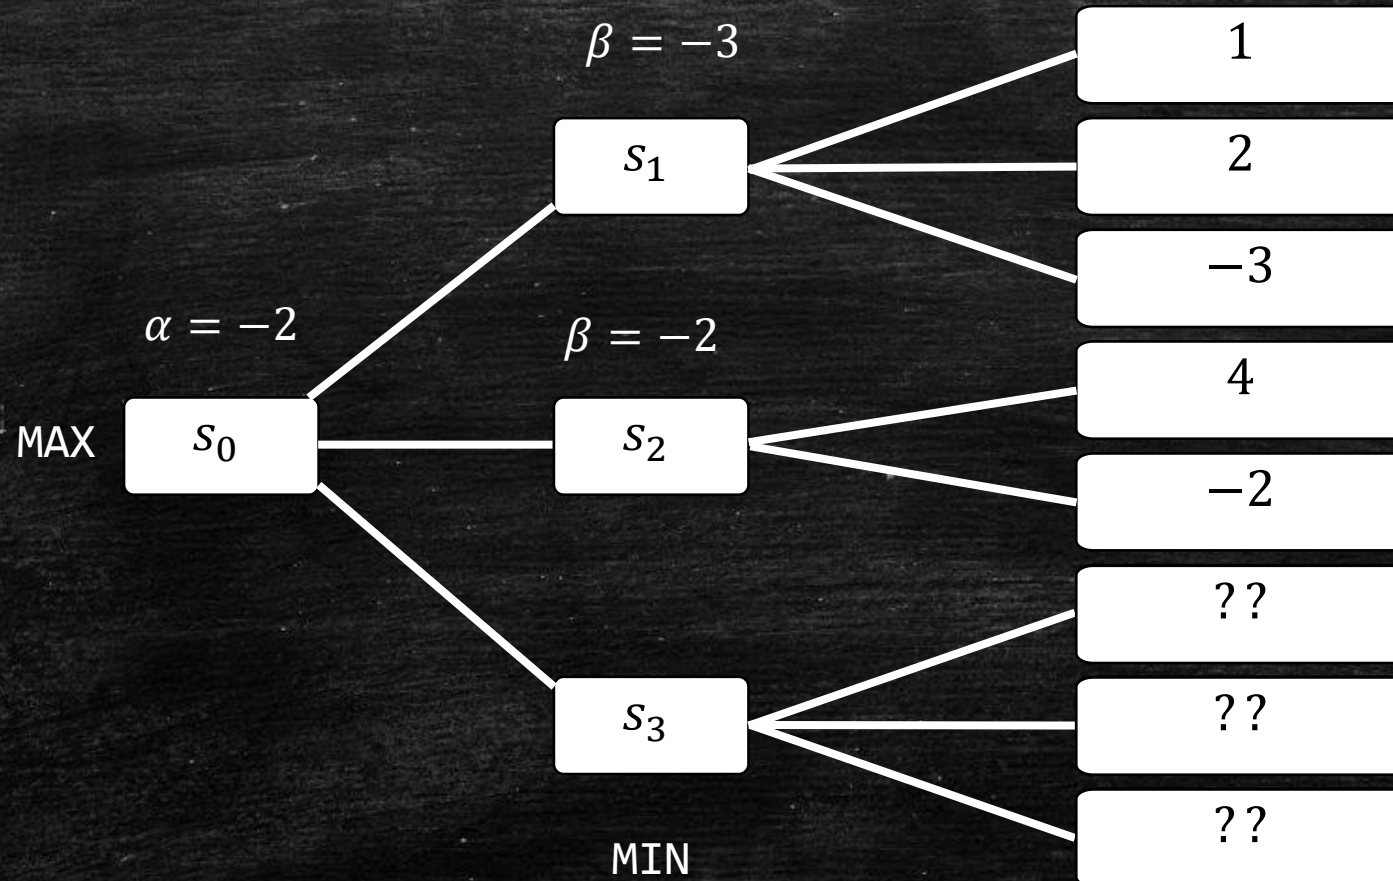
Note that from this point, all utility values will be specified in terms of the MAX player.

α - β Pruning – Example Trace




Note that from this point, all utility values will be specified in terms of the MAX player.

α - β Pruning – Example Trace



Note that from this point, all utility values will be specified in terms of the MAX player.


α - β Pruning – Example Trace



Note that from this point, all utility values will be specified in terms of the MAX player.

No point exploring the last two nodes! Choosing s_3 results in a loss of at least 7...

α - β Pruning – Tic-Tac-Toe Example



α - β Pruning

- MAX node n
 - $\alpha(n)$ = highest observed value found on path from n
 - Initially $\alpha(n) = -\infty$
- MIN node n
 - $\beta(n)$ = lowest observed value found on path from n
 - Initially $\beta(n) = +\infty$
- Pruning rules
 - Given a MIN node n , stop searching below n if
 - Some MAX ancestor i (of n) with $\alpha(i) \geq \beta(n)$
 - Given a MAX node n , stop searching below n if
 - Some MIN ancestor i (of n) with $\beta(i) \leq \alpha(n)$

MIN will choose $\beta(n)$ or lower at n , but ancestor MAX will NEVER choose the subtree at n since at i , there is a better option with higher value $\alpha(i)$

MAX will choose $\alpha(n)$ or higher at n , but ancestor MIN will NEVER choose the subtree at n since at i , there is a better option with lower value $\beta(i)$

α - β Pruning Analysis

- Pruning a branch never affects the final outcome
- Good move ordering improves effectiveness of pruning
 - “Perfect” ordering
 - Time complexity $O(b^{m/2})$
 - Good pruning strategies allow us to search twice as deep!
 - Example: Chess
 - Simple ordering gets you close to best-case result
 - Checks
 - Take pieces
 - Forward moves
 - Backwards move
- Expansion-order heuristics will improve the search
- Random ordering gives complexity $O(b^{3m/4})$ for $b < 1000$

Issue with α - β Pruning

- Original Problem
 - Most games have very large game trees
- Solution
 - α - β pruning can remove large parts of search trees
- Unresolved Issue
 - Maximum depth of tree
 - Backwards induction works backwards from terminal states
 - Still have to traverse to a terminal states
 - Standard solution – Heuristic Minimax
 - Cutoff test – e.g., depth limit (DLS)
 - Evaluation function – estimates expected utility of state

Heuristic Minimax


Heuristic Minimax Value

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s, \text{MAX}) & \text{if Is-Terminal}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if To-Move}(s) = \text{MIN} \end{cases}$$


$$\text{H-Minimax}(s, d) = \begin{cases} \text{Eval}(s, \text{MAX}) & \text{if Cutoff-Test}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-Minimax}(\text{Result}(s, a), d + 1) & \text{if To-Move}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-Minimax}(\text{Result}(s, a), d + 1) & \text{if To-Move}(s) = \text{MIN} \end{cases}$$

Run **Minimax** until **depth d**; then start using the **evaluation function** to choose nodes

Evaluation Functions – Checkers Example




How good is this move?


To know for sure, we need to explore entire subtree up to terminal states.

Unrealistic (even for 'simple' game of checkers)

Evaluation Functions – Checkers Example



Evaluation Functions – Checkers Example



Pretend these
are **terminal**
nodes with
values given
by $f(\cdot)$

Evaluation Functions

- An evaluation function is a mapping from game states to real values
 - $f: \mathcal{S} \rightarrow \mathbb{R}$

- Default evaluation function:

$$f(s) = \begin{cases} \text{Utility}(s, \text{MAX}) & \text{if } \text{Is-Terminal}(s) \\ 0 & \text{otherwise} \end{cases}$$

No information on quality of non-terminal nodes

- Determine a function to estimate value that is strongly correlated to actual chances of winning
 - Modelling problem (similar to heuristic design problem from informed/local search)

Evaluation Functions

- Determine important features/variables
- Chess example
 - # of pieces (*NPcs*)
 - # of queens (*NQns*)
 - # of controlled squares (*CtlSqs*)
 - # of threatened opponent pieces (*ThrPcs*)
 - ...
- $f(n) = w_1 \times (NPcs) + w_2 \times (NQns) + w_3 \times (CtlSqs) + w_4 \times (ThrPcs)$

Determine values for w_1, \dots, w_4

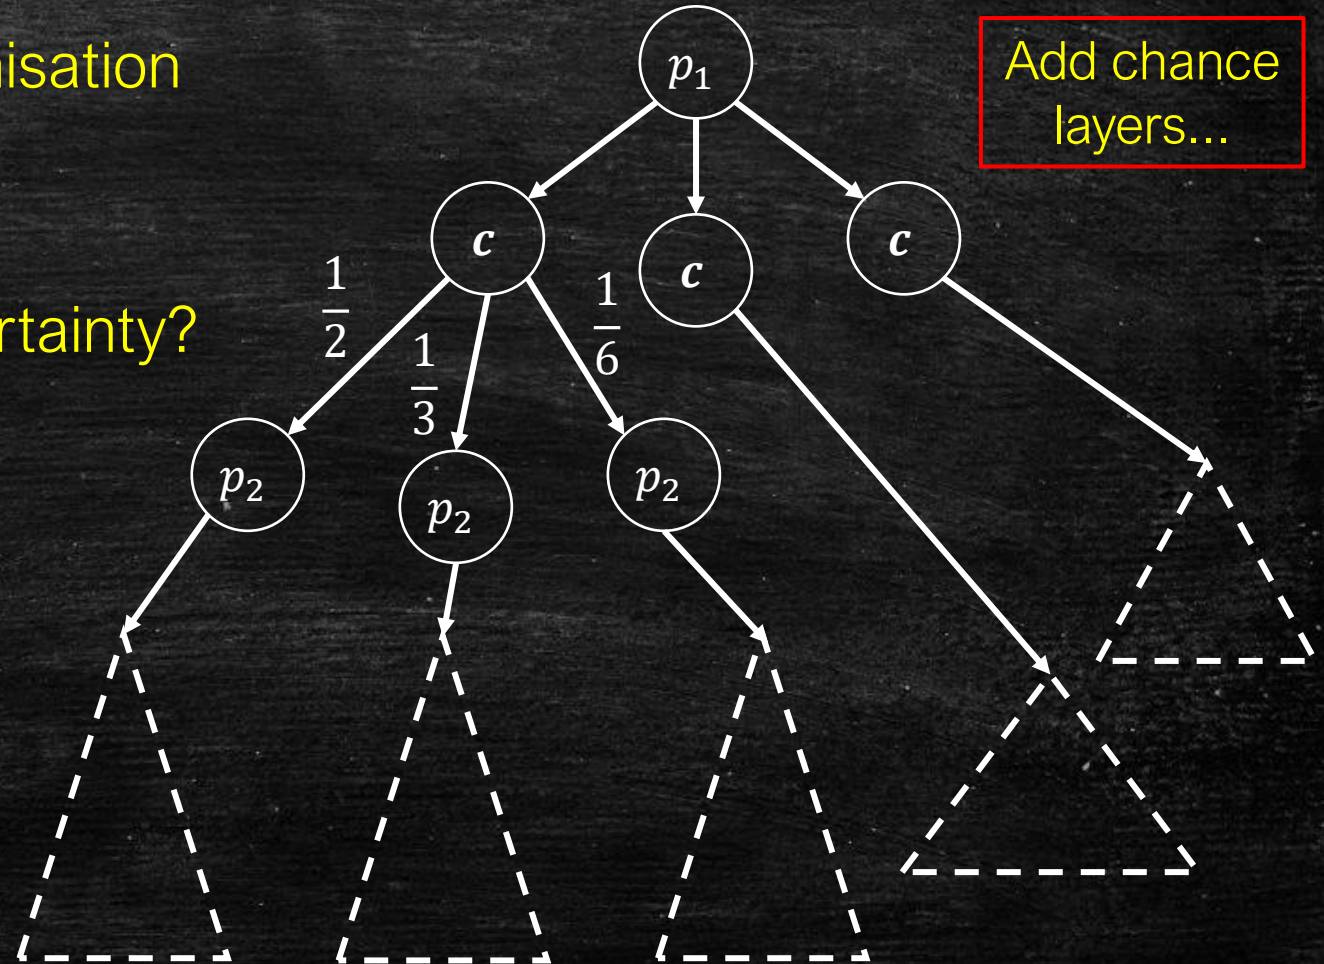
Cutting Off Search

- Modify **Minimax** or **α - β Pruning** algorithms by replacing
 - **Is-Terminal(s)** with **Cutoff-Test(s, d)**
 - **Utility(s, p)** with **Eval(s, p)**
- Can replace **DLS** strategy with **IDS**

Stochastic Games

- Many games have randomisation
 - Settlers of Catan
 - Poker
- How do we deal with uncertainty?
 - Can still use **Minimax**
 - Search tree is larger

Calculate expected value of a state (MUCH harder than deterministic games)



Questions about the Lecture?

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
 - Specify a question
 - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)
<https://archipelago.rocks/app/resend-invite/64238273017>