

# Logical Agents: Knowledge Representation II

---

CS3243: Introduction to Artificial Intelligence – Lecture 9a

20 March 2023

# Contents

---

1. Administrative Matters
2. Recap on Logical Agents
3. Theorem-Proving Methods
4. Resolution
5. Uncertainty & Recap on Probability

# Administrative Matters

---

# Upcoming...

---

- Deadlines

- TA7 (released last week)
    - *Due in your Week 10 tutorial session*
    - *Submit the a physical copy (more instructions on the Tutorial Worksheet)*
  - Prepare for the tutorial!
    - Participation marks = 5%
  - Project 2 (released Week 6)
    - *Was due yesterday, Sunday (19 March), 2359 hrs*
    - *Late penalties now apply*
  - Project 3 (released last week)
    - *Due Week 12 Sunday (9 April), 2359 hrs*
- Week 11: Final Content Lecture
  - Week 12: Lecture Slot = Project 3 Consultation (conducted over Zoom)

# Recap on Logical Agents

---

# Logical Agents

---

- Agent contains
  - Knowledge Base (KB)
    - Specified in some language (e.g., propositional logic)
  - Inference Engine (IE)
    - Determines sentences that will guide action choice,  $\alpha_1, \alpha_2, \dots, \alpha_k$
    - Uses an algorithm that infers  $\alpha_i$  such the  $KB \models \alpha_i$
- General algorithm
  - Pre-populate KB with domain knowledge
  - Each time step  $t$ :
    - Update KB with percepts
    - Use IE to make inferences
      - Update KB with inferences
      - Select action based on inferences
  - Take action and update KB with new state (current truth value assignments)

# Logical Agents

Logical Agent **cannot** plan an entire path to goal  
Environment is only **Partially Observable**

Assume operation as follows

Make inferences about environment

Assume query ( $\alpha$ ) to action ( $a_i$ ) mappings  
e.g., reflex agent with conditions based on KB

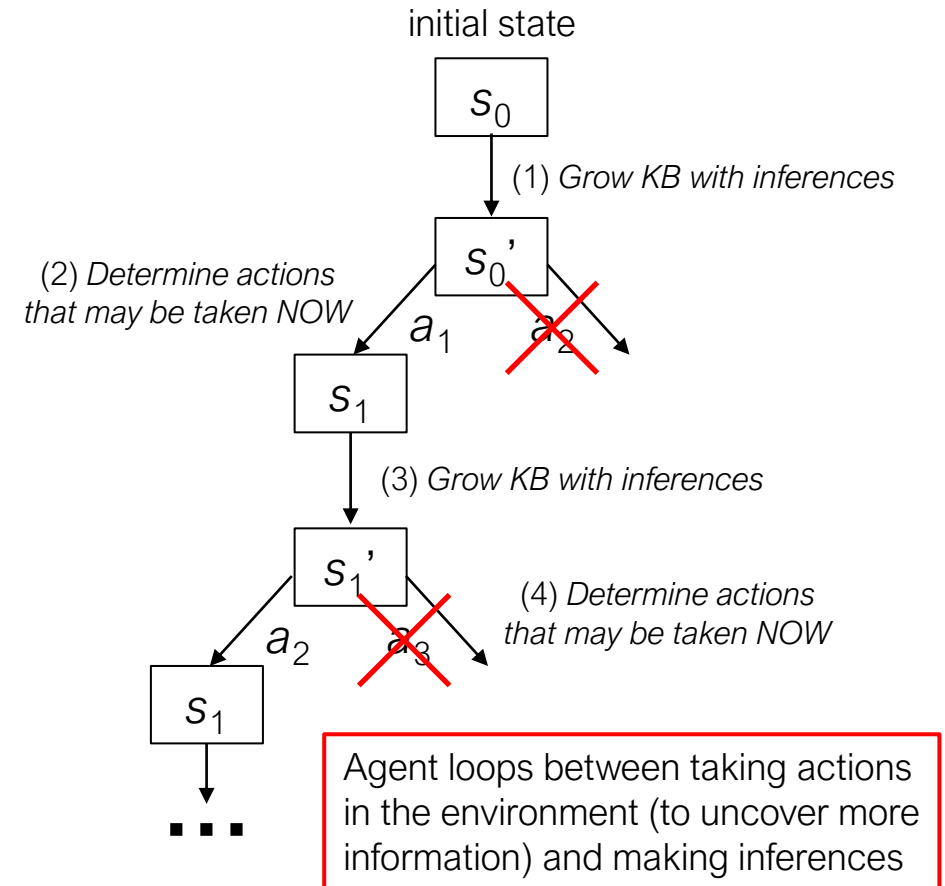
Developer designing the agent must  
formulate these mappings

Suppose  $\alpha \Rightarrow a_i$

Agent uses the KB & IE to determine if a query ( $\alpha$ )  
may be inferred, and if so, takes associated action ( $a_i$ )

Chosen action ( $a_i$ ) is EXECUTED in environment  
(as we are no longer planning)

Loops...



# Making Inferences

- Entailment ( $\models$ )
  - $KB \models \alpha$  means that  $M(KB) \subseteq M(\alpha)$ 
    - This says that all value assignments that satisfy the KB will also satisfy  $\alpha$
    - i.e., whenever KB is true,  $\alpha$  is true
- Inference algorithm ( $\mathcal{A}$ )
  - Sound:  $(KB \vdash_{\mathcal{A}} \alpha) \Rightarrow (KB \models \alpha)$ 
    - $\mathcal{A}$  only infers  $\alpha$  that are valid
  - Complete:  $(KB \models \alpha) \Rightarrow (KB \vdash_{\mathcal{A}} \alpha)$ 
    - $\mathcal{A}$  is able to infer all valid  $\alpha$
- Inference algorithm example: Truth Table Enumeration
  - Construct entire truth table for KB and  $\alpha$
  - Check (via DFS) that  $M(KB) \subseteq M(\alpha)$
  - i.e., every model of KB is a model of  $\alpha$

Truth Table Enumeration:

- Sound and Complete
- Time complexity  $O(2^n)$
- Space complexity  $O(n)$



# Theorem Proving Methods

---

# Proof Methods

---

- Model checking (special case of CSPs where domains are T/F)
  - Truth Table Enumeration (time complexity exponential in  $n$ )
  - Resolution (inference via proof)
- Applying inference rules (i.e., theorem proving)
  - Generate new sentences from old
  - Proof = sequential application of inference rules
    - Inference rules help deduce valid actions
    - Proof facilitates efficiency – ignores irrelevant propositions

# Validity & Satisfiability

- A sentence  $\alpha$  is valid if it is *true for ALL possible truth value assignments*
  - e.g., True,  $A \vee \neg A$ ,  $A \Rightarrow A$ ,  $(A \wedge (A \Rightarrow B)) \Rightarrow B$
  - i.e., tautologies
- Validity is connected to entailment via the Deduction Theorem:
  - $(KB \models \alpha) \Leftrightarrow ((KB \Rightarrow \alpha) \text{ is valid})$
- A sentence is satisfiable if it is *true for SOME truth value assignment*
  - e.g.,  $A \vee B$ ,  $C$
- A sentence is unsatisfiable if it is *true for NO truth value assignments*
  - e.g.,  $A \wedge \neg A$
  - i.e., contradictions
- Satisfiability is connected to entailment via the following:
  - $(KB \models \alpha) \Leftrightarrow ((KB \wedge \neg \alpha) \text{ is unsatisfiable})$
  - i.e., definition of Proof by Contradiction

i.e., a model exists  
for that sentence

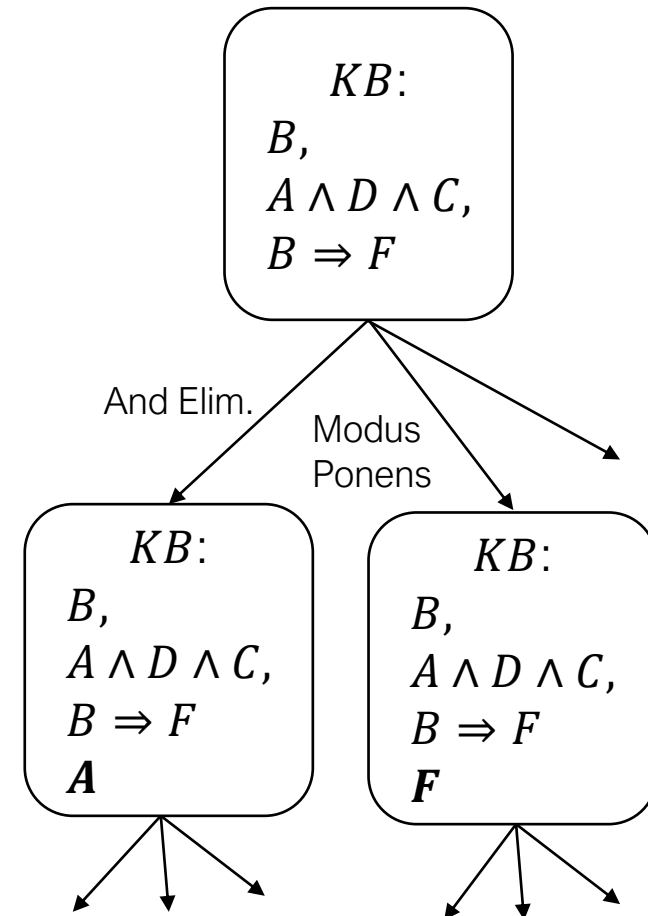
i.e., no model exists  
for that sentence

- $A \Rightarrow B \equiv \neg A \vee B$
- $\neg(\neg A \vee B) \equiv A \wedge \neg B$
- Showing that  $(A \wedge \neg B)$  is unsatisfiable shows that the negation,  $\neg A \vee B$  is valid!

# Inference Algorithms: Application of Inference Rules

- Search for more knowledge (growing the KB)
  - Equivalent to a search problem
    - States: Versions of the KB (e.g., initial state is initial KB)
    - Actions: Application of inference rules
    - Transition: Update KB with an inferred sentence (may not be target one)
    - Goal: KB contains sentence to (dis)prove (e.g., given query  $\alpha$ )
- Examples of inference rules
  - And-Elimination (AE): e.g.,  $a \wedge b \models a$ ;  $a \wedge b \models b$
  - Modus Ponens (MP): e.g.,  $a \wedge (a \Rightarrow b) \models b$
  - Logical Equivalences: e.g.,  $(a \vee b) \models \neg(\neg a \wedge \neg b)$

How does this relate to Truth Table Enumeration?



# Resolution

---

# Resolution for Conjunctive Normal Form (CNF)

- CNF = conjunction of disjunctive sentences
  - e.g.,  $(x_1 \vee \neg x_2) \wedge (x_2 \vee x_3 \vee \neg x_4)$
- Resolution
  - Method of simplifying KB to prove entailment of query  $\alpha$
  - Specifically
    - Given KB:  $R_1 \wedge R_2 \wedge \dots \wedge R_n$
    - If a literal,  $x$ , appears in  $R_i$  and its negation,  $\neg x$ , appears in  $R_j$ , where  $R_i, R_j \in \text{KB}$ , it can be removed from both

$$\text{resolvent} \quad \frac{(x_1 \vee \dots \vee x_m \vee x) \wedge (y_1 \vee \dots \vee y_k \vee \neg x)}{(x_1 \vee \dots \vee x_m \vee y_1 \vee \dots \vee y_k)}$$

- Resolution under propositional logic
  - Sound
  - Complete

KB:  $(P \text{ or } x) \text{ and } (Q \text{ or } (\text{not } x))$   
 $\alpha$ :  $(P \text{ or } Q)$  must hold?

$x=T$ ,  $Q$  must be True for KB to hold  
 $x=F$ ,  $P$  must be True for KB to hold

So  $(P \text{ or } Q)$  must hold

Verify with truth table as an exercise  
(note: we want  $M(\text{KB}) \subseteq M(\alpha)$ )

# Some Rules for Conversion to CNF

1. Convert  $\alpha \Leftrightarrow \beta$  to  $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$
2. Convert  $\alpha \Rightarrow \beta$  to  $\neg \alpha \vee \beta$
3. Expand  $\neg$  inwards using De Morgan and double negation
  - a. Convert  $\neg(\alpha \vee \beta)$  to  $\neg \alpha \wedge \neg \beta$
  - b. Convert  $\neg(\alpha \wedge \beta)$  to  $\neg \alpha \vee \neg \beta$
  - c. Convert  $\neg(\neg \alpha)$  to  $\alpha$
4. Convert  $(\alpha \vee (\beta \wedge \gamma))$  to  $(\alpha \vee \beta) \wedge (\alpha \vee \gamma)$

Each of these conversions produces two rules in the KB (the others just one)

# Resolution Algorithm

- Utilises proof by contradiction – tries to show that  $KB \wedge \neg\alpha$  is unsatisfiable

**function** PL-RESOLUTION( $KB, \alpha$ ) **returns** *true* or *false*

**inputs:**  $KB$ , the knowledge base, a sentence in propositional logic  
 $\alpha$ , the query, a sentence in propositional logic

$clauses \leftarrow$  the set of clauses in the CNF representation of  $KB \wedge \neg\alpha$

$new \leftarrow \{\}$

**while true do**

**for each** pair of clauses  $C_i, C_j$  **in**  $clauses$  **do**

$resolvents \leftarrow$  PL-RESOLVE( $C_i, C_j$ )

**if**  $resolvents$  contains the empty clause **then return true**

$new \leftarrow new \cup resolvents$

**if**  $new \subseteq clauses$  **then return false**

$clauses \leftarrow clauses \cup new$

If cannot be resolved further and not empty clause – cannot infer  $\alpha$

What does an empty clause imply??

Suppose we have a KB as follows:

$$\frac{(x_1 \vee \dots \vee x_m \vee x) \wedge (y_1 \vee \dots \vee y_k \vee \neg x)}{(x_1 \vee \dots \vee x_m \vee y_1 \vee \dots \vee y_k)}$$

And the algorithm slowly removes literals:

$$\frac{(\cancel{x_1} \vee \dots \vee \cancel{x_m} \vee x) \wedge (\cancel{y_1} \vee \dots \vee \cancel{y_k} \vee \neg x)}{(\cancel{x_1} \vee \dots \vee \cancel{x_m} \vee \cancel{y_1} \vee \dots \vee \cancel{y_k})}$$

Eventually, there is nothing in the KB.

KB indicates the disjunction of no literals holds. A disjunction is True only when at least one literal is true. So, whole KB is False here – i.e., the query  $\neg\alpha$  is unsatisfiable.

We may infer  $\alpha$  (via proof by contradiction).



# Resolution Algorithm

- Summary
  - Make a clause list – i.e., copy of KB specified in CNF including negation of query,  $\neg\alpha$ 
    - Use conversion rules to convert KB to CNF
  - Repeatedly resolve two clauses from clause list
    - Add resolvent to clause list
  - Keep doing this till empty clause found or no more resolutions possible
    - If empty clause then can infer  $\alpha$
    - If no more resolutions and not empty clause then cannot infer  $\alpha$

Why is Resolution under Propositional Logic Sound and Complete?

Soundness:

- Each resolvent is implied by generating clauses
- If  $\emptyset$  is found, then  $(KB \wedge \neg\alpha)$  is unsatisfiable

Completeness:

- Based on the idea of resolution closure – set of all clauses derivable
- Not covered in CS3243
- Refer to AIMA 4<sup>th</sup> Edition pp. 228-229

# Resolution Example

---

# Resolution Example: Back to Wumpus World

- Assume that agent is at (1,1) in Wumpus World
  - And we wish to make inferences about a pit at (1,2)

- KB

- $$(B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge (\neg B_{1,1})$$
- i.e., we know
  - $R_a$ : rule for breezes
  - $R_b$ : no breeze at (1,1)

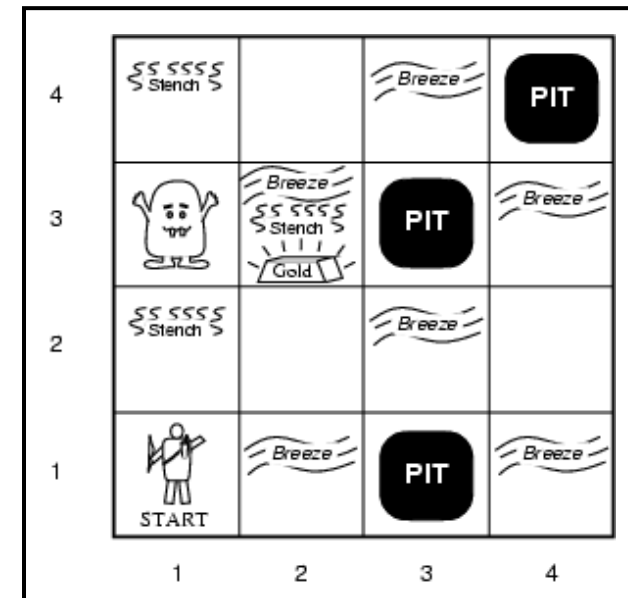
- $\alpha$

- $\neg P_{1,2}$
- i.e., want to know if we can move to (1,2)

- Can we infer  $\alpha$ ?

- Use the resolution algorithm to determine if  $(KB \Rightarrow \alpha)$
- i.e., use  $(KB \Rightarrow \alpha) \Leftrightarrow (KB \wedge \neg \alpha)$  is unsatisfiable

1,2	2,2
P?	
1,1	2,1
A OK	



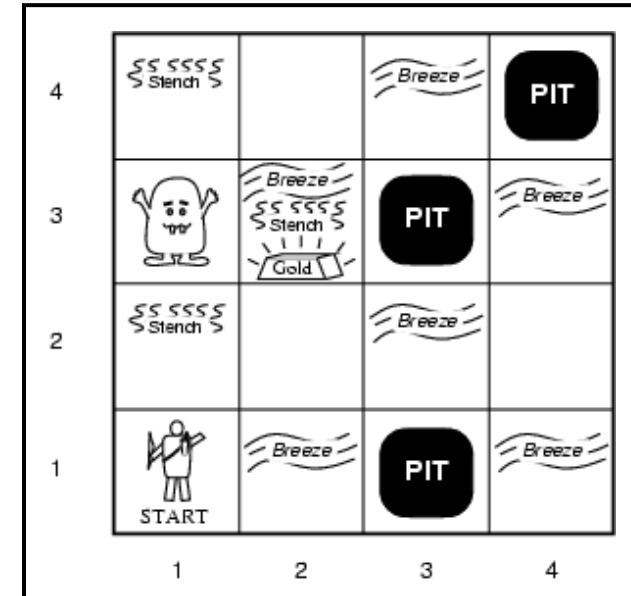
# Resolution Example: Back to Wumpus World

- Given
  - $KB = \neg B_{1,1} \wedge B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$
  - $\alpha = \neg P_{1,2}$
- Step 1 – Form clause list (over  $KB \wedge \neg\alpha$ )
  - $(\neg B_{1,1}) \wedge (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge (P_{1,2})$
- Step 2 – Convert clause list to CNF
  - $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$ 
    - $B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})$
    - $(P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1}$
- Step 2a
  - $B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})$ 
    - $\neg B_{1,1} \vee (P_{1,2} \vee P_{2,1})$
    - $\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}$

$(\neg B_{1,1}), (P_{1,2})$   
already in  
CNF (literals)

Now in CNF

1,2	2,2
P?	
1,1	2,1
A OK	



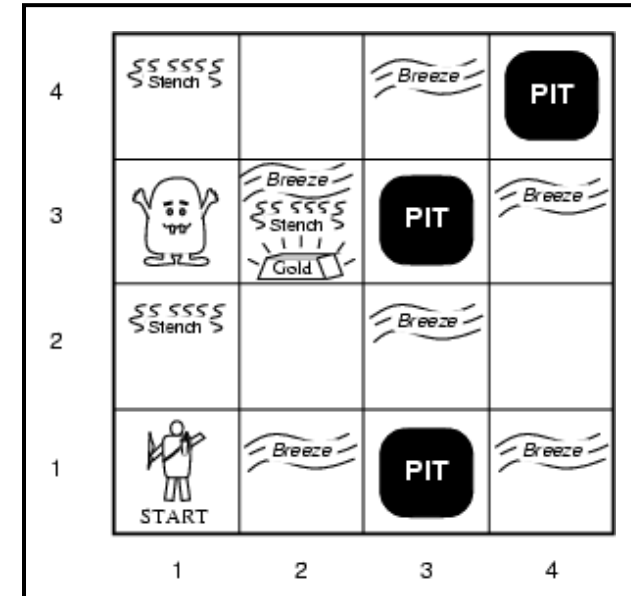
- Step 2b
  - $(P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1}$ 
    - $\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1}$
    - $(\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1}$
    - $(\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$

Now in CNF  
– as 2 rules

# Resolution Example: Back to Wumpus World

- Clause list (in CNF)
  - $R_1: \neg B_{1,1}$
  - $R_2: P_{1,2}$
  - $R_3: \neg B_{1,1} \vee P_{1,2} \vee P_{2,1}$
  - $R_4: \neg P_{1,2} \vee B_{1,1}$
  - $R_5: \neg P_{2,1} \vee B_{1,1}$
- Step 3 – Pick two rules and resolve via
 
$$\frac{(x_1 \vee \dots \vee x_m \vee x) \wedge (y_1 \vee \dots \vee y_k \vee \neg x)}{(x_1 \vee \dots \vee x_m \vee y_1 \vee \dots \vee y_k)}$$
- Step 3a – Reduce  $R_2$  and  $R_4$ 
  - $R_6: B_{1,1}$
- Step 3b – Reduce  $R_1$  and  $R_6$ 
  - $\emptyset$

1,2 <b>P?</b>	2,2
1,1 <b>A</b> <b>OK</b>	2,1



Proof by contradiction that  $KB \models \alpha$   
 i.e.,  $\alpha$  holds when  $KB$  holds; we can infer  $\alpha = \neg P_{1,2}$

# Where Does $\alpha$ Come From?

---

# Regarding the Query $\alpha$

- Inference algorithms show that we can infer  $\alpha$

- Where do we get  $\alpha$ ?

- Recall that Logical Agent program

**function** **KB-AGENT**(*percept*) **returns** an *action*

**persistent:** *KB*, a knowledge base

*t*, a counter, initially 0, indicating time

TELL(*KB*, MAKE-PERCEPT-SENTENCE(*percept*, *t*))

*action*  $\leftarrow$  ASK(*KB*, MAKE-ACTION-QUERY(*t*))

TELL(*KB*, MAKE-ACTION-SENTENCE(*action*, *t*))

*t*  $\leftarrow$  *t* + 1

**return** *action*

reasoning on what should be done

construct a sentence  
relating to an action to take

- Inference algorithms ( $\mathcal{A}$ ) assumes  $\alpha$  is given and decides if  $KB \models \alpha$
  - When discussing soundness and completeness of  $\mathcal{A}$ , we consider which among any given/input  $\alpha$  that will satisfy  $KB \models \alpha$

# Questions about the Lecture?

---

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
  - Specify a question
  - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)  
<https://archipelago.rocks/app/resend-invite/73861543030>



# Uncertainty



CS3243: Introduction to Artificial Intelligence – Lecture 9b

20 March 2023

# Logical Agents & Uncertainty

---

# Dealing with Uncertainty

- Example – Let  $A_t$  denote an autonomous taxi agent's action
  - $A_t$ : leave for airport  $t$  minutes before a flight
  - Will  $A_t$  get me to the airport on time?
- Sources of uncertainty
  - Partial observability (e.g., road state, other drivers' plans)
  - Noisy sensors (e.g., traffic reports, fuel sensor)
  - Uncertainty in action outcomes (e.g., flat tire, accident)
  - Complexity in modelling and predicting traffic (e.g., congestion)
- Logical agent will either
  1. Risk Falsehood – e.g.,  $A_{25}$  **will** get me there on time
  2. Reaches weaker conclusion – e.g.,  $A_{25}$  **will** get me there on time if
    - a. There is no accident on the bridge
    - b. It does not rain
    - c. I do not get a flat tire

Under logic (certainty), you may require  $A_{1440}$  to reach the airport on time (i.e., stay overnight)

Better to consider the probability of being on time with more reasonable  $t$  ...

# Probability

---

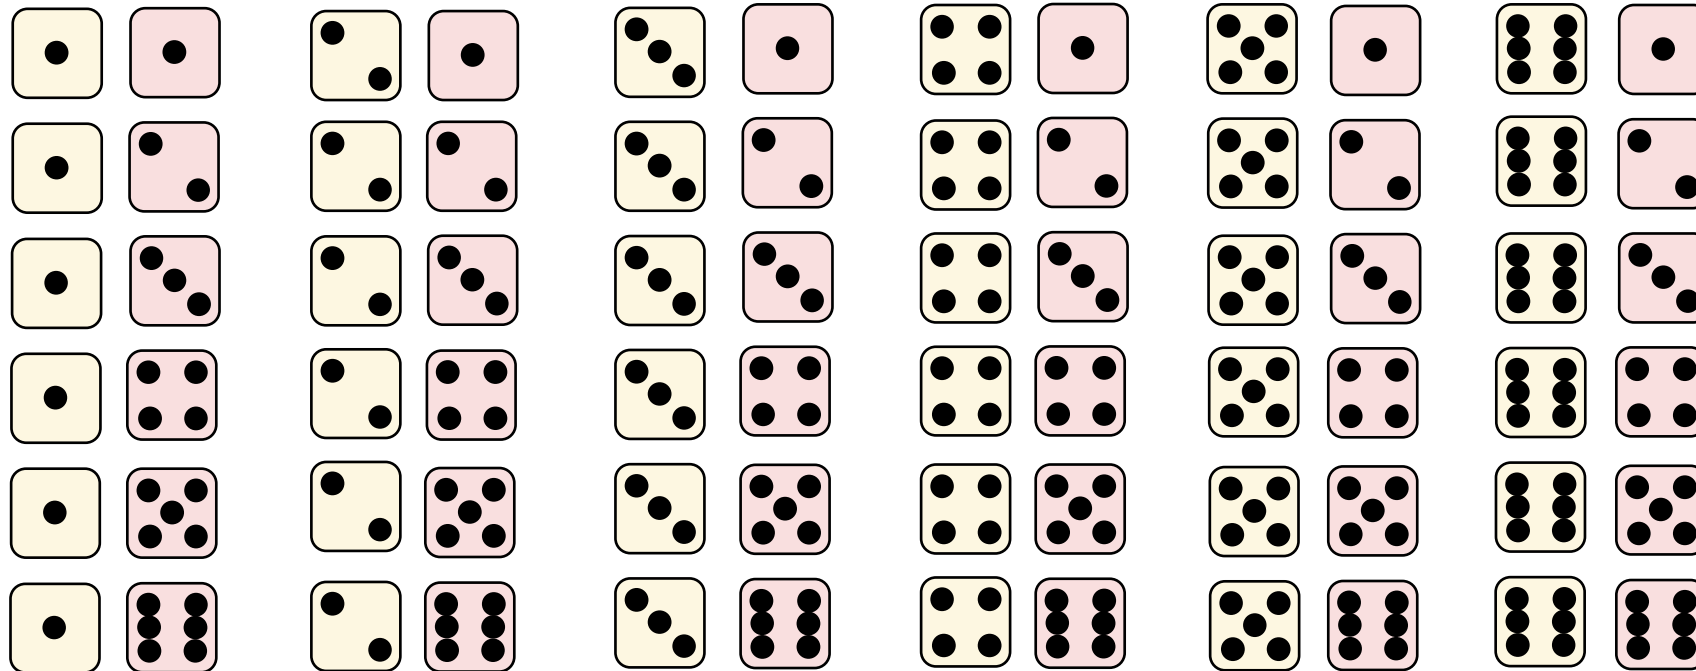
# Random Variables

---

- Random variable ( $X$ )
  - Quantifies an outcome of a random occurrence
    - e.g., outcome of a coin toss, die roll, or COVID-19 ART
- Domains ( $D_X$ )
  - Boolean: coin is either heads or tails (i.e., True or False)
  - Discrete: a die can have values  $\{1, \dots, 6\}$
- Events
  - Subsets of domains
    - e.g.,  $\text{Heads}(X)$ : coin flipped to heads
    - e.g.,  $\text{Even}(X)$ : die has value  $\in \{2, 4, 6\}$

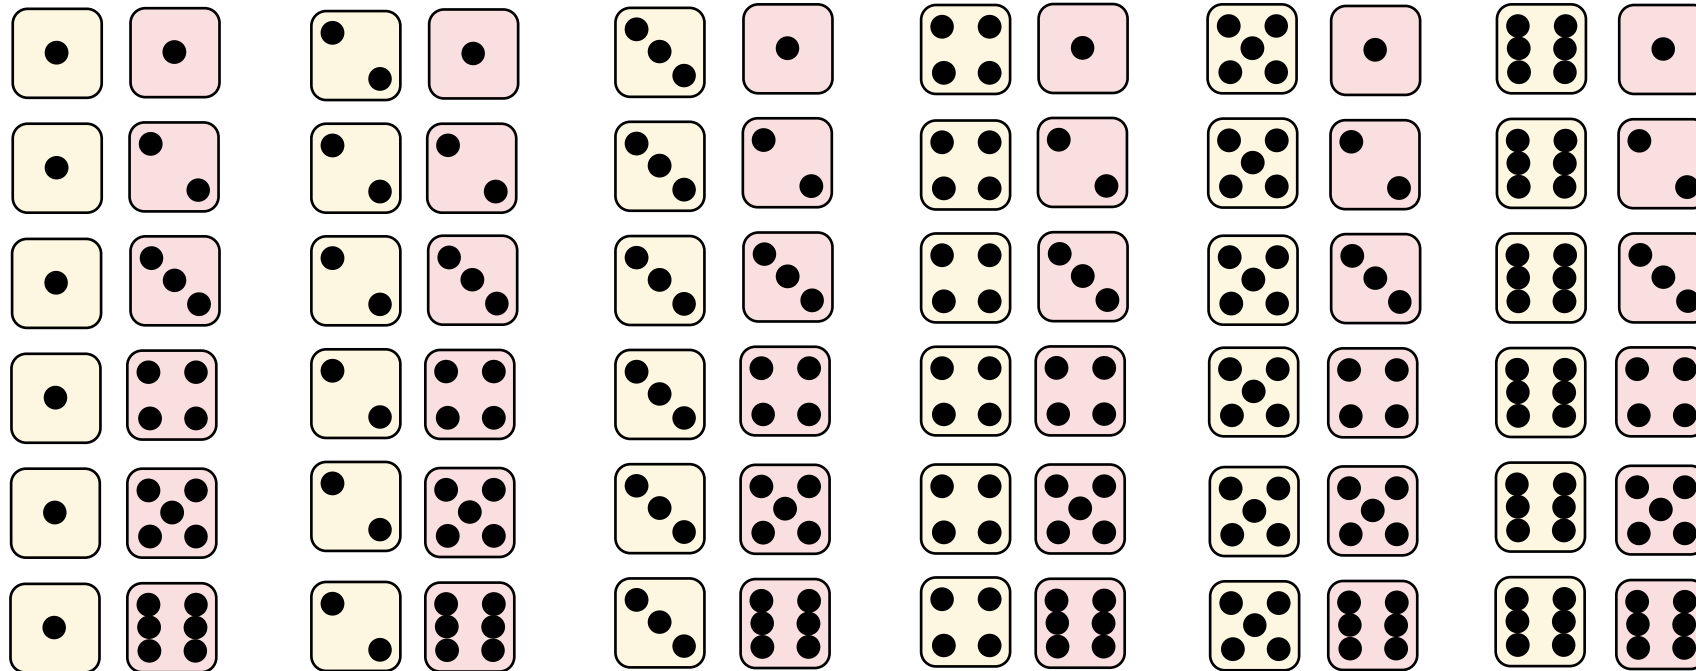
# Events

- Atomic / singleton event (possible world)
  - An assignment of a value to each random variable
- Example – we roll two different dice



# Events

- Yellow die =  $X_1$
- Pink die =  $X_2$
- Event:  $X_1 + X_2 = 8$



# Axioms of Probability

---

- Let  $X$  be a random variable with finite domain  $D_X$
- A probability distribution over  $D_X$  assigns a value  $p_X(v) \in [0,1]$  to every  $v \in D_X$  such that

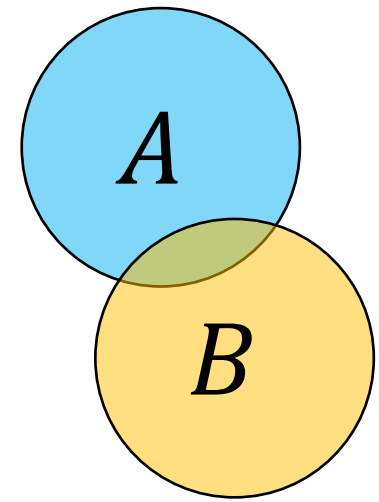
$$\sum_{v \in D_X} p_X(v) = 1$$

- For any event  $A \subseteq D_X$ , we have

$$\Pr[X \subseteq A] \equiv \Pr_X[A] = \sum_{v \in A} p_X(v)$$

- In particular

$$\Pr[A] + \Pr[B] = \Pr[A \cap B] + \Pr[A \cup B]$$





# Joint Probability

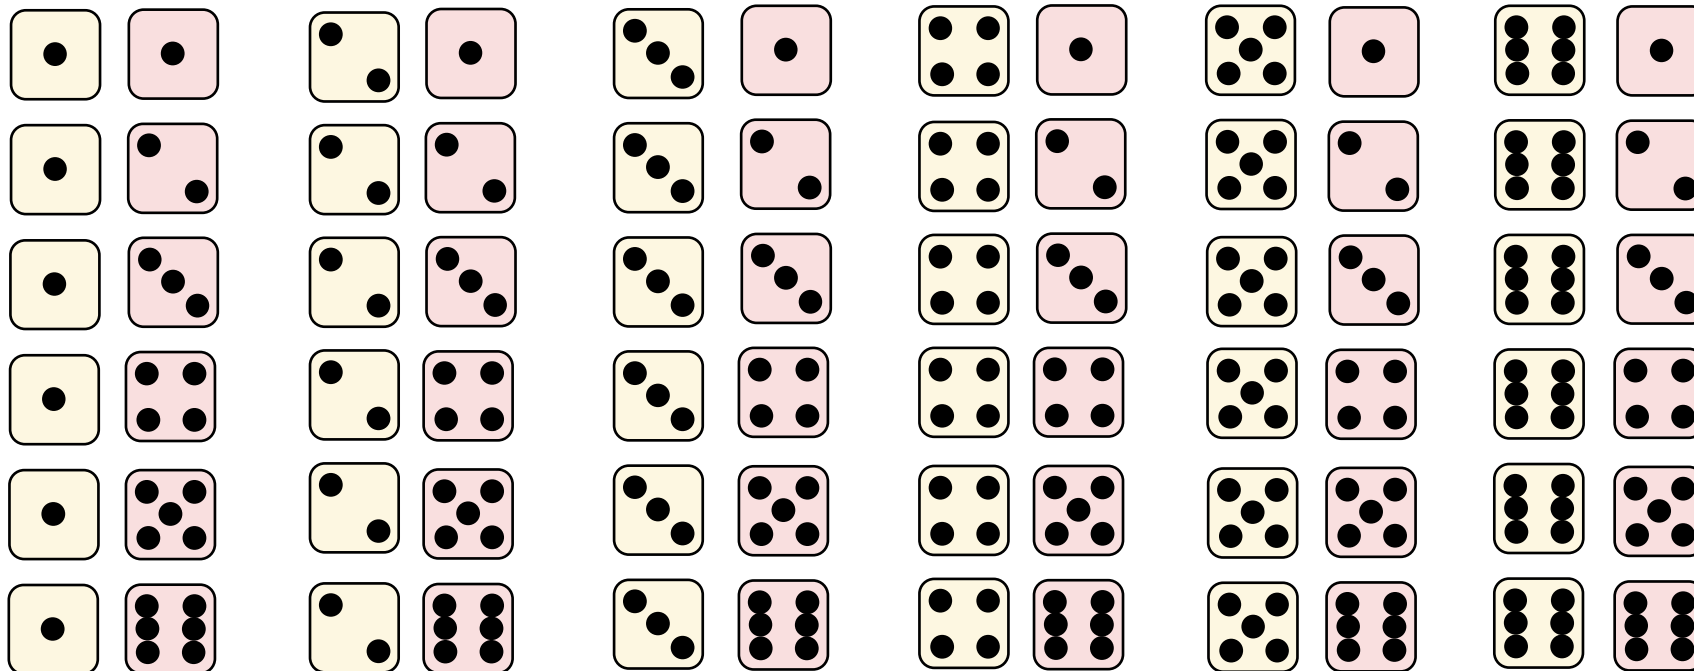
- Given two random variables  $X$  and  $Y$ 
  - The joint probability of an atomic event  $(x, y) \in D_X \times D_Y$  is  $p_{X,Y}(x, y) = \Pr[X = x \wedge Y = y]$
- In particular  $p_X(x) = \sum_{y \in D_Y} p_{X,Y}(x, y)$
- Example

Income (in SGD) / AGE	15-24	25-34	35-44	45-54	55-64	65+
< S\$2500	0.062	0.051	0.037	0.019	0.015	0.039
S\$2500 – S\$5000	0.078	0.068	0.061	0.057	0.031	0.053
> S\$5000	0.015	0.051	0.094	0.119	0.111	0.039

$$\Pr[Age = (25 - 34)] = 0.051 + 0.068 + 0.051 = 0.17$$

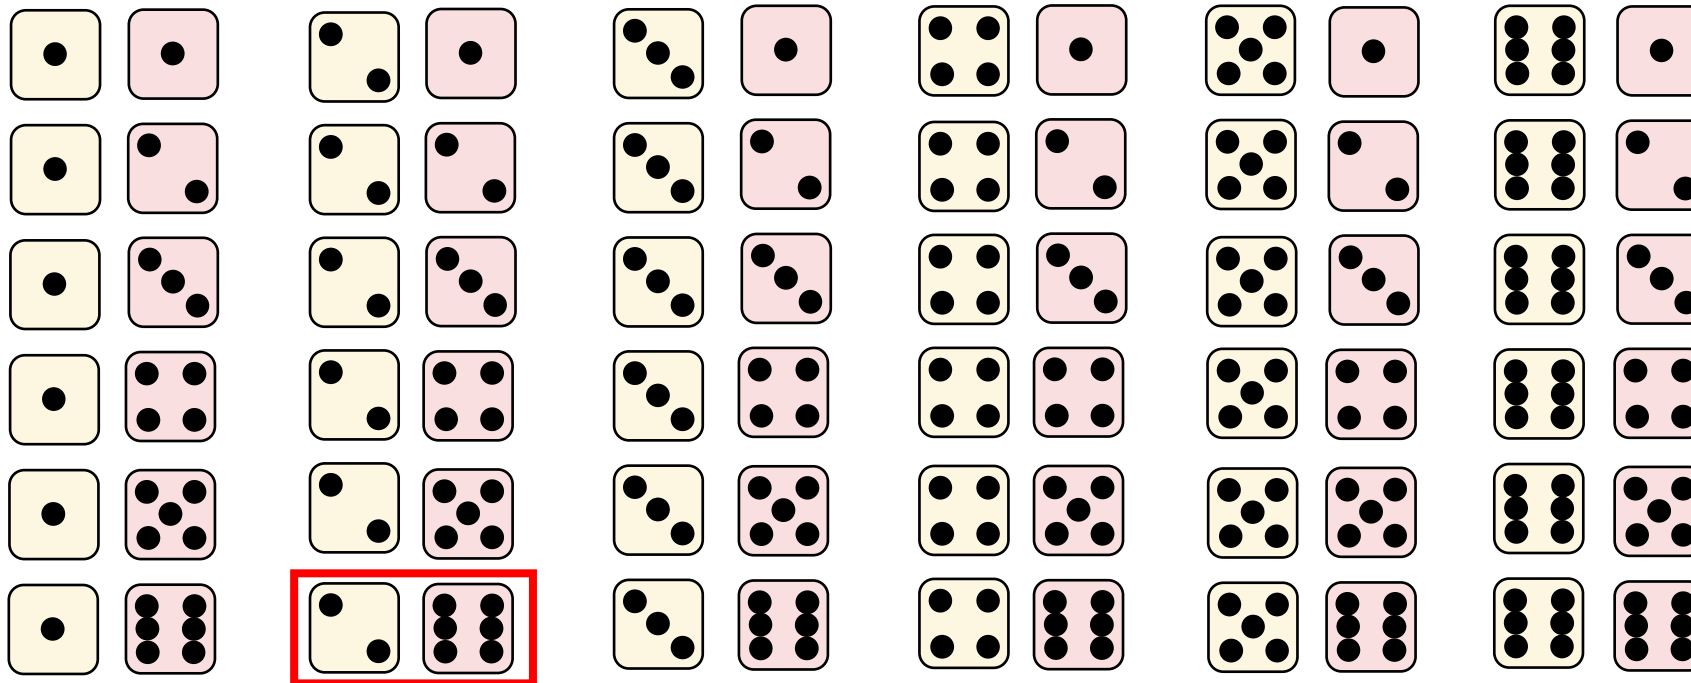
# Conditional Probability

- The probability that an event occurs, given that some other event occurs
- Example – rolling 2 dice;  $\Pr[X_1 = 2] = \frac{6}{36}$



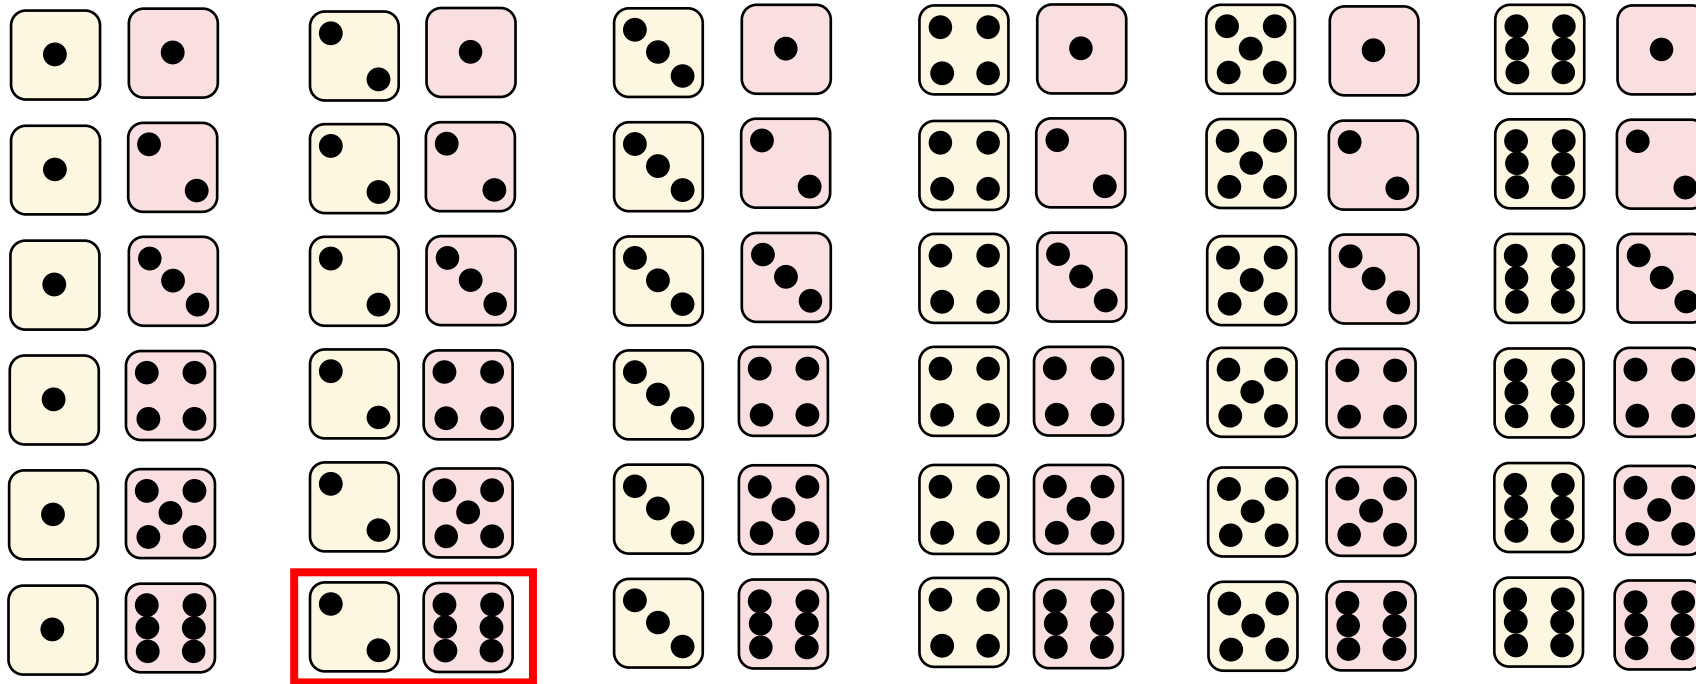
# Conditional Probability

- The probability that an event occurs, given that some other event occurs
- Example – rolling 2 dice;  $\Pr[X_1 = 2 \mid X_1 + X_2 = 8] = \frac{1}{5}$



# Conditional Probability

- The probability that an event occurs, given that some other event occurs
- Example – rolling 2 dice;  $\Pr[X_1 + X_2 = 8 \mid X_1 = 2] = \frac{1}{6}$



# Conditional Probabilities & Bayes Rule

- $\Pr[A | B] = \frac{\Pr[A \wedge B]}{\Pr[B]}$  assuming that  $\Pr[B] > 0$

Note:

$$\Pr[A | B] = \Pr[A \wedge B] / \Pr[B] \text{ --- (1)}$$

$$\Pr[B | A] = \Pr[B \wedge A] / \Pr[A] \text{ --- (2)}$$

From (2) and (3), we have:

$$\Pr[A \wedge B] = \Pr[B | A] \cdot \Pr[A] \text{ --- (4)}$$

Also, we know:

$$\Pr[A \wedge B] = \Pr[B \wedge A] \text{ --- (3)}$$

And thus from (4) and the definition above, we have Bayes Rule:

$$\Pr[A|B] = (\Pr[B|A] \cdot \Pr[A]) / \Pr[B]$$

- Bayes rule:  $\Pr[A | B] = \frac{\Pr[B|A] \Pr[A]}{\Pr[B]}$
- Example:  $\Pr[X_1 = 2 | X_1 + X_2 = 8] = ?$

$$= \frac{\overset{1/6}{\Pr[X_1 + X_2 = 8 | X_1 = 2]} \cdot \overset{1/6}{\Pr[X_1 = 2]}}{\underset{5/36}{\Pr[X_1 + X_2 = 8]}} = \frac{1}{5}$$

Next week, we will look at various applications of Bayes Rule

# Questions about the Lecture?

---

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
  - Specify a question
  - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)  
<https://archipelago.rocks/app/resend-invite/73861543030>