

Logical Agents: Knowledge Representation II

CS3243: Introduction to Artificial Intelligence – Lecture 9a

20 March 2023

Contents

1. Administrative Matters
2. Recap on Logical Agents
3. Theorem-Proving Methods
4. Resolution
5. Uncertainty & Recap on Probability

Reference: AIMA 4th Edition, Section 7.5

Administrative Matters

Upcoming...

- Deadlines

- TA7 (released last week)
 - *Due in your Week 10 tutorial session*
 - *Submit the a physical copy (more instructions on the Tutorial Worksheet)*
 - Prepare for the tutorial!
 - *Participation marks = 5%*
 - Project 2 (released Week 6)
 - *Was due yesterday, Sunday (19 March), 2359 hrs*
 - *Late penalties now apply*
 - Project 3 (released last week)
 - *Due Week 12 Sunday (9 April), 2359 hrs*
- Week 11: Final Content Lecture
 - Week 12: Lecture Slot = Project 3 Consultation (conducted over Zoom)

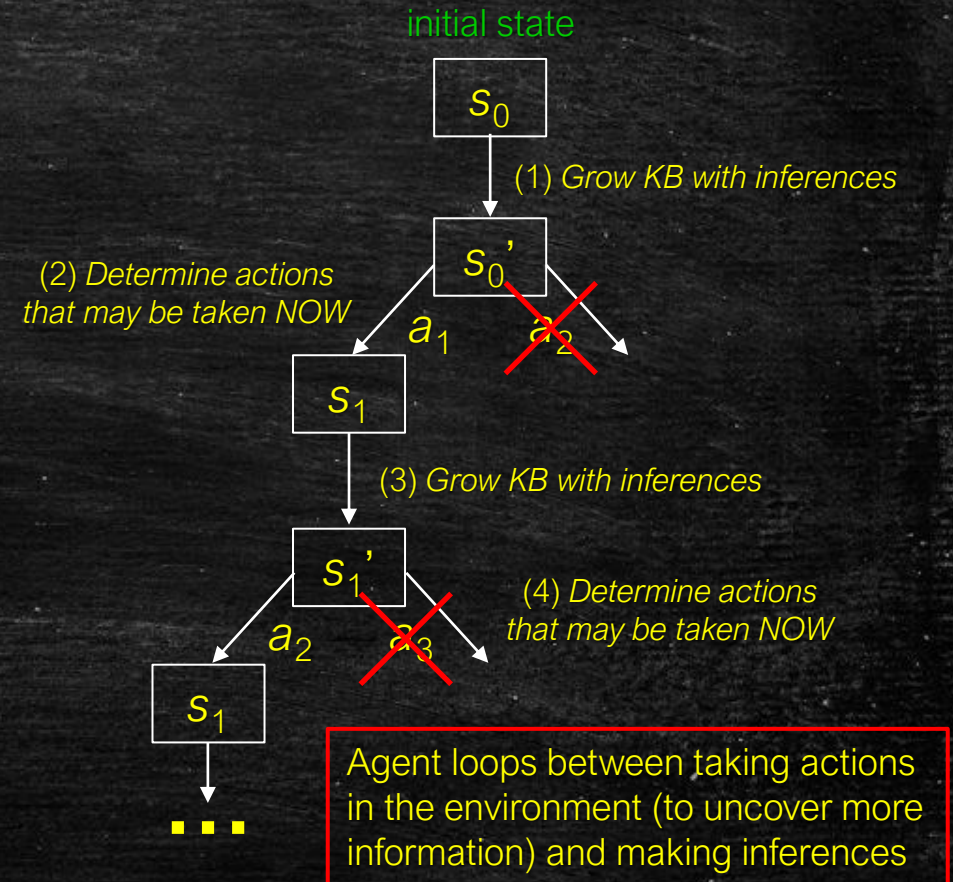
Recap on Logical Agents

Logical Agents

- Agent contains
 - Knowledge Base (KB)
 - Specified in some language (e.g., propositional logic)
 - Inference Engine (IE)
 - Determines sentences that will guide action choice, $\alpha_1, \alpha_2, \dots, \alpha_k$
 - Uses an algorithm that infers α_i such the $KB \models \alpha_i$
- General algorithm
 - Pre-populate KB with domain knowledge
 - Each time step t :
 - Update KB with percepts
 - Use IE to make inferences
 - Update KB with inferences
 - Select action based on inferences
 - Take action and update KB with new state (current truth value assignments)

Logical Agents

- Logical Agent *cannot* plan an entire path to goal
 - Environment is only **Partially Observable**
- Assume operation as follows
 - Make **inferences** about environment
 - Assume query (α) to action (a_i) mappings
 - e.g., reflex agent with conditions based on KB
 - Developer designing the agent must formulate these mappings
 - Suppose $\alpha \Rightarrow a_i$
 - Agent uses the **KB** & **IE** to determine if a query (α) may be inferred, and if so, takes associated action (a_i)
 - Chosen action (a_i) is **EXECUTED** in environment (as we are no longer planning)
 - Loops...



Making Inferences

- Entailment (\models)
 - $KB \models \alpha$ means that $M(KB) \subseteq M(\alpha)$
 - This says that all value assignments that satisfy the KB will also satisfy α
 - i.e., whenever KB is true, α is true
- Inference algorithm (\mathcal{A})
 - Sound: $(KB \vdash_{\mathcal{A}} \alpha) \Rightarrow (KB \models \alpha)$
 - \mathcal{A} only infers α that are valid
 - Complete: $(KB \models \alpha) \Rightarrow (KB \vdash_{\mathcal{A}} \alpha)$
 - \mathcal{A} is able to infer all valid α
- Inference algorithm example: Truth Table Enumeration
 - Construct entire truth table for KB and α
 - Check (via DFS) that $M(KB) \subseteq M(\alpha)$
 - i.e., every model of KB is a model of α

Truth Table Enumeration:

- Sound and Complete
- Time complexity $O(2^n)$
- Space complexity $O(n)$

Theorem Proving Methods

Proof Methods

- Model checking (special case of CSPs where domains are T/F)
 - Truth Table Enumeration (time complexity exponential in n)
 - Resolution (inference via proof)
- Applying inference rules (i.e., theorem proving)
 - Generate new sentences from old
 - Proof = sequential application of inference rules
 - Inference rules help deduce valid actions
 - Proof facilitates efficiency – ignores irrelevant propositions

Validity & Satisfiability

- A sentence α is **valid** if it is *true for ALL possible truth value assignments*
 - e.g., True, $A \vee \neg A$, $A \Rightarrow A$, $(A \wedge (A \Rightarrow B)) \Rightarrow B$
 - i.e., **tautologies**
- **Validity** is connected to **entailment** via the **Deduction Theorem**:
 - $(KB \models \alpha) \Leftrightarrow ((KB \Rightarrow \alpha) \text{ is valid})$
- A sentence is **satisfiable** if it is *true for SOME truth value assignment*
 - e.g., $A \vee B$, C
- A sentence is **unsatisfiable** if it is *true for NO truth value assignments*
 - e.g., $A \wedge \neg A$
 - i.e., **contradictions**
- **Satisfiability** is connected to **entailment** via the following:
 - $(KB \models \alpha) \Leftrightarrow ((KB \wedge \neg \alpha) \text{ is unsatisfiable})$
 - i.e., definition of **Proof by Contradiction**

i.e., a model exists
for that sentence

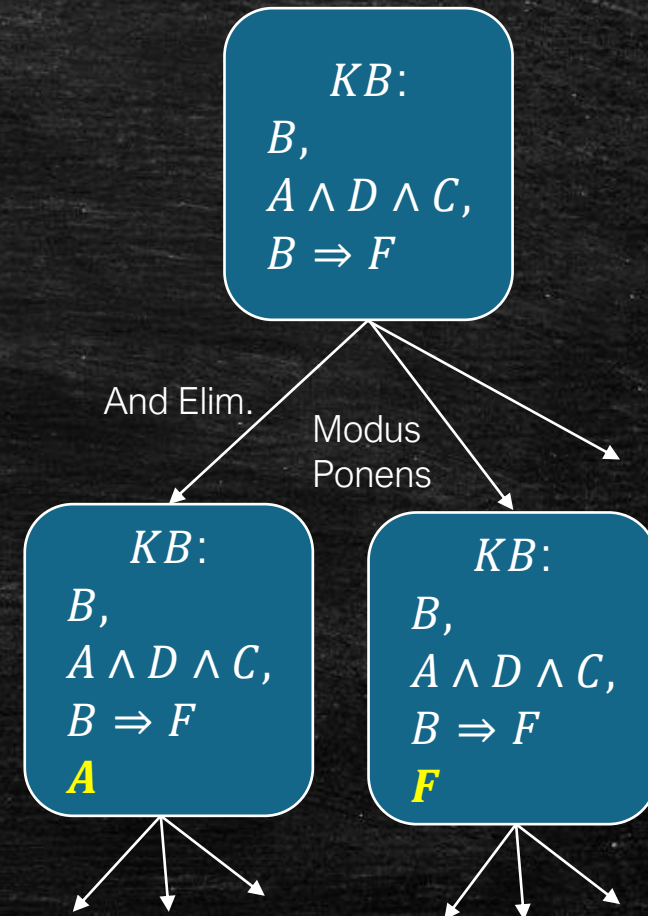
i.e., no model exists
for that sentence

- $A \Rightarrow B \equiv \neg A \vee B$
- $\neg(\neg A \vee B) \equiv A \wedge \neg B$
- Showing that $(A \wedge \neg B)$ is unsatisfiable shows that the negation, $\neg A \vee B$ is valid!

Inference Algorithms: Application of Inference Rules

- Search for more knowledge (growing the KB)
 - Equivalent to a search problem
 - States: Versions of the KB (e.g., initial state is initial KB)
 - Actions: Application of inference rules
 - Transition: Update KB with an inferred sentence (may not be target query α)
 - Goal: KB contains sentence to (dis)prove (e.g., given query α)
- Examples of inference rules
 - And-Elimination (AE): e.g., $a \wedge b \models a$; $a \wedge b \models b$
 - Modus Ponens (MP): e.g., $a \wedge (a \Rightarrow b) \models b$
 - Logical Equivalences: e.g., $(a \vee b) \models \neg(\neg a \wedge \neg b)$

How does this relate to Truth Table Enumeration?



Resolution

Resolution for Conjunctive Normal Form (CNF)

- CNF = conjunction of disjunctive sentences

- e.g., $(x_1 \vee \neg x_2) \wedge (x_2 \vee x_3 \vee \neg x_4)$

- Resolution

- Method of simplifying KB to prove entailment of query α

- Specifically

- Given KB: $R_1 \wedge R_2 \wedge \dots \wedge R_n$

- If a literal, x , appears in R_i and its negation, $\neg x$, appears in R_j , where $R_i, R_j \in \text{KB}$, it can be removed from both

$$\text{resolvent} \quad \frac{(x_1 \vee \dots \vee x_m \vee x) \wedge (y_1 \vee \dots \vee y_k \vee \neg x)}{(x_1 \vee \dots \vee x_m \vee y_1 \vee \dots \vee y_k)}$$

- Resolution under propositional logic

- Sound
 - Complete

KB: (P or x) and (Q or (not x))

α : (P or Q) must hold?

x=T, Q must be True for KB to hold

x=F, P must be True for KB to hold

So (P or Q) must hold

Verify with truth table as an exercise
(note: we want $M(\text{KB}) \subseteq M(\alpha)$)

Some Rules for Conversion to CNF

1. Convert $\alpha \Leftrightarrow \beta$ to $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$
2. Convert $\alpha \Rightarrow \beta$ to $\neg \alpha \vee \beta$
3. Expand \neg using De Morgan and double negation
 - a. Convert $\neg(\alpha \vee \beta)$ to $\neg \alpha \wedge \neg \beta$
 - b. Convert $\neg(\alpha \wedge \beta)$ to $\neg \alpha \vee \neg \beta$
 - c. Convert $\neg(\neg \alpha)$ to α
4. Convert $(\alpha \vee (\beta \wedge \gamma))$ to $(\alpha \vee \beta) \wedge (\alpha \vee \gamma)$

Each of these conversions produces two rules in the KB (the others just one)

Resolution Algorithm

- Utilises proof by contradiction – tries to show that $KB \wedge \neg\alpha$ is unsatisfiable

```
function PL-RESOLUTION( $KB, \alpha$ ) returns true or false
  inputs:  $KB$ , the knowledge base, a sentence in propositional logic
            $\alpha$ , the query, a sentence in propositional logic

   $clauses \leftarrow$  the set of clauses in the CNF representation of  $KB \wedge \neg\alpha$ 
   $new \leftarrow \{\}$ 
  while true do
    for each pair of clauses  $C_i, C_j$  in  $clauses$  do
       $resolvents \leftarrow$  PL-RESOLVE( $C_i, C_j$ )
      if  $resolvents$  contains the empty clause then return true
       $new \leftarrow new \cup resolvents$ 
    if  $new \subseteq clauses$  then return false
     $clauses \leftarrow clauses \cup new$ 
```

If cannot be resolved further and not empty clause – cannot infer α

What does an empty clause imply??

Suppose we have a KB as follows:

$$\frac{(x_1 \vee \dots \vee x_m \vee x) \wedge (y_1 \vee \dots \vee y_k \vee \neg x)}{(x_1 \vee \dots \vee x_m \vee y_1 \vee \dots \vee y_k)}$$

And the algorithm slowly removes literals:

$$\frac{(\cancel{x_1} \vee \dots \vee \cancel{x_m} \vee x) \wedge (\cancel{y_1} \vee \dots \vee \cancel{y_k} \vee \neg x)}{(\cancel{x_1} \vee \dots \vee \cancel{x_m} \vee \cancel{y_1} \vee \dots \vee \cancel{y_k})}$$

Eventually, there is nothing in the KB.

KB indicates the disjunction of no literals holds. A disjunction is True only when at least one literal is true. So, whole KB is False here – i.e., the query $\neg\alpha$ is unsatisfiable.

We may infer α (via proof by contradiction).

Resolution Algorithm

- Summary

- Make a **clause list** – i.e., copy of **KB** specified in **CNF** including negation of **query**, $\neg\alpha$
 - Use conversion rules to convert **KB** to **CNF**
- Repeatedly **resolve two clauses** from **clause list**
 - Add **resolvent** to **clause list**
- Keep doing this till **empty clause** found or **no more resolutions** possible
 - If **empty clause** then can infer α
 - If **no more resolutions** and not **empty clause** then cannot infer α

Why is Resolution under Propositional Logic Sound and Complete?

Soundness:

- Each resolvent is implied by generating clauses
- If \emptyset is found, then $(KB \wedge \neg\alpha)$ is unsatisfiable

Completeness:

- Based on the idea of resolution closure – set of all clauses derivable
- Not covered in CS3243
- Refer to AIMA 4th Edition pp. 228-229

Resolution Example

Resolution Example: Back to Wumpus World

- Assume that agent is at (1,1) in Wumpus World
 - And we wish to make inferences about a pit at (1,2)

KB

- $(B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge (\neg B_{1,1})$
- i.e., we know
 - R_a : rule for breezes
 - R_b : no breeze at (1,1)


α

- $\neg P_{1,2}$
- i.e., want to know if we can move to (1,2)

Can we infer α ?

- Use the resolution algorithm to determine if $(KB \models \alpha)$
- i.e., use $(KB \models \alpha) \Leftrightarrow ((KB \wedge \neg \alpha) \text{ is unsatisfiable})$

1,2	2,2
P?	
1,1	2,1
A OK	




Resolution Example: Back to Wumpus World

- Given
 - $KB = \neg B_{1,1} \wedge B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$
 - $\alpha = \neg P_{1,2}$
- Step 1 – Form clause list (over $KB \wedge \neg \alpha$)
 - $(\neg B_{1,1}) \wedge (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge (P_{1,2})$
- Step 2 – Convert clause list to CNF
 - $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$
 - $B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})$
 - $(P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1}$
- Step 2a
 - $B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})$
 - $\neg B_{1,1} \vee (P_{1,2} \vee P_{2,1})$
 - $\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}$

$(\neg B_{1,1}), (P_{1,2})$
already in
CNF (literals)

Now in CNF

1,2	2,2
P?	
1,1	2,1
A	
OK	



- Step 2b
 - $(P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1}$
 - $\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1}$
 - $(\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1}$
 - $(\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$

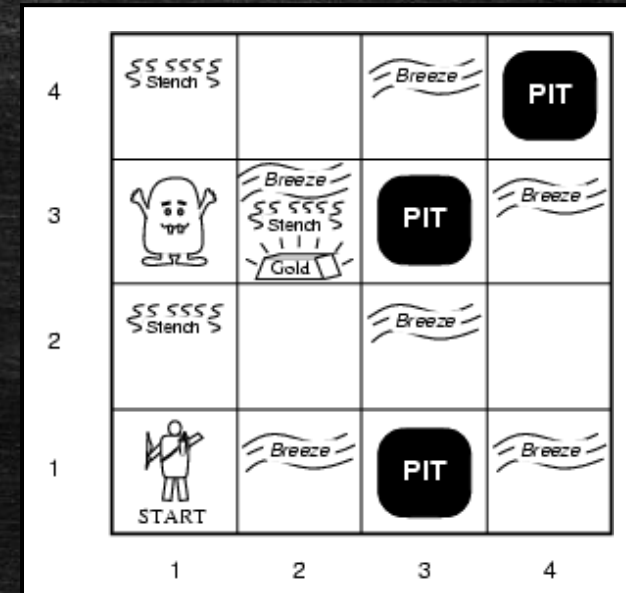
Now in CNF
– as 2 rules

Resolution Example: Back to Wumpus World

- Clause list (in CNF)
 - $R_1: \neg B_{1,1}$
 - $R_2: P_{1,2}$
 - $R_3: \neg B_{1,1} \vee P_{1,2} \vee P_{2,1}$
 - $R_4: \neg P_{1,2} \vee B_{1,1}$
 - $R_5: \neg P_{2,1} \vee B_{1,1}$
- Step 3 – Pick two rules and resolve via

$$\frac{(x_1 \vee \dots \vee x_m \vee x) \wedge (y_1 \vee \dots \vee y_k \vee \neg x)}{(x_1 \vee \dots \vee x_m \vee y_1 \vee \dots \vee y_k)}$$
- Step 3a – Reduce R_2 and R_4
 - $R_6: B_{1,1}$
- Step 3b – Reduce R_1 and R_6
 - \emptyset

1,2	2,2
P?	
1,1	2,1
A	
OK	



Proof by contradiction that $KB \models \alpha$
 i.e., α holds when KB holds; we can infer $\alpha = \neg P_{1,2}$

Where Does α Come From?

Regarding the Query α

- Inference algorithms show that we can infer α
- Where do we get α ?
 - Recall that Logical Agent program

```
function KB-AGENT(percept) returns an action  
  persistent: KB, a knowledge base  
               t, a counter, initially 0, indicating time  
  
  TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))  
  action  $\leftarrow$  ASK(KB, MAKE-ACTION-QUERY(t))  
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))  
  t  $\leftarrow$  t + 1  
  return action
```

reasoning on what should be done

construct a sentence
relating to an action to take

- Inference algorithms (\mathcal{A}) assumes α is given and decides if $KB \models \alpha$
- When discussing **soundness** and **completeness** of \mathcal{A} , we consider which among any given/input α that will satisfy $KB \models \alpha$

Questions about the Lecture?

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
 - Specify a question
 - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)
<https://archipelago.rocks/app/resend-invite/73861543030>

Uncertainty

CS3243: Introduction to Artificial Intelligence – Lecture 9b

20 March 2023

Logical Agents & Uncertainty

Dealing with Uncertainty

- Example – Let A_t denote an autonomous taxi agent's action
 - A_t : leave for airport t minutes before a flight
 - Will A_t get me to the airport on time?
- Sources of uncertainty
 - Partial observability (e.g., road state, other drivers' plans)
 - Noisy sensors (e.g., traffic reports, fuel sensor)
 - Uncertainty in action outcomes (e.g., flat tire, accident)
 - Complexity in modelling and predicting traffic (e.g., congestion)
- Logical agent will either
 1. Risk Falsehood – e.g., A_{25} **will** get me there on time
 2. Reaches weaker conclusion – e.g., A_{25} **will** get me there on time if
 - a. There is no accident on the bridge
 - b. It does not rain
 - c. I do not get a flat tire

Under logic (certainty), you may require A_{1440} to reach the airport on time (i.e., stay overnight)

Better to consider the probability of being on time with more reasonable t ...


Probability

Random Variables

- Random variable (X)
 - Quantifies an outcome of a random occurrence
 - e.g., outcome of a coin toss, die roll, or COVID-19 ART
- Domains (D_X)
 - Boolean: coin is either heads or tails (i.e., True or False)
 - Discrete: a die can have values $\{1, \dots, 6\}$
- Events
 - Subsets of domains
 - e.g., Heads(X): coin flipped to heads
 - e.g., Even(X): die has value $\in \{2, 4, 6\}$


Events

- Atomic / singleton event (possible world)
 - An assignment of a value to each random variable
- Example – we roll two different dice



Events

- Yellow die = X_1
- Pink die = X_2
- Event: $X_1 + X_2 = 8$



Axioms of Probability

- Let X be a random variable with finite domain D_X
- A probability distribution over D_X assigns a value $p_X(v) \in [0,1]$ to every $v \in D_X$ such that


$$\sum_{v \in D_X} p_X(v) = 1$$

- For any event $A \subseteq D_X$, we have

$$\Pr[X \subseteq A] \equiv \Pr_X[A] = \sum_{v \in A} p_X(v)$$

- In particular

$$\Pr[A] + \Pr[B] = \Pr[A \cap B] + \Pr[A \cup B]$$



Joint Probability


- Given two random variables X and Y
 - The joint probability of an atomic event $(x, y) \in D_X \times D_Y$ is $p_{X,Y}(x, y) = \Pr[X = x \wedge Y = y]$
- In particular $p_X(x) = \sum_{y \in D_Y} p_{X,Y}(x, y)$
- Example

Income (in SGD) / AGE	15-24	25-34	35-44	45-54	55-64	65+
< S\$2500	0.062	0.051	0.037	0.019	0.015	0.039
S\$2500 – S\$5000	0.078	0.068	0.061	0.057	0.031	0.053
> S\$5000	0.015	0.051	0.094	0.119	0.111	0.039

$$\Pr[Age = (25 - 34)] = 0.051 + 0.068 + 0.051 = 0.17$$


Conditional Probability

- The probability that an event occurs, given that some other event occurs
- Example – rolling 2 dice; $\Pr[X_1 = 2] = \frac{8}{36}$




Conditional Probability

- The probability that an event occurs, given that some other event occurs
- Example – rolling 2 dice; $\Pr[X_1 = 2 \mid X_1 + X_2 = 8] = \frac{1}{5}$



Conditional Probability

- The probability that an event occurs, given that some other event occurs
- Example – rolling 2 dice; $\Pr[X_1 + X_2 = 8 \mid X_1 = 2] = \frac{1}{6}$



Conditional Probabilities & Bayes Rule

- $\Pr[A | B] = \frac{\Pr[A \wedge B]}{\Pr[B]}$ assuming that $\Pr[B] > 0$

Note:

$$\Pr[A | B] = \Pr[A \wedge B] / \Pr[B] \text{ --- (1)}$$

$$\Pr[B | A] = \Pr[B \wedge A] / \Pr[A] \text{ --- (2)}$$

From (2) and (3), we have:

$$\Pr[A \wedge B] = \Pr[B | A] \cdot \Pr[A] \text{ --- (4)}$$

Also, we know:

$$\Pr[A \wedge B] = \Pr[B \wedge A] \text{ --- (3)}$$

And thus from (4) and the definition above, we have **Bayes Rule**:

$$\Pr[A|B] = (\Pr[B|A] \cdot \Pr[A]) / \Pr[B]$$

- Bayes rule: $\Pr[A | B] = \frac{\Pr[B|A] \Pr[A]}{\Pr[B]}$
- Example: $\Pr[X_1 = 2 | X_1 + X_2 = 8] = ?$

$$= \frac{\overset{1/6}{\Pr[X_1 + X_2 = 8 | X_1 = 2]} \cdot \overset{1/6}{\Pr[X_1 = 2]}}{\underset{5/36}{\Pr[X_1 + X_2 = 8]}} = \frac{1}{5}$$

Next week, we will look at various applications of Bayes Rule

Questions about the Lecture?

- Was anything unclear?
- Do you need to clarify anything?
- Ask on Archipelago
 - Specify a question
 - Upvote someone else's question



Invitation Link (Use NUS Email --- starts with E)
<https://archipelago.rocks/app/resend-invite/73861543030>